

Negative Perceptions of Outsourcing to Artificial Intelligence

Scott Claessens, Pierce Veitch, and Jim A.C. Everett

School of Psychology, University of Kent

Author Note

Scott Claessens  <https://orcid.org/0000-0002-3562-6981>

Pierce Veitch  <https://orcid.org/0009-0005-3364-7470>

Jim A.C. Everett  <https://orcid.org/0000-0003-2801-5426>

This pre-print is currently not yet peer-reviewed, and may differ from the final version.

Word count = 9119 words. Author roles were classified using the Contributor Role Taxonomy (CRediT; <https://credit.niso.org/>) as follows: Scott Claessens: conceptualization, data curation, formal analysis, investigation, methodology, visualization, writing, editing; Pierce Veitch: formal analysis, methodology, editing; Jim A.C. Everett: conceptualization, funding acquisition, methodology, supervision, editing

Correspondence concerning this article should be addressed to Jim A.C. Everett, School of Psychology, University of Kent, Keynes College, Canterbury CT2 7NP, UK, Email: j.a.c.everett@kent.ac.uk

Abstract

As artificial intelligence (AI) tools become increasingly integrated into daily life, people are beginning to outsource not only professional tasks but also socio-relational ones. Large language models like ChatGPT can generate wedding vows, speeches, and personal messages, raising questions about how individuals who use AI for such tasks are perceived by others. In this paper, we conduct five pre-registered studies with British participants ($N = 3,649$) to understand how people view those who outsource tasks to AI, and how this depends on how socio-relational the task is, whether AI is used as a tool or fully delegated to, and the acknowledgment of the AI use. We find negative perceptions of outsourcing, particularly for socio-relational tasks. We show that outsourcing makes us think more negatively about not only the person and their motivations, but also the outsourced work itself. Moreover, we provide insight into why this occurs: the reduced effort from outsourcing socio-relational tasks to AI signals that the output is less authentically one's own and that the person cares less about the task. Our research highlights the way that AI use shapes our perceptions of people, raising key philosophical questions about efficiency, authenticity, and social ties in a world filled with AI-mediated interactions.

Significance Statement

People form negative impressions of individuals who use artificial intelligence tools to complete tasks for them, especially interpersonal tasks like writing love letters or apology notes. This occurs even when they use AI as a tool instead of a replacement, are honest about it, and use AI because they care about doing the task well. We show that doing something oneself, even if AI could do it quicker and easier, signals that one is authentic and cares about the task. Our results highlight that people do not only care about whether something is done, but how it is done too.

Keywords: artificial intelligence, person perception, outsourcing, effort, trust

Negative Perceptions of Outsourcing to Artificial Intelligence

1 The widespread release of generative AI language models has transformed daily life,
2 offering the potential to perform a variety of tasks more efficiently and, in some cases, with
3 greater effectiveness than by doing them oneself. But as AI becomes more widely available,
4 people are not only using it to assist them with things like preparing dinner recipes, writing data
5 analysis code, and planning daily schedules. Increasingly, AI might be used beyond routine or
6 technical domains to instead assist in tasks that are more socio-relational in nature, like writing
7 wedding vows, apology notes, and love letters. Anecdotal evidence suggests that not only is
8 AI-outsourcing of this kind already happening, but that it potentially has serious effects on how
9 we judge others. In a recent Reddit post, a disgruntled newlywed tells the story of her husband
10 using ChatGPT to write his wedding vows, expressing her discomfort with outsourcing something
11 to AI that, to her, is deeply meaningful and a reflection of their love for one another ([miramar0, 2024](#)). Other reports describe people's negative reactions when they learn that their romantic
12 partner used ChatGPT to write them an apology note ([Tait, 2024](#)) or even to break up with them
13 ([Anderson, 2025](#)). Outsourcing tasks – especially socio-relational ones – to AI tools may be
14 efficient, but could have negative consequences for person perception.

16 There is nothing new, in principle, about outsourcing tasks. For hundreds of years,
17 personal assistants have organised daily schedules, recipe-books have provided meal plans, and
18 guidebooks have created travel itineraries. In the socio-relational domain, ghostwriters have
19 long-existed, and the internet is abound with professional paid services for writing wedding vows
20 and personal speeches. AI merely supercharges what is an ancient human impulse: the push to
21 reduce mental energy by outsourcing parts of our work onto people, books, tools, or systems. But
22 even if outsourcing is an old phenomenon, the rapid shift in availability and use of AI models has
23 fundamentally changed the ease with which people can outsource work, what kinds of tasks they
24 can outsource, and the way in which they can outsource. These new developments in society
25 mean that even as an old phenomenon in new clothes, there is much we still need to know about
26 outsourcing.

27 First, we need to know how people who outsource tasks to AI are perceived. We know that
28 people are increasingly using large language models (LLMs) for a wide variety of tasks
29 (Department for Science, Innovation & Technology, 2024). Due to their ubiquity, perhaps
30 outsourcing to LLMs might not lead to negative perceptions? We are sceptical. We know that
31 people dislike it when others “free ride” or reduce effort while benefiting from collective resources
32 (e.g. Cubitt et al., 2011; Kerr, 1983) and that people’s outputs are perceived as more valuable the
33 more effort was ostensibly put into them (Kruger et al., 2004). Moreover, exertion of effort is
34 deemed morally admirable and is rewarded, even in situations where effort does not directly
35 generate additional product, quality, or economic value, suggesting that effort itself is moralised
36 (Celniker et al., 2023). Given this, even if AI tools are widely available and pitched as improving
37 efficiency, the core social psychological processes are likely to remain: someone is expending less
38 effort to achieve a task, and people value effort. Indeed, some work shows that describing
39 someone as using AI for a relational task led to the perception they expended less effort and were
40 less satisfied with their relationship (Liu et al., 2024) and other unpublished work looking at
41 perceptions of people using AI to complete academic assignments finds that using AI leads to
42 more negative perceptions of moral character and suitability as a partner (Roth & Tissot, 2025).

43 Second, we need to know whether the *type* of task that people are outsourcing matters.
44 One might expect outsourcing to be perceived negatively regardless of the type of task being
45 outsourced – if effort is generally moralised, then the domain in which it is expended (or not)
46 should have little impact. However, there are reasons to expect differences between social tasks
47 like writing vows and non-social tasks like writing computer code. We know that different norms,
48 standards, and expectations can be applied to social and non-social tasks and exchanges (e.g. A. P.
49 Fiske, 1992; Heider, 1958; Malle, 2022). Moreover, from a philosophical perspective, it often
50 matters not only *whether* something is done, but *how* it is done (Aristotle, 2009; Hursthouse &
51 Pettigrove, 2023; Stohr, 2006). An apology is not just about hearing someone say “I am sorry”,
52 but seeing genuine regret; a love letter is not just about hearing someone say “I love you”, but
53 seeing depth of emotion; and a bereavement letter is not just about hearing someone say “I am

54 sorry for your loss”, but seeing an understanding for the powerful human experience of loss.
55 There is, perhaps especially for social tasks, value not only in the outcome of doing something,
56 but the *process* too (Goodman, 2010). To understand any potential negative effects of outsourcing
57 to AI, we must therefore look at a broad range of non-social and social tasks, rather than draw
58 broad conclusions based on a few use cases.

59 Third, we need to know how different ways of outsourcing to AI influence negative
60 perceptions. Someone who “fully” outsources a task to AI by simply giving it a prompt and
61 copying the output word-for-word might be perceived very differently to someone who gives the
62 AI a prompt, revises the work accordingly, and finishes it themselves – using AI as a *collaborative*
63 *tool*, rather than as a replacement. Similarly, someone could deceive others about their use of AI
64 or be perfectly honest about it. While it seems reasonable to assume that “fully” outsourcing
65 would be perceived worse than using AI as a tool, and that not acknowledging AI use would be
66 perceived worse than being honest about it, it remains unclear how much this reduces negative
67 perceptions: if someone uses AI in the “best” way, by using it as a collaborative tool and being
68 open about this use, would they still suffer negative social consequences from doing so?

69 Fourth, we need to know how outsourcing to AI, in different kinds of tasks and in different
70 ways, may shape different *kinds* of social perceptions. People can judge others on separate
71 dimensions of warmth and competence (e.g. Abele et al., 2021; S. T. Fiske et al., 2007) as well as
72 on dimensions of morality and trustworthiness (Goodwin et al., 2014). It remains unclear how
73 outsourcing to AI might lead to differential character judgments across these different dimensions.

74 Fifth, we need to understand *why* outsourcing to AI, and therefore expending less effort,
75 might have these effects. Previous work has focused on how expending less effort leads to
76 negative perceptions of others (Celniker et al., 2023). But this raises the question of *why* effort is
77 seen as important and what exactly it is signalling to others, beyond one’s general cooperative
78 intent. It is possible that outsourcing leads to negative perceptions because the lack of effort spent
79 on the task signals something more fundamental about how authentic one is and how much one
80 cares about the task: when someone chooses to outsource a love letter to an AI, they might be

81 seen as valuing that love letter and what it represents less. It could be this second-step order of
82 perceptions that is the key driver of negative perceptions, especially for socio-relational tasks.

83 **Present Research**

84 In this paper, we build on classic social psychological work on character inferences from
85 reduced effort to understand how people view others who outsource different kinds of tasks, in
86 different ways, for different reasons, to AI. Across five pre-registered experiments with British
87 participants, we seek not only to understand how reduced effort through AI-outsourcing might
88 shape perceptions of others, but also to understand in more depth *why* it is that reduced effort has
89 the effect that it does.

90 In our initial pilot studies to motivate this work, we found that people who outsource a
91 range of tasks to AI or another person are perceived more negatively than people who complete
92 the tasks by themselves (see Supplementary Materials). In Study 1, we look at the effects of task
93 type, AI use, and honesty. We explore how people perceive others who outsource different kinds
94 of tasks with different levels of social relevance (e.g., from daily schedules, computer code and
95 dinner recipes to wedding vows, apology letters, or bereavement cards), manipulating whether
96 people use AI as a collaborative tool or “fully” outsource to AI and whether they are honest or
97 deceptive about their use of AI. After turning to look at perceptions of both outsourcers and the
98 outsourced work in Study 2, in Studies 3-5 we probe why outsourcing may have negative effects
99 on how we evaluate others. In Study 3, we test potential mechanisms of perceived effort and
100 authenticity by looking at how people evaluate others who either spend a lot or little time crafting
101 the AI prompts, and who either outsource to a generic or personalised AI. In Study 4, we test the
102 potential mechanism of perceived importance in the task by manipulating people’s reasons for
103 using AI – either because they wanted to save time or because they cared about the task and
104 thought that AI would improve their work. Finally, in Study 5, we bring these different potential
105 mechanisms together to explore the different pathways that influence the relationship between
106 outsourcing and negative perceptions, focusing on perceived effort, authenticity, and care in the
107 task.

Study 1**109 Methods****110 Ethical Approval**

111 Ethical approval was granted for all studies in this paper by the REDACTED Psychology
112 Research Ethics Panel. Participants in all studies provided informed consent and were debriefed
113 after the study.

114 Participants

115 We conducted a power simulation to determine our target sample size. The simulation
116 suggested that a sample size of 150 participants per condition (overall $n = 750$ for five conditions)
117 would be required to detect a small difference between conditions (Cohen's $d \approx 0.20$) with above
118 80% power.

119 We recruited a convenience sample of 800 participants from the United Kingdom through
120 the online platform Prolific (<https://www.prolific.com/>). After excluding participants who failed
121 our pre-treatment attention check, we were left with a final sample of 762 participants (438
122 female; 316 male; 4 non-binary / third gender; 4 undisclosed gender; mean age = 42.16 years).
123 78% of these participants reported having used ChatGPT before.

124 Design

125 We used a “control plus 2x2” between-subjects design. Participants were randomly
126 allocated to either the control condition, in which people in the scenarios complete the tasks
127 themselves, or one of four experimental conditions, in which people in the scenarios use AI to
128 complete the tasks. In the experimental conditions, we manipulated whether people in the
129 scenarios used AI as a collaborative tool or “fully” outsourced to AI, and whether people were
130 honest or deceptive about their use of AI. This resulted in five conditions overall: (i) the control
131 condition, (ii) the tool-honest condition, (iii) the tool-deception condition, (iv) the full-honest
132 condition, and (v) the full-deception condition.

133 **Procedure**

134 We presented participants with six scenarios. Each scenario described a person
135 completing a task, such as writing computer code or writing a love letter. The six tasks were
136 randomly drawn from a larger set of 16 tasks (see Supplementary Table 1 for the full list of tasks).

137 For each scenario, we first told participants:

- 138 • *Control condition*: “In order to complete this task, [the person] works on it by themselves
139 from start to finish.”
- 140 • *Tool outsourcing conditions*: “In order to complete this task, [the person] uses the AI tool
141 ChatGPT. They ask ChatGPT to provide ideas, inspiration, and feedback, but they edit and
142 rewrite the suggestions and finish the task themselves.”
- 143 • *Full outsourcing conditions*: “In order to complete this task, [the person] uses the AI tool
144 ChatGPT. They copy ChatGPT’s output word-for-word, rather than doing it themselves.”

145 We then told participants in the experimental conditions:

- 146 • *Honest conditions*: “After completing the task, [the person] is asked how they came up with
147 their ideas. [The person] acknowledges that they used ChatGPT as a tool / got ChatGPT to
148 do the task for them.”
- 149 • *Deception conditions*: “After completing the task, [the person] is asked how they came up
150 with their ideas. [The person] does not acknowledge that they used ChatGPT as a tool / got
151 ChatGPT to do the task for them.”

152 We then asked participants how well each of the following words described the person in
153 the scenario: competent, warm, moral, lazy, and trustworthy. Participants answered these
154 questions on 7-point Likert scales, ranging from “does not describe [the person] well” to
155 “describes [the person] extremely well”.

156 After the six scenarios, we asked participants several questions about the AI tool
157 ChatGPT, including their familiarity with ChatGPT, whether they had used ChatGPT before, how
158 frequently they used ChatGPT, and how trustworthy they thought ChatGPT was.

159 ***Pre-registration***

160 We pre-registered the study on the Open Science Framework

161 (https://osf.io/xhmzk/?view_only=a4da193574d7410ba4d2aa3945a28b05).

162 ***Statistical Analysis***

163 We fitted Bayesian multivariate multilevel cumulative-link ordinal models to the data
164 using the *brms* R package (Bürkner, 2017). We modelled each character evaluation – competence,
165 warmth, morality, laziness, and trustworthiness – as a separate response variable and included
166 fixed effects for conditions, varying intercepts for participants, and varying intercepts and slopes
167 for tasks. We used regularising priors for all parameters to impose conservatism on parameter
168 estimates. All models converged normally ($\hat{R} \leq 1.01$).

169 ***Transparency and Openness***

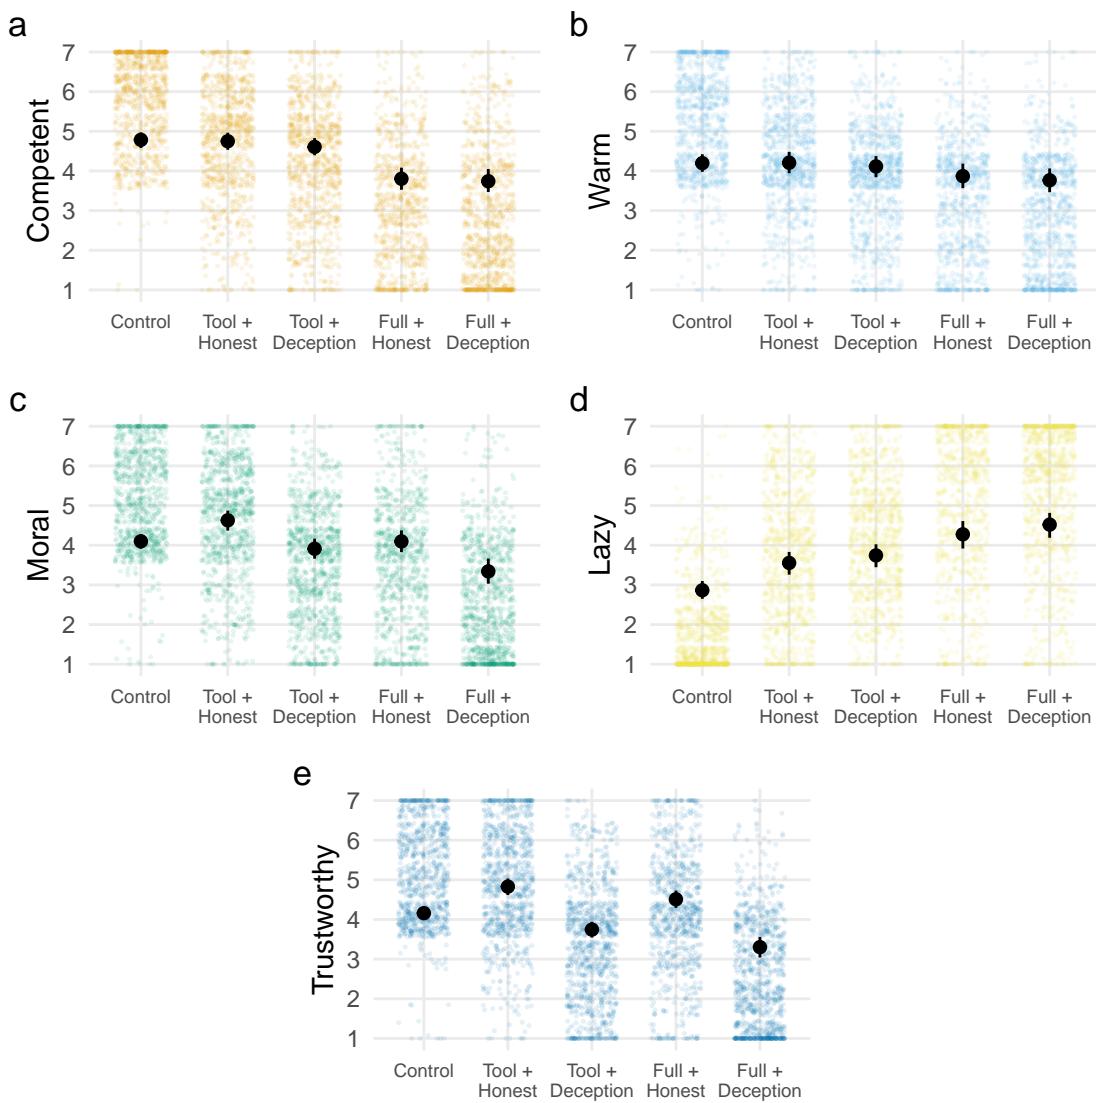
170 For all studies in this paper, we report how we determined our sample size, all data
171 exclusions, all manipulations, and all measures in the studies. All studies were pre-registered.
172 Analyses for all studies were conducted in R v4.4.2 (R Core Team, 2022). Visualisations were
173 produced using the *ggplot2* and *patchwork* packages (Pedersen, 2025; Wickham, 2016). The
174 manuscript was reproducibly generated using the *targets* package (Landau, 2021) and *quarto*
175 (Allaire et al., 2024). All data and code to reproduce the analyses and figures in this paper can be
176 found here: https://osf.io/xhmzk/?view_only=a4da193574d7410ba4d2aa3945a28b05

177 **Results**

178 We first looked at the overall results averaging over tasks. Across all five character
179 evaluations, we found that fully outsourcing to AI (i.e., copying the AI output verbatim) was
180 perceived more negatively than using AI as a collaborative tool (Figure 1; Table 1). By contrast,
181 we found that deception about AI usage had specific negative effects on perceptions of morality
182 and trustworthiness: people who did not acknowledge their use of AI were perceived as less moral
183 and less trustworthy. We did not find any interaction effects between full outsourcing and
184 deception.

Figure 1

Overall Character Evaluations in Study 1



Note. Participants in the control condition and four AI outsourcing conditions evaluated people in the scenarios on (a) competence, (b) warmth, (c) morality, (d) laziness, and (e) trustworthiness. Coloured points represent participant responses to the questions, jittered for easier viewing. Black points are estimated marginal means from the fitted model, pooling over participants and tasks. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.

Table 1*Overall Pairwise Contrasts in Study 1*

	Response				
	Competent	Warm	Moral	Lazy	Trustworthy
Comparison to control					
Tool Honest - Control	-0.04 [-0.25 0.19]	0.01 [-0.22 0.25]	0.53 [0.28 0.76]	0.69 [0.37 1.01]	0.67 [0.45 0.88]
Tool Deception - Control	-0.18 [-0.42 0.07]	-0.09 [-0.31 0.15]	-0.19 [-0.41 0.05]	0.88 [0.55 1.20]	-0.42 [-0.62 -0.20]
Full Honest - Control	-0.98 [-1.25 -0.71]	-0.33 [-0.59 -0.07]	-0.01 [-0.27 0.27]	1.42 [1.03 1.77]	0.35 [0.13 0.58]
Full Deception - Control	-1.05 [-1.33 -0.75]	-0.43 [-0.69 -0.16]	-0.76 [-1.04 -0.47]	1.66 [1.28 2.00]	-0.86 [-1.11 -0.59]
Effect of full outsourcing					
Full Honest - Tool Honest	-0.94 [-1.26 -0.63]	-0.34 [-0.64 -0.03]	-0.53 [-0.87 -0.20]	0.72 [0.29 1.14]	-0.32 [-0.60 -0.03]
Full Deception - Tool Deception	-0.87 [-1.18 -0.52]	-0.36 [-0.66 -0.02]	-0.57 [-0.91 -0.24]	0.78 [0.35 1.17]	-0.44 [-0.73 -0.15]
Effect of deception					
Tool Deception - Tool Honest	-0.14 [-0.42 0.13]	-0.09 [-0.39 0.20]	-0.72 [-1.01 -0.40]	0.19 [-0.19 0.58]	-1.08 [-1.35 -0.81]
Full Deception - Full Honest	-0.06 [-0.41 0.31]	-0.11 [-0.43 0.24]	-0.76 [-1.12 -0.38]	0.24 [-0.21 0.69]	-1.21 [-1.52 -0.89]
Interaction effect					
Interaction effect	0.08 [-0.36 0.53]	-0.01 [-0.45 0.43]	-0.03 [-0.51 0.42]	0.06 [-0.55 0.64]	-0.12 [-0.53 0.28]

Note. Numbers reflect differences in marginal means on a 7-point Likert scale, pooling over participants and tasks. The bottom row represents the interaction between full outsourcing and deception (i.e., the difference between the differences in the rows above). Main numbers are posterior medians, numbers in the square brackets are 95% credible intervals.

The effects of outsourcing to AI varied across the different tasks, especially for perceptions of warmth and morality (Figure 2). For example, people who used AI for social tasks, such as writing an apology letter or a bereavement card, were perceived as less warm, less moral, and lazier compared to people who completed the task themselves. This was true even if the person used AI as a tool and was honest about their usage of AI. By contrast, we observed weaker effects of outsourcing for non-social tasks like writing computer code or planning a syllabus.

185

186

187

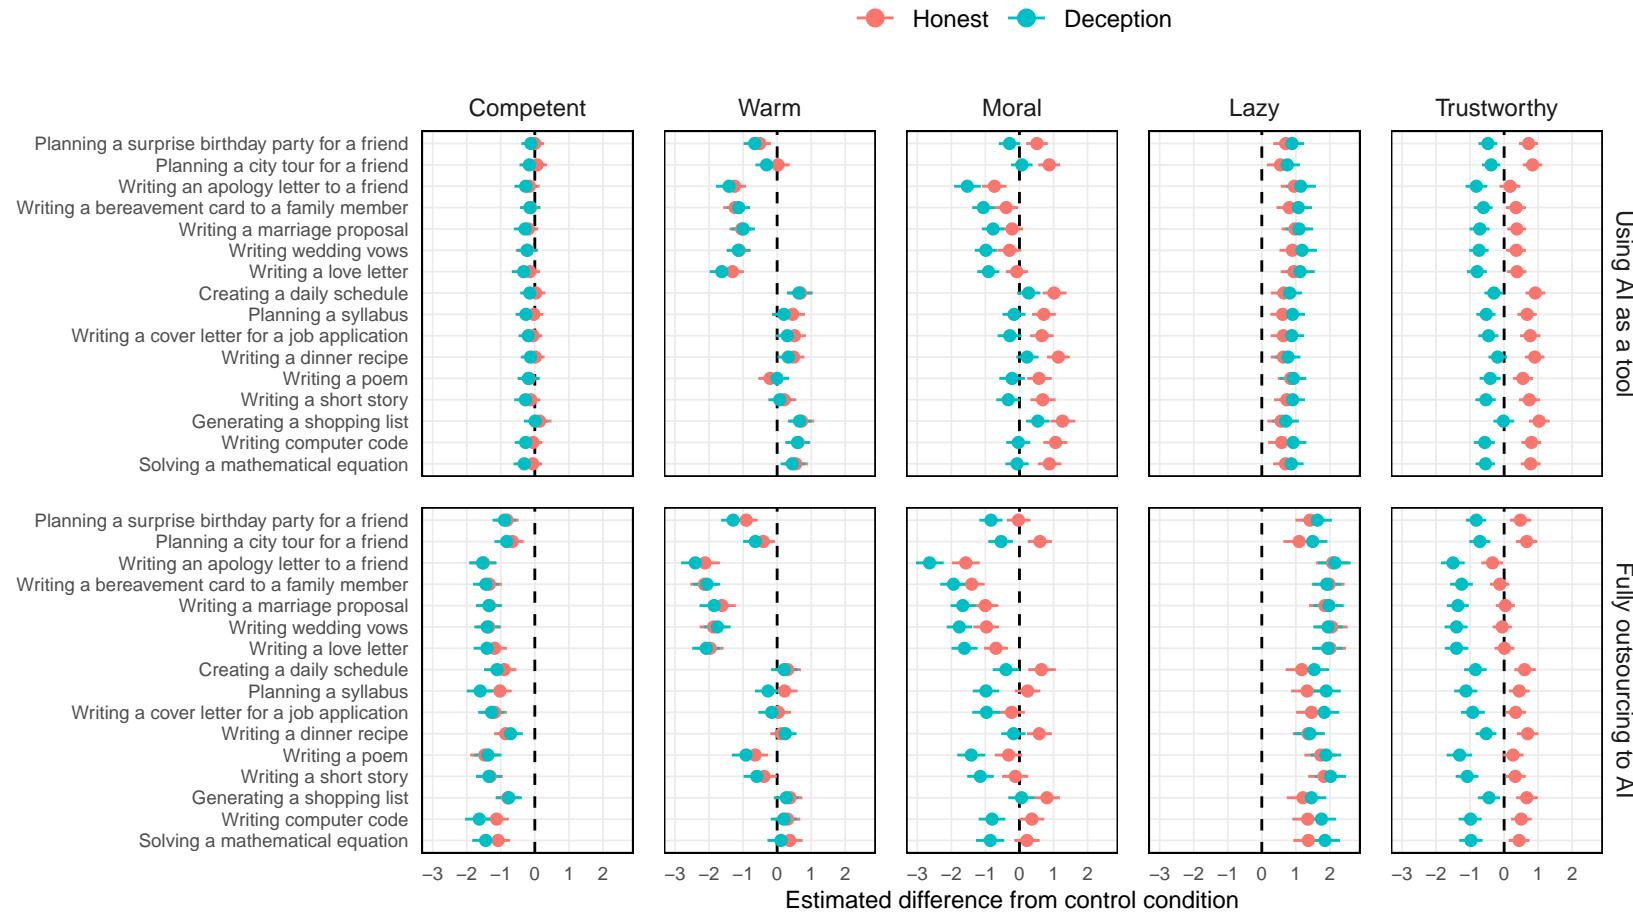
188

189

190

Figure 2

Variation in the Effects of Outsourcing across Tasks in Study 1



Note. Tasks are ordered from most social (top) to least social (bottom) according to ratings from a pilot study. Point ranges are differences in marginal means on a 7-point Likert scale for the honest conditions (red) and deception conditions (blue) compared to the control condition. Upper panels refer to the tool outsourcing conditions, and lower panels refer to the full outsourcing conditions. Points and ranges represent posterior medians and 95% credible intervals, respectively.

191 To determine the factors that predict variation across tasks, we incorporated ratings of
192 tasks from a pilot study (see Supplementary Materials for details). Participants were asked to rate
193 the different tasks on several features: whether the task is social, requires social skills, impacts
194 others, has important consequences, and requires effort. All of these features predicted stronger
195 causal effects of outsourcing to AI compared to the control (Supplementary Figures 1 and 2). In
196 other words, outsourcing to AI is perceived more negatively for tasks that have these features,
197 compared to tasks without these features.

198 Discussion

199 In Study 1, we looked at how people who outsourced to AI in different ways were
200 perceived across a broad range of social and non-social tasks. In line with our predictions, we
201 found that “fully” outsourcing to AI was perceived more negatively than using AI as a
202 collaborative tool, particularly for socio-relational tasks. We also found, predictably, that people
203 were seen as less moral and less trustworthy if they did not acknowledge their use of AI.
204 Importantly, though, we show that even using AI in the “best” way – only as a tool and being
205 honest about one’s usage – still led to negative social perceptions for the more socio-relational
206 tasks like writing a love letter, an apology, or wedding vows.

207 In Study 2, we investigate whether these negative perceptions extend to the work itself and
208 remain after seeing the output. It could be, for example, that someone is perceived badly for using
209 ChatGPT to write their bereavement card, but the writing itself is seen as equally well-written and
210 authentic, if not more so, than if the person had written the card themselves. Indeed, evidence
211 suggests that text generated by ChatGPT is rated as higher quality than human-written text (Noy
212 & Zhang, 2023). Moreover, it is possible that seeing appropriate output could mitigate negative
213 perceptions by highlighting how the AI can in fact perform the task well. We explored these
214 possibilities in Study 2.

Study 2**216 Methods****217 Participants**

218 We conducted a power simulation to determine our target sample size. The simulation
219 suggested that a sample size of 125 participants per condition (overall $n = 750$ for six conditions)
220 would be required to detect a small-to-medium difference between conditions (Cohen's $d \approx 0.40$)
221 with above 80% power.

222 We recruited a convenience sample of 800 participants from the United Kingdom through
223 Prolific. After excluding participants who failed our pre-treatment attention check, we were left
224 with a final sample of 766 participants (425 female; 337 male; 3 non-binary / third gender; 1
225 undisclosed gender; mean age = 41.93 years). 72% of these participants reported having used
226 ChatGPT before.

227 Design

228 We randomly allocated participants to one of six conditions in a 3x2 between-subjects
229 design. We manipulated the type of outsourcing: (i) no outsourcing control, (ii) using AI as a tool,
230 and (iii) fully outsourcing to AI. Here, in contrast to Study 1, we also explicitly manipulated
231 whether the task prompt was social or non-social.

232 Procedure

233 We told participants that they would read and evaluate a short piece of writing from
234 "another participant". In reality, we generated the writing using ChatGPT version 4.0. We asked
235 ChatGPT to generate a 300 word response to the prompt and to write convincingly like a real
236 human. We then edited the text to appear more human-like by, for example, removing classic AI
237 markers like dashes and concluding sentences and ensuring that the information was not too
238 generic, such that the writing could reasonably be attributed to both a human and AI.

239 The prompt for the piece of writing varied between conditions:

- 240 • *Social conditions*: "Please write a description of a close family member or friend,

241 explaining why they are special to you.”

- 242 • *Non-social conditions*: “Please write a short description of a book, TV show, or film of
243 your choice.”

244 We explained that the “other participant” was asked several questions about how they

245 produced their answer, including whether or not they used an AI tool like ChatGPT. We explained

246 that the participant was encouraged to be honest and told that they would be paid regardless. The

247 response from the “other participant” varied between conditions:

- 248 • *Control conditions*: “The participant reported that they did not use any AI tool like
249 ChatGPT. Instead, they worked on the response themselves from start to finish.”
- 250 • *Tool outsourcing conditions*: “The participant reported using ChatGPT to provide ideas,
251 inspiration, and feedback. The participant told us that they edited and rewrote ChatGPT’s
252 suggestions and finished writing the response themselves.”
- 253 • *Full outsourcing conditions*: “The participant reported using ChatGPT to complete the task.
254 The participant told us that they copied ChatGPT’s output word-for-word, rather than
255 producing the response themselves.”

256 Next, we presented participants with a randomly-chosen pre-generated essay answer to the

257 prompt (see Supplementary Tables 2 and 3 for full essay answers). In the social conditions, the

258 answer either referred to the participant’s father, their sister, or their best friend. In the non-social

259 conditions, the answer either referred to the book The Hobbit, the TV show Buffy the Vampire

260 Slayer, or the film Titanic. Reading times and responses to a follow-up comprehension question

261 suggested that participants read the essay answers in sufficient detail (see Supplementary Table 4).

262 Finally, we asked participants about their perceptions of the essay answer and the “other

263 participant”. We asked how well-written, meaningful, and authentic they thought the answer was

264 (7-point Likert scales), what letter grade they would give the answer (A-E), and how much they

265 would hypothetically reward the other participant for their work (from £0.00 to £1.00). We also

266 asked how well each of the following words described the other participant: competent, warm,
267 moral, lazy, and trustworthy (7-point Likert scales).

268 At the end of the study, we gave participants a manipulation check and asked them
269 whether they believed the manipulation. Almost all participants correctly reported the condition
270 that they were in and most participants stated that they believed the essay response was written in
271 the way we described, suggesting that the manipulation was successful (see Supplementary
272 Table 5). We also asked participants several questions about ChatGPT.

273 ***Pre-registration***

274 We pre-registered the study on the Open Science Framework¹.

275 ***Statistical Analysis***

276 We fitted two Bayesian multilevel models to the data. The first model was a multivariate
277 cumulative-link ordinal model including all Likert scales as separate response variables. The
278 second model was a zero-one-inflated-beta model applied specifically to the reward question,
279 which was a slider scale varying between 0 and 1. For both models, we included fixed effects for
280 the interaction between outsourcing type and task type and varying intercepts and slopes for essay
281 answers. We used regularising priors for all parameters to impose conservatism on parameter
282 estimates. All models converged normally ($\hat{R} \leq 1.01$).

283 **Results**

284 We first looked at character evaluations. We found that even when provided with concrete
285 output, people were still perceived more negatively across all character evaluations if they
286 outsourced the writing task to AI, either by using ChatGPT as a collaborative tool or by copying
287 the text from ChatGPT verbatim (Supplementary Figure 3; Supplementary Table 6). In contrast to
288 Study 1, however, we did not find any differences in character evaluations between the tool

¹ Due to a technical error with archiving this pre-registration on the Open Science Framework, the timestamp for the registration was lost. However, on our OSF project

(https://osf.io/xhmzk/?view_only=a4da193574d7410ba4d2aa3945a28b05), it is possible to view our pre-registration document file and its timestamped upload date.

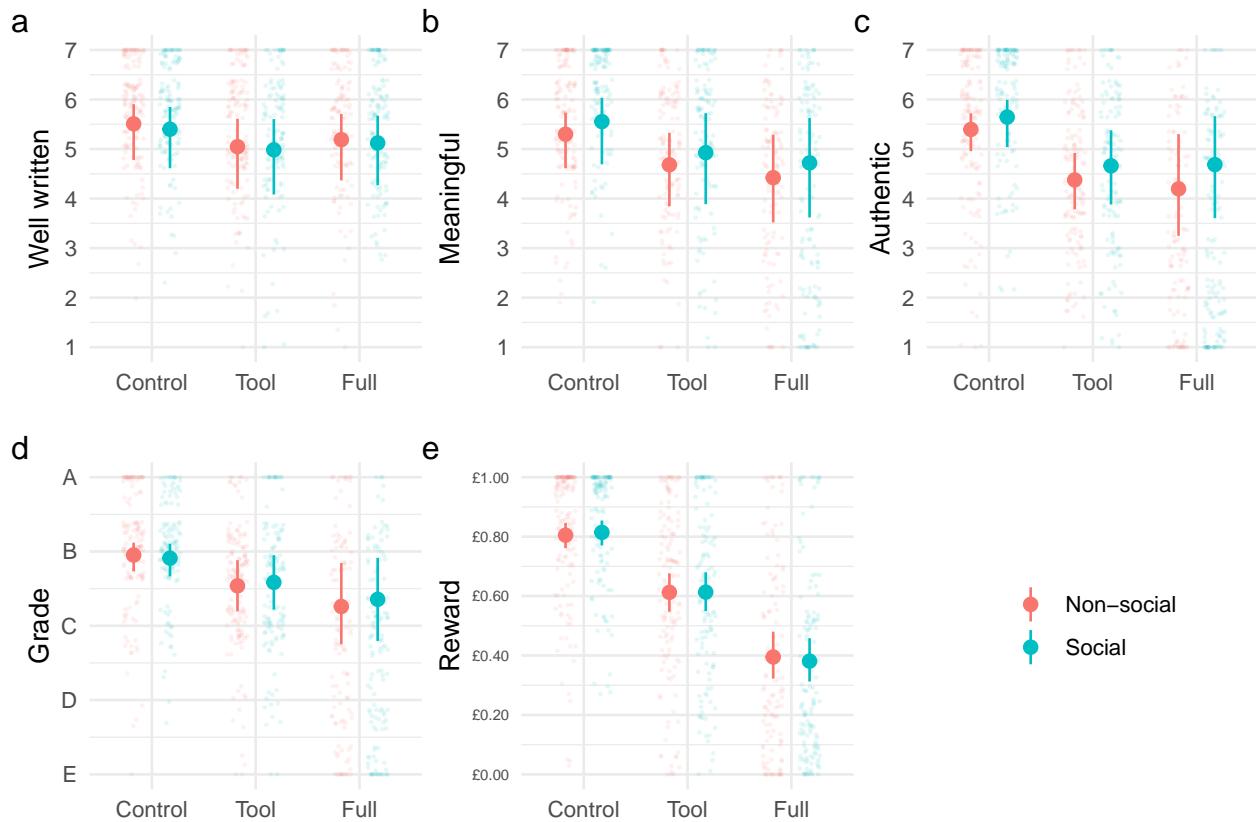
289 outsourcing and full outsourcing conditions. We did not find differences in character evaluations
290 between social and non-social tasks and did not find any interaction effects.

291 Turning to evaluations of the work itself, we found that the AI-outsourced work (either
292 outsourced by using AI as a collaborative tool or fully outsourced) was judged as being equally
293 well written to the work in the control condition (Figure 3; Table 2). This is in line with the
294 writing indeed being identical in all conditions. Interestingly, however, we found that essay
295 responses that were ostensibly generated using AI were perceived as less meaningful and less
296 authentic compared to essay responses ostensibly written by a human. Participants also marked
297 AI-generated essays with a lower grade and rewarded AI-generated essays with a lower
298 hypothetical monetary bonus. In contrast to Study 1, we did not find differences in perceptions of
299 the work between the tool outsourcing and full outsourcing conditions, except for the reward
300 question, where fully outsourced essays (i.e., essays copied verbatim from ChatGPT) were
301 rewarded £0.23 less than essays generated using AI as a collaborative tool. We did not find any
302 differences between social and non-social tasks and did not find any interaction effects.

303 **Discussion**

304 In Study 2, we turned to look at how people perceived both the outsourcer and the
305 outsourced work when given specific output in a social or non-social task that was described as
306 being produced independently by a person, produced by a person in collaboration with AI as a
307 tool, or outsourced in full to AI. We find that our results generalise from character judgments to
308 perceptions of the work itself: text purportedly generated using AI was perceived to be less
309 meaningful, less authentic, and less reward-worthy compared to the same text described as
310 human-generated.

311 Surprisingly, we found no differences in the effect of AI-outsourcing between social and
312 non-social tasks. This may be due to the particular tasks we chose. Writing *about* someone close
313 to you is not quite the same as writing something *for* someone close to you, as is the case with
314 wedding vows, love letters, and bereavement cards. We also found no differences between the tool
315 and full outsourcing conditions, aside from the lower rewards given to participants in the latter

Figure 3*Perceptions of the Work in Study 2*

Note. Participants in the control condition, the tool outsourcing condition, and the full outsourcing condition evaluated the essay response to the writing task. Participants rated whether the essay response was (a) well-written, (b) meaningful, and (c) authentic. Participants also (d) graded the work and (e) rewarded the work with a hypothetical monetary bonus. Jittered points represent participant responses to the questions, split by whether the writing task was a non-social task (red) or a social task (blue). Point ranges are estimated marginal means from the fitted model, pooling over essay answers. Points and line ranges represent posterior medians and 95% credible intervals, respectively.

Table 2

Pairwise Contrasts for Perceptions of the Work in Study 2

	Response				
	Well written	Meaningful	Authentic	Grade	Reward
Effect of outsourcing type					
Task type = Social					
Tool Social - Control Social	-0.42 [-0.84 0.03]	-0.64 [-1.11 -0.08]	-1.01 [-1.47 -0.35]	-0.33 [-0.60 -0.02]	-0.20 [-0.27 -0.13]
Full Social - Control Social	-0.28 [-0.63 0.09]	-0.84 [-1.55 -0.06]	-0.95 [-1.93 -0.07]	-0.55 [-1.07 -0.04]	-0.43 [-0.50 -0.36]
Full Social - Tool Social	0.14 [-0.39 0.67]	-0.21 [-1.13 0.74]	0.04 [-1.12 1.12]	-0.23 [-0.84 0.36]	-0.23 [-0.32 -0.14]
Task type = Non-social					
Tool Non-social - Control Non-social	-0.46 [-0.92 -0.02]	-0.62 [-1.09 -0.08]	-1.03 [-1.47 -0.49]	-0.42 [-0.71 -0.07]	-0.19 [-0.26 -0.12]
Full Non-social - Control Non-social	-0.32 [-0.70 0.03]	-0.90 [-1.56 -0.08]	-1.23 [-2.09 -0.12]	-0.71 [-1.16 -0.13]	-0.41 [-0.49 -0.32]
Full Non-social - Tool Non-social	0.14 [-0.39 0.71]	-0.28 [-1.06 0.62]	-0.21 [-1.19 0.97]	-0.29 [-0.86 0.34]	-0.22 [-0.31 -0.12]
Effect of task type					
Control Social - Control Non-social	-0.10 [-0.56 0.33]	0.28 [-0.36 0.73]	0.27 [-0.21 0.64]	-0.04 [-0.27 0.17]	0.01 [-0.05 0.06]
Tool Social - Tool Non-social	-0.06 [-0.67 0.56]	0.27 [-0.61 0.99]	0.29 [-0.41 0.95]	0.04 [-0.30 0.41]	0.00 [-0.09 0.09]
Full Social - Full Non-social	-0.07 [-0.64 0.51]	0.30 [-0.59 1.15]	0.48 [-0.37 1.38]	0.09 [-0.33 0.57]	-0.02 [-0.12 0.09]
Interaction effect					
Interaction: Tool - Control	0.05 [-0.39 0.51]	-0.02 [-0.52 0.56]	0.02 [-0.51 0.62]	0.09 [-0.22 0.39]	-0.01 [-0.10 0.08]
Interaction: Full - Control	0.04 [-0.34 0.46]	0.03 [-0.59 0.73]	0.23 [-0.50 1.02]	0.13 [-0.23 0.55]	-0.03 [-0.12 0.08]
Interaction: Full - Tool	0.00 [-0.58 0.55]	0.05 [-0.74 0.89]	0.21 [-0.66 1.12]	0.05 [-0.43 0.53]	-0.01 [-0.13 0.10]

Note. Numbers reflect differences in marginal means on either a 7-point Likert scale (well-written, meaningful, authentic), a 5-point ordinal grade scale (grade), or a 0-1 sliding scale (reward). Estimates are pooled over essay answers. The bottom rows represent the interactions between outsourcing type and task type. Main numbers are posterior medians, numbers in the square brackets are 95% credible intervals.

316 condition. It is possible that because the set-up described to participants was of another
 317 participant who was asked to produce work on Prolific and then admitted they used AI,
 318 participants saw any kind of AI use as violating an implicit contract between the survey requester
 319 and respondent and judged them negatively accordingly.

320 In Study 3, we turn to explore potential mechanisms driving our effects. We assume that
 321 effort may play a role, since perceived effort is often used as a signal of one's moral character

322 (Cubitt et al., 2011) and cooperative intent (Celniker et al., 2023). Study 2 also suggested a role of
323 authenticity: in line with work on the psychological importance of authenticity (Newman, 2019),
324 people who outsource to AI may be perceived as producing work that is less authentically their
325 own, leading to negative evaluations. To explore these potential mechanisms, we experimentally
326 manipulate (1) how much effort someone puts into the task and (2) whether they outsource the
327 task to a standard LLM like ChatGPT or a personalised LLM trained specifically on their own
328 prior writings (and so therefore producing work that is more authentically “theirs”). We expected
329 negative perceptions of outsourcing to be mitigated when the person uses a personalised LLM and
330 expends significant effort on formulating prompts for the AI.

331 **Study 3**

332 **Methods**

333 **Participants**

334 We used the same power estimate from Study 1 to determine our target sample size of $n =$
335 750 (150 participants in each of five conditions). We recruited a convenience sample of 802
336 participants from the United Kingdom through Prolific. After excluding participants who failed
337 our pre-treatment attention check, we were left with a final sample of 753 participants (462
338 female; 278 male; 9 non-binary / third gender; 4 undisclosed gender; mean age = 44.29 years).
339 74% of these participants reported having used ChatGPT before.

340 **Design**

341 We used the same “control plus 2x2” between-subjects design as in Study 1. In the
342 experimental conditions, we manipulated whether people in the scenarios used a standard or
343 personalised AI model, and whether people put more or less effort into the task. This resulted in
344 five conditions overall: (i) the control condition, (ii) the standard-low-effort condition, (iii) the
345 standard-high-effort condition, (iv) the personalised-low-effort condition, and (v) the
346 personalised-high-effort condition. Our authenticity manipulation was inspired by recent
347 psychological work looking at the credit-blame asymmetry in AI use (Earp et al., 2024), showing

348 that people receive more personal credit for their work when they use an AI model trained on their
349 own prior writings.

350 ***Procedure***

351 The procedure was mostly identical to Study 1 to allow us to explore effects across a range
352 of tasks, but we updated the study preamble and the presentation of the scenarios. For participants
353 in the personalised AI conditions, we expanded the study preamble to explain that personalised AI
354 models were trained on people's own prior writings and "tailored to each specific person and their
355 own thoughts, feelings, and values". Then in the scenarios, we told participants in the
356 experimental conditions:

- 357 • *Standard AI conditions*: "In order to complete this task, [the person] uses the AI tool
358 ChatGPT."
- 359 • *Personalised AI conditions*: "In order to complete this task, [the person] uses a personalised
360 AI tool."

361 We then told participants:

- 362 • *Low effort conditions*: "[The person] quickly gives the AI a rushed prompt and uses its first
363 output."
- 364 • *High effort conditions*: "[The person] carefully gives the AI several detailed prompts and,
365 after multiple rounds of changes, uses its resulting output."

366 At the end of the study, we asked participants to choose which of these was more authentic
367 and effortful, respectively. 94% of participants stated that the personalised AI was more authentic
368 and 99% of participants stated that giving the AI several detailed prompts was more effortful.
369 This suggests that even if participants might not have felt the output was meaningfully authentic in
370 the way that mattered (see Discussion), our participants agreed that using a personalised AI was at
371 least more authentic than using a generic one.

372 ***Pre-registration***

373 We pre-registered the study on the Open Science Framework
374 (https://osf.io/ac9g3/?view_only=912d9b57023d49baa87eea999574f0ce).

375 ***Statistical Analysis***

376 We fitted the same Bayesian multivariate multilevel cumulative-link ordinal models as in
377 Studies 1 and 2. All models converged normally ($\hat{R} \leq 1.01$).

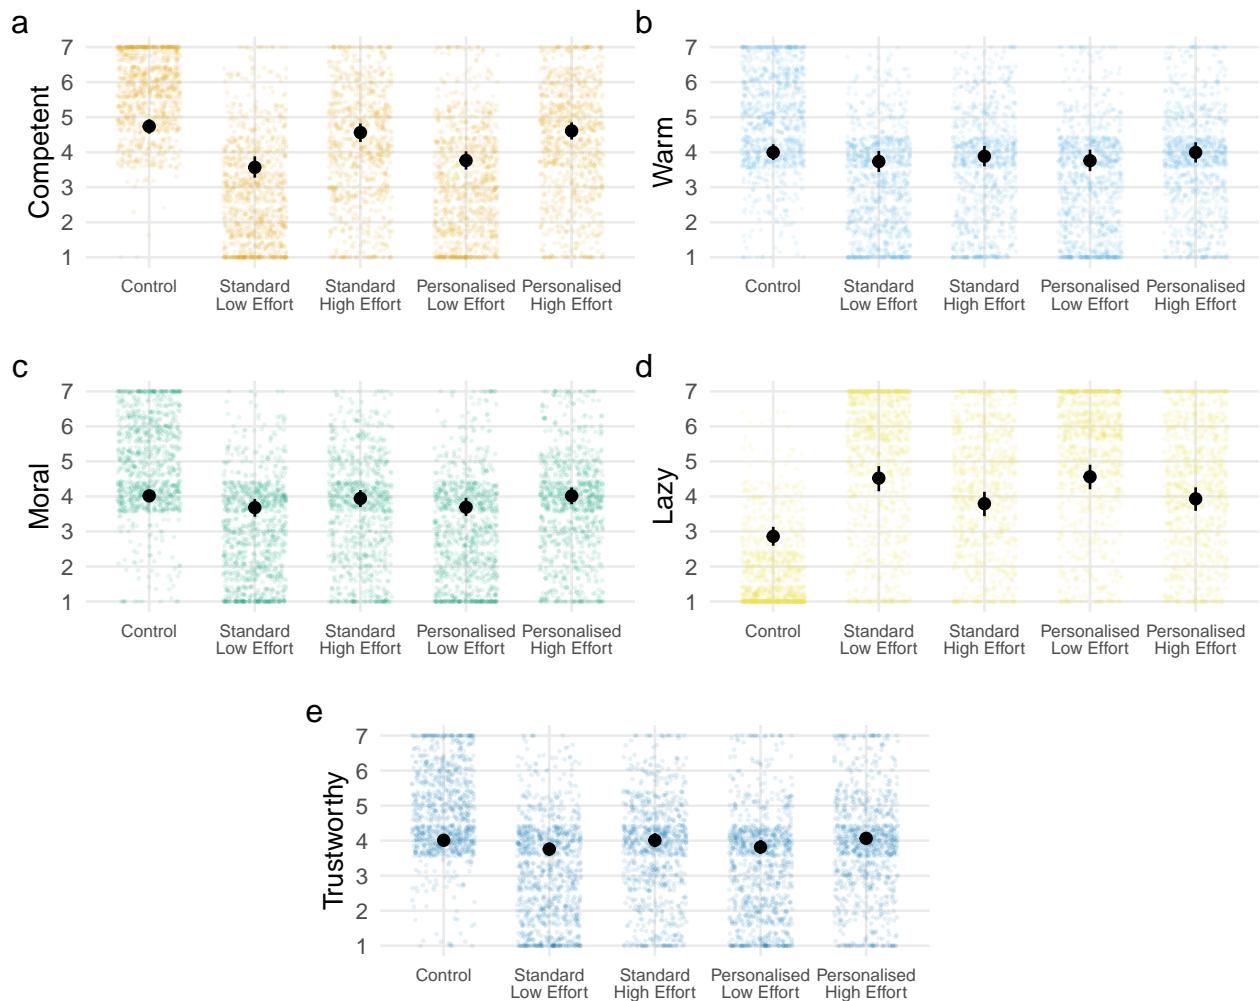
378 **Results**

379 We first looked across all the tasks. On average, we found that people who outsourced to
380 AI in a low effort way were perceived as less competent, less moral, lazier, and less trustworthy
381 than people who put more effort into their use of AI (Figure 4; Table 3). By contrast, we found
382 that character evaluations did not differ between people who used a standard AI model rather than
383 a personalised AI model. We also found no interaction effects between effort and the type of AI
384 used.

385 As in Study 1, the effects of outsourcing to AI varied across the different tasks, especially
386 for perceptions of warmth and morality (Figure 5). We again found that the negative causal effects
387 of outsourcing to AI were particularly strong for tasks that are social, require social skills, impact
388 others, have important consequences, and require effort (Supplementary Figures 4 and 5). Indeed,
389 for tasks like writing wedding vows or writing a love letter, outsourcing to a personalised AI in a
390 high effort way was still perceived more negatively than the control condition for all five character
391 dimensions.

Figure 4

Overall Character Evaluations in Study 3



Note. Participants in the control condition and four AI outsourcing conditions evaluated people in the scenarios on (a) competence, (b) warmth, (c) morality, (d) laziness, and (e) trustworthiness. Coloured points represent participant responses to the questions, jittered for easier viewing. Black points are estimated marginal means from the fitted model, pooling over participants and tasks. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.

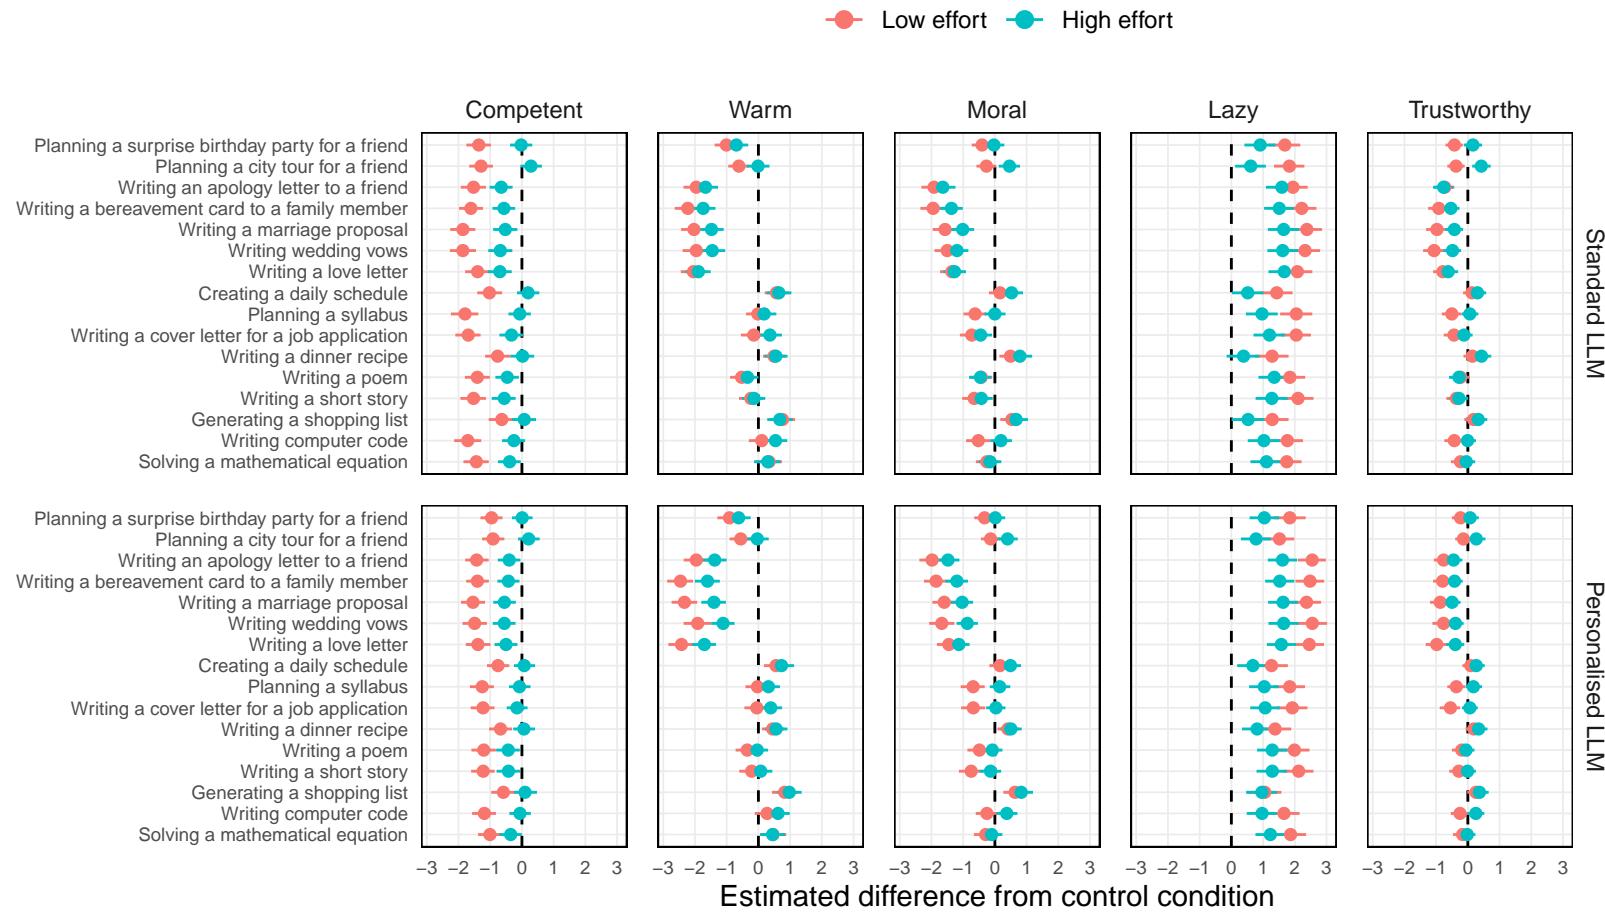
Table 3*Overall Pairwise Contrasts in Study 3*

	Response				
	Competent	Warm	Moral	Lazy	Trustworthy
Comparison to control					
Standard Low Effort - Control	-1.17 [-1.49 -0.85]	-0.26 [-0.53 0.00]	-0.34 [-0.58 -0.10]	1.67 [1.26 2.05]	-0.25 [-0.44 -0.06]
Standard High Effort - Control	-0.18 [-0.44 0.09]	-0.11 [-0.35 0.12]	-0.08 [-0.31 0.14]	0.94 [0.54 1.32]	0.00 [-0.19 0.19]
Personalised Low Effort - Control	-0.98 [-1.24 -0.70]	-0.24 [-0.50 0.02]	-0.33 [-0.56 -0.08]	1.71 [1.29 2.08]	-0.19 [-0.38 -0.01]
Personalised High Effort - Control	-0.13 [-0.38 0.13]	0.00 [-0.24 0.23]	0.00 [-0.22 0.23]	1.08 [0.69 1.44]	0.05 [-0.12 0.23]
Effect of AI type					
Standard Low Effort - Personalised Low Effort	-0.19 [-0.54 0.17]	-0.03 [-0.35 0.31]	-0.01 [-0.32 0.30]	-0.04 [-0.51 0.43]	-0.06 [-0.30 0.18]
Standard High Effort - Personalised High Effort	-0.04 [-0.36 0.27]	-0.11 [-0.41 0.20]	-0.07 [-0.36 0.22]	-0.14 [-0.61 0.34]	-0.05 [-0.28 0.17]
Effect of effort					
Standard Low Effort - Standard High Effort	-1.00 [-1.34 -0.62]	-0.15 [-0.47 0.17]	-0.26 [-0.56 0.04]	0.73 [0.25 1.21]	-0.25 [-0.50 -0.02]
Personalised Low Effort - Personalised High Effort	-0.84 [-1.16 -0.53]	-0.24 [-0.56 0.10]	-0.33 [-0.63 -0.03]	0.63 [0.15 1.08]	-0.25 [-0.48 -0.03]
Interaction effect					
Interaction effect	-0.15 [-0.61 0.32]	0.09 [-0.36 0.53]	0.06 [-0.36 0.49]	0.10 [-0.55 0.76]	0.00 [-0.32 0.32]

Note. Numbers reflect differences in marginal means on a 7-point Likert scale, pooling over participants and tasks. The bottom row represents the interaction between AI type and effort (i.e., the difference between the differences in the rows above). Main numbers are posterior medians, numbers in the square brackets are 95% credible intervals.

Figure 5

Variation in the Effects of Outsourcing across Tasks in Study 3



Note. Tasks are ordered from most social (top) to least social (bottom) according to ratings from a pilot study. Point ranges are differences in marginal means on a 7-point Likert scale for the low effort conditions (red) and high effort conditions (blue) compared to the control condition. Upper panels refer to the standard LLM conditions, and lower panels refer to the personalised LLM conditions. Points and ranges represent posterior medians and 95% credible intervals, respectively.

392 **Discussion**

393 In Study 3, we found that effort is an important mechanism by which outsourcing to AI
394 leads to negative character evaluations. People who engaged in effortless copying of the AI's first
395 output were perceived more negatively than people who spent time and effort crafting the AI's
396 output with multiple prompts. Nevertheless, for social tasks like writing wedding vows or love
397 letters, outsourcing to AI in a high effort way was still perceived more negatively than completing
398 the task oneself.

399 Interestingly, we found no effect of authenticity as proxied by the use of a personalised AI
400 that is trained on one's own prior writings compared to a standard AI like ChatGPT. This could
401 indicate that authenticity is not an important mechanism underlying the effect of outsourcing on
402 negative character evaluations. However, our specific manipulation may not have moved the
403 needle on authenticity enough to impact character evaluations. While previous work has found an
404 effect of personalised AI models on perceived credit (Earp et al., 2024), and the majority of
405 participants in our study stated that the personalised AI was more authentic than a standard model
406 like ChatGPT, it is possible that perceptions of *meaningful* authenticity in our study remained low
407 even with the personalised AI model. An AI could be perfectly trained on all apologies that a
408 person has ever written, but one might still think that a specific apology it then generates in a new
409 instance is not an *authentic* apology. Therefore, even if people were described as outsourcing to
410 an AI that was trained on their own writing and therefore personalised, participants still may not
411 have seen the specific output as being meaningfully authentic in the way that matters for character
412 judgments.

413 In Study 4, we turn to look at a third potential mechanism: a perceived lack of importance
414 attached to the task. When participants read about someone who outsources to AI in our studies,
415 they may be inferring that they simply did not care enough about the task – “If this was important
416 to them, they would do it themselves!”. To the extent that we especially want people to care about
417 their relationships with others – the kind of things demonstrated through love letters, apology
418 notes, and gift-giving – this could explain the particular negativity we see for social tasks

419 compared to tasks like writing daily schedules, recipes, or computer code. To test this, in Study 4,
420 we attempted to counteract inferences about care for the task by explicitly telling participants that
421 someone had a good reason for using AI: that they really cared about the task and used AI because
422 they wanted to get it right.

423 **Study 4**

424 **Methods**

425 **Participants**

426 We used the same power estimate from Study 1 to determine our target sample size of $n =$
427 750 (150 participants in each of five conditions). We recruited a convenience sample of 800
428 participants from the United Kingdom through Prolific. After excluding participants who failed
429 our pre-treatment attention check, we were left with a final sample of 758 participants (398
430 female; 346 male; 8 non-binary / third gender; 6 undisclosed gender; mean age = 41.72 years).
431 80% of these participants reported having used ChatGPT before.

432 **Design**

433 We used the same “control plus 2x2” between-subjects design as in Studies 1 and 3. In the
434 experimental conditions, we manipulated whether people in the scenarios used AI as a tool or
435 “fully” outsourced to AI, and whether people had bad or good reasons for using AI. This resulted
436 in five conditions overall: (i) the control condition, (ii) the tool-bad-reason condition, (iii) the
437 tool-good-reason condition, (iv) the full-bad-reason condition, and (v) the full-good-reason
438 condition.

439 **Procedure**

440 The procedure was mostly identical to Study 3, with two changes. First, we reduced the
441 number of tasks, focusing on eight tasks (four “social” tasks and four “non-social” tasks) that fit
442 with the manipulation of the updated design (since, for example, participants might find it difficult
443 to see how someone could deeply value a shopping list and want to get it right). Second, we
444 updated the presentation of the scenarios. We told participants in the experimental conditions:

- 445 • *Bad reason conditions*: “Because they are really short on time and want to complete the
446 task quickly, [the person] uses the AI tool ChatGPT.”
447 • *Good reason conditions*: “Because this task is really important to them and they want to
448 make sure they get it right, [the person] uses the AI tool ChatGPT.”

449 We then told participants:

- 450 • *Tool outsourcing conditions*: “[The person] asks ChatGPT to provide ideas, inspiration, and
451 feedback, but they edit and rewrite the suggestions and finish the task themselves.”
452 • *Full outsourcing conditions*: “[The person] copies ChatGPT’s output word-for-word, rather
453 than doing it themselves.”

454 In addition to the five character evaluations, on each page we also asked participants, on a
455 7-point Likert scale, how much they thought the person cared about the task.

456 ***Pre-registration***

457 We pre-registered the study on the Open Science Framework

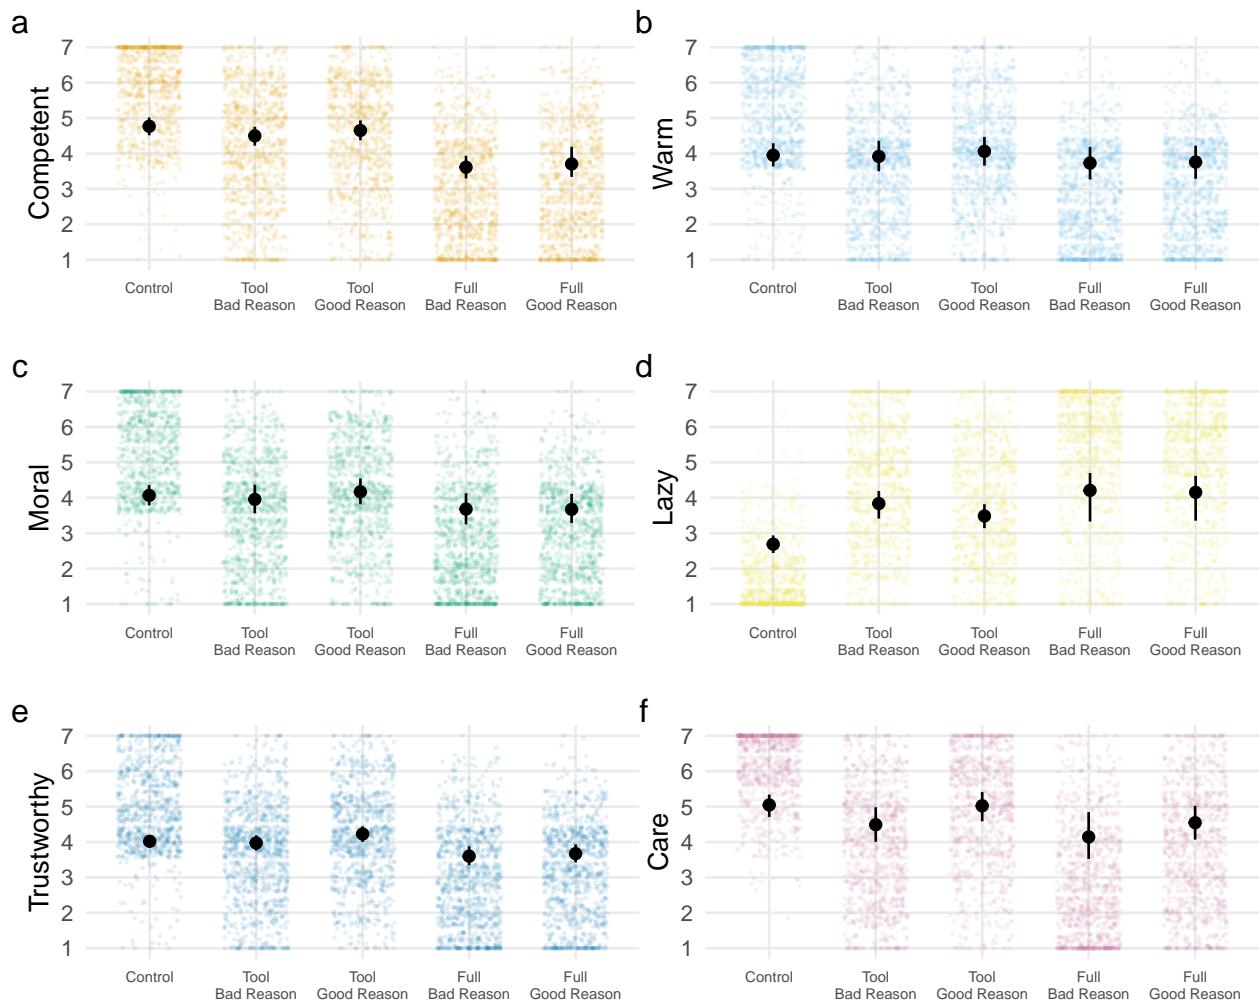
458 (https://osf.io/ac9g3/?view_only=912d9b57023d49baa87eea999574f0ce).

459 ***Statistical Analysis***

460 We fitted the same Bayesian multivariate multilevel cumulative-link ordinal models as in
461 Studies 1 and 3. All models converged normally ($\hat{R} \leq 1.01$).

462 **Results**

463 We first looked across all the tasks. In line with our previous results, we found that people
464 who fully outsourced to AI by copying the output verbatim were perceived as less competent, less
465 moral, and less trustworthy than people who used AI as a collaborative tool (Figure 6; Table 4).
466 Perhaps surprisingly, though, people’s reasons for outsourcing to AI did not appear to influence
467 character evaluations when pooling across all the tasks. When looking at the tasks overall,
468 character evaluations did not differ between people who really cared about the task and wanted to
469 get it right and people who used AI because they were short on time and wanted to complete the
470 task quickly. This was true both when using the AI as a tool or outsourcing in full.

Figure 6*Overall Character Evaluations in Study 4*

Note. Participants in the control condition and four AI outsourcing conditions evaluated people in the scenarios on (a) competence, (b) warmth, (c) morality, (d) laziness, and (e) trustworthiness. Coloured points represent participant responses to the questions, jittered for easier viewing. Black points are estimated marginal means from the fitted model, pooling over participants and tasks. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.

Table 4*Overall Pairwise Contrasts in Study 4*

	Response					
	Competent	Warm	Moral	Lazy	Trustworthy	Care
Comparison to control						
Tool Bad Reason - Control	-0.27 [-0.53 -0.02]	-0.03 [-0.37 0.30]	-0.11 [-0.45 0.24]	1.16 [0.73 1.51]	-0.05 [-0.27 0.16]	-0.56 [-0.97 -0.11]
Tool Good Reason - Control	-0.12 [-0.38 0.13]	0.11 [-0.21 0.43]	0.11 [-0.19 0.39]	0.80 [0.43 1.17]	0.21 [-0.02 0.42]	-0.02 [-0.37 0.31]
Full Bad Reason - Control	-1.16 [-1.45 -0.83]	-0.22 [-0.59 0.14]	-0.39 [-0.75 0.00]	1.58 [0.75 2.01]	-0.42 [-0.68 -0.15]	-0.92 [-1.46 -0.29]
Full Good Reason - Control	-1.08 [-1.40 -0.62]	-0.19 [-0.54 0.16]	-0.40 [-0.73 -0.05]	1.50 [0.78 1.94]	-0.35 [-0.59 -0.09]	-0.50 [-0.91 -0.07]
Effect of outsourcing type						
Full Bad Reason - Tool Bad Reason	-0.89 [-1.19 -0.55]	-0.19 [-0.66 0.27]	-0.27 [-0.77 0.21]	0.42 [-0.53 0.98]	-0.37 [-0.68 -0.05]	-0.36 [-1.04 0.39]
Full Good Reason - Tool Good Reason	-0.96 [-1.32 -0.46]	-0.30 [-0.78 0.16]	-0.50 [-0.93 -0.04]	0.70 [-0.11 1.24]	-0.56 [-0.85 -0.23]	-0.48 [-0.99 0.07]
Effect of reasons for outsourcing						
Tool Bad Reason - Tool Good Reason	-0.15 [-0.44 0.14]	-0.14 [-0.59 0.28]	-0.22 [-0.63 0.23]	0.36 [-0.16 0.82]	-0.25 [-0.53 0.03]	-0.54 [-1.04 0.01]
Full Bad Reason - Full Good Reason	-0.07 [-0.60 0.31]	-0.03 [-0.54 0.45]	0.01 [-0.48 0.49]	0.08 [-0.85 0.85]	-0.07 [-0.41 0.26]	-0.42 [-1.10 0.34]
Interaction effect						
Interaction effect	0.07 [-0.50 0.55]	0.11 [-0.57 0.75]	0.22 [-0.41 0.86]	-0.28 [-1.28 0.64]	0.18 [-0.26 0.61]	0.13 [-0.73 1.01]

Note. Numbers reflect differences in marginal means on a 7-point Likert scale, pooling over participants and tasks. The bottom row represents the interaction between outsourcing type and the reasons for outsourcing (i.e., the difference between the differences in the rows above). Main numbers are posterior medians, numbers in the square brackets are 95% credible intervals.

471 Importantly, though, as in our previous studies, the type of task mattered (Figure 7).

472 Perceptions of outsourcing were particularly negative for tasks that are social, require social skills,

473 impact others, have important consequences, and require effort (Supplementary Figures 6 and 7).

474 Indeed, for socio-relational tasks like writing an apology letter and writing wedding vows, people

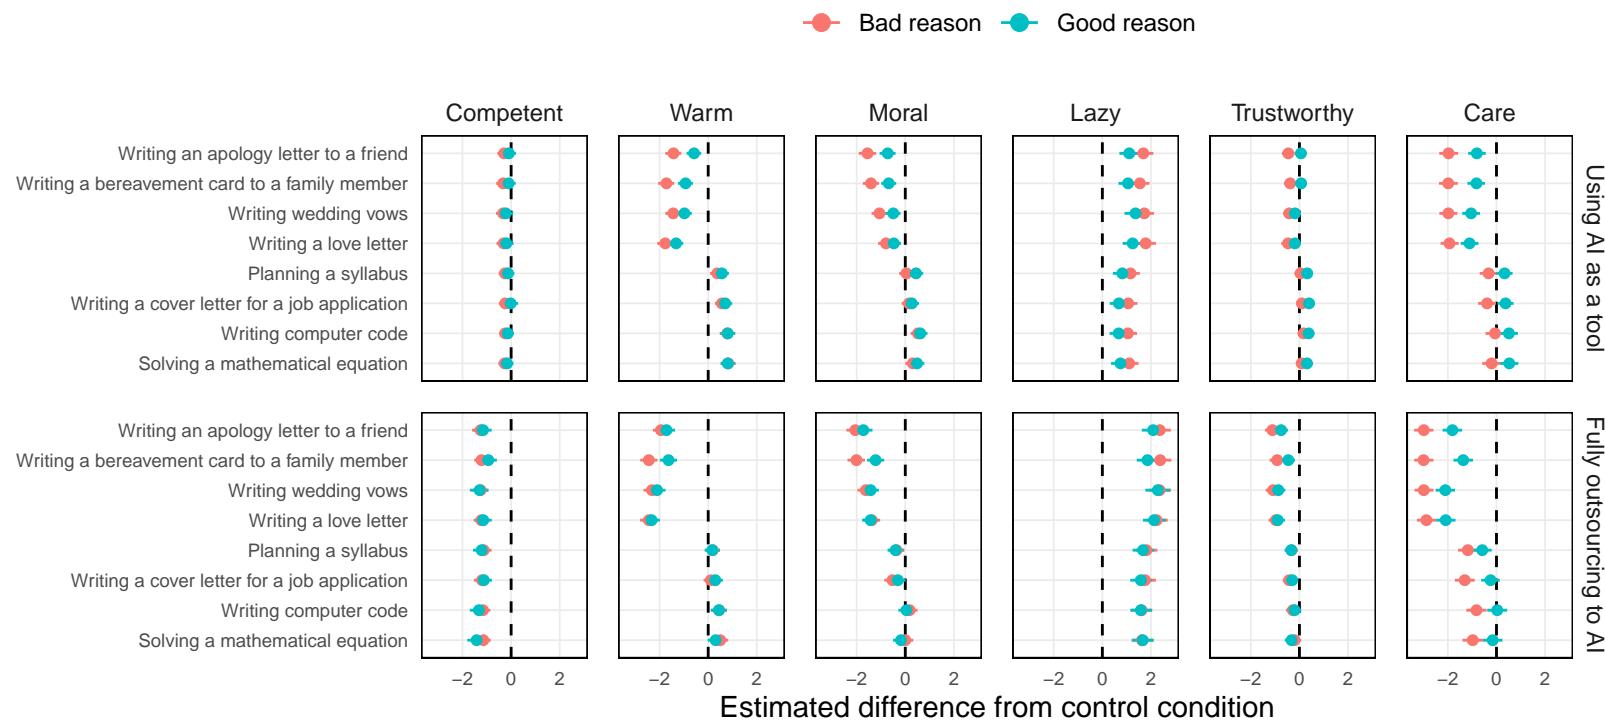
475 using AI as a tool for good reasons were still perceived more negatively than the control condition

476 on the dimensions of warmth, morality, laziness, and care, though not on the dimensions of

477 competence or trustworthiness.

Figure 7

Variation in the Effects of Outsourcing across Tasks in Study 4



Note. Tasks are ordered from most social (top) to least social (bottom) according to ratings from a pilot study. Point ranges are differences in marginal means on a 7-point Likert scale for the bad reason conditions (red) and good reason conditions (blue) compared to the control condition. Upper panels refer to the tool outsourcing conditions, and lower panels refer to the full outsourcing conditions. Points and ranges represent posterior medians and 95% credible intervals, respectively.

478 Moreover, when we delved further into the task-specific estimates, we found that the
479 reasons manipulation did indeed have an effect on character evaluations for social tasks – but not
480 non-social tasks (Supplementary Figure 8). When writing a bereavement card, for example,
481 people were perceived as less warm, less moral, lazier, and less trustworthy when they used AI to
482 save time compared to when they used it because they cared about doing the task well. The same
483 was not true for non-social tasks like writing computer code or solving a mathematical equation.

484 Discussion

485 In Study 4, we attempted to counteract the potential perception that outsourcing to AI
486 reflects caring less about the task by explicitly informing participants about the person's reason for
487 outsourcing: they outsourced to AI because they really cared about the task and wanted to get it
488 right. As well as replicating our finding that fully outsourcing to AI is perceived more negatively
489 than using AI as a tool, we also found an important effect of the reasons for outsourcing, but only
490 for socio-relational tasks. When writing a bereavement card or an apology letter, for example,
491 people were perceived more negatively if they used an AI tool to produce a quick output in a rush,
492 rather than to ensure they got it right. Nonetheless, for socio-relational tasks, the “best” use of AI
493 in this study – using AI as a tool because they cared about the task and wanted to get it right —
494 still led to targets being perceived more negatively than if they had completed the task themselves.

495 While we have so far shown varying evidence for three different mechanisms that might
496 underlie the negative perceptions of outsourcing to AI – effort, authenticity, and caring about the
497 task – it is likely that these mechanisms are related. For example, outsourcing to AI might indicate
498 a lack of effort, which then might signal a lack of authenticity and reduced care in the task,
499 leading to negative character evaluations. Our previous studies have been unable to test causal
500 models like these as we manipulated the mechanisms separately and independently. In Study 5,
501 therefore, we bring all three mechanisms together and test their combined associations with
502 character evaluations. To do this, we focus on a single socio-relational task — writing a love letter
503 — which we elaborate for participants with a more detailed vignette.

504

Study 5

505 **Methods**506 **Participants**

507 We conducted a power simulation to determine our target sample size. The simulation
508 suggested that a sample size of 200 participants per condition (overall $n = 600$ for three
509 conditions) would be required to detect a small-to-medium difference between conditions
510 (Cohen's $d \approx 0.30$) with above 80% power.

511 We recruited a convenience sample of 651 participants from the United Kingdom through
512 Prolific. After excluding participants who failed our pre-treatment attention check, we were left
513 with a final sample of 610 participants (371 female; 233 male; 4 non-binary / third gender; 2
514 undisclosed gender; mean age = 42.85 years). 82% of these participants reported having used AI
515 tools like ChatGPT before.

516 **Design**

517 We randomly allocated participants into one of three conditions in a between-subjects
518 design: (i) the control condition, (ii) the tool outsourcing condition, or (iii) the full outsourcing
519 condition. These conditions determined how the scenario was presented to participants.

520 **Procedure**

521 We presented participants with a vignette about a person, Adam, who is writing a love
522 letter in a Valentine's Day card to his partner (see Supplementary Materials for full vignette
523 wording). We told participants in each of the conditions:

- 524 • *Control condition*: “Adam decides to write the love letter in the card by himself.”
- 525 • *Tool outsourcing condition*: “Adam decides to use AI to help write the love letter in the
526 card. He asks ChatGPT to provide ideas, inspiration, and feedback, but he edits and
527 rewrites the suggestions and finishes writing the love letter himself.”
- 528 • *Full outsourcing condition*: “Adam decides to use AI to write the love letter in the card. He
529 asks ChatGPT to write the love letter and copies the output word-for-word, rather than

530 writing it himself.”

531 We then presented participants with the love letter that Adam wrote (in reality, this was
532 written by ChatGPT version 4o; see Supplementary Materials for wording). On the following
533 page, we asked participants what Adam wrote and whether he used AI to help. 95% of
534 participants answered both of these comprehension questions correctly.

535 Using 7-point Likert scales, we then asked participants how much effort they thought
536 Adam put into the love letter, how authentic they thought the love letter was, how much they
537 thought Adam cared about the love letter, and the same five character evaluations as in our
538 previous studies. In additional free response questions, we asked participants to explain how they
539 felt towards Adam and how they would feel if they were Adam’s partner. Finally, we asked
540 participants several questions about AI tools like ChatGPT.

541 ***Pre-registration***

542 We pre-registered the study on the Open Science Framework
543 (https://osf.io/ac9g3/?view_only=912d9b57023d49baa87eea999574f0ce).

544 ***Statistical Analysis***

545 We fitted two Bayesian regression models to the data. The first model was a multivariate
546 cumulative-link ordinal model including all Likert scales as separate response variables. The
547 second model was a path model capturing the effect of outsourcing on character evaluations, both
548 directly and indirectly through perceptions of effort, authenticity, and care. In this second model,
549 we included ordinal predictors as monotonic effects and modelled the five character evaluations as
550 a single latent variable. We used regularising priors for all parameters to impose conservatism on
551 parameter estimates. All models converged normally ($\hat{R} \leq 1.01$).

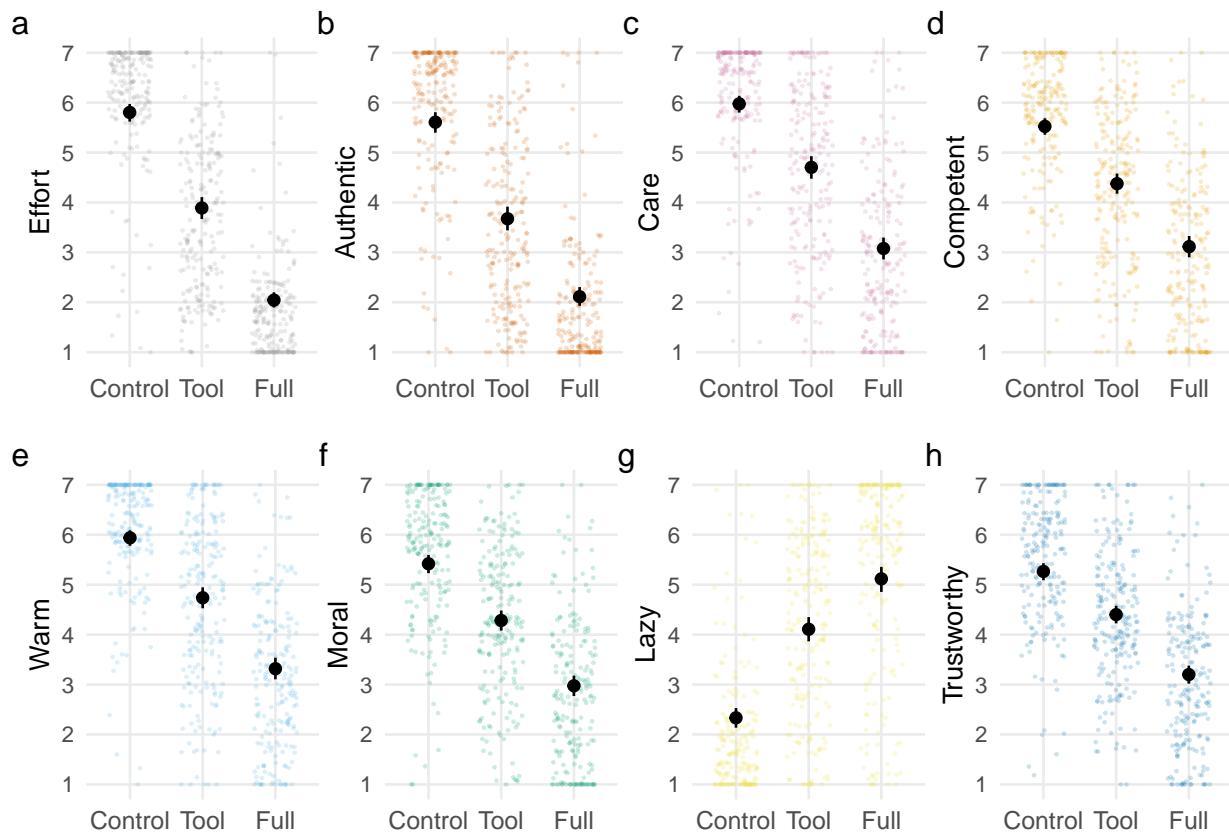
552 **Results**

553 Across all measures, we found that outsourcing the love letter to AI was perceived more
554 negatively compared to the control condition and that fully outsourcing to AI was perceived more
555 negatively than using AI as a collaborative tool (Figure 8; Table 5). Not only did outsourcing the

- 556 love letter lead to more negative character evaluations, but outsourcing to AI was also seen as less
 557 effortful, less authentic, and indicative of caring less about the task.

Figure 8

Perceptions of the Person and the Love Letter in Study 5



Note. Participants in the control, tool outsourcing, and full outsourcing conditions rated (a) the amount of effort put into the love letter, (b) how authentic the love letter was, (c) how much the person cared about the love letter, and (d-h) five character evaluation measures. Coloured points represent participant responses to the questions, jittered for easier viewing. Black points are estimated marginal means from the fitted model. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.

Table 5*Pairwise Contrasts in Study 5*

	Response							
	Effort	Authentic	Care	Competent	Warm	Moral	Lazy	Trustworthy
Tool - Control	-1.91 [-2.18 -1.64]	-1.94 [-2.23 -1.61]	-1.27 [-1.54 -1.00]	-1.15 [-1.40 -0.89]	-1.20 [-1.45 -0.95]	-1.14 [-1.39 -0.88]	1.77 [1.47 2.07]	-0.86 [-1.11 -0.62]
Full - Control	-3.76 [-3.98 -3.52]	-3.50 [-3.77 -3.22]	-2.90 [-3.16 -2.63]	-2.41 [-2.67 -2.14]	-2.62 [-2.88 -2.37]	-2.44 [-2.71 -2.17]	2.78 [2.47 3.09]	-2.07 [-2.30 -1.82]
Full - Tool	-1.85 [-2.11 -1.58]	-1.56 [-1.87 -1.26]	-1.63 [-1.93 -1.31]	-1.26 [-1.55 -0.97]	-1.42 [-1.72 -1.12]	-1.31 [-1.60 -1.02]	1.01 [0.65 1.35]	-1.20 [-1.45 -0.95]

Note. Numbers reflect differences in marginal means on a 7-point Likert scale. Main numbers are posterior medians, numbers in the square brackets are 95% credible intervals.

558 Exploratory text analysis of participants' free responses supported this quantitative pattern

559 (see Supplementary Materials for methodology and Supplementary Table 7 for results). When

560 comparing word frequencies between conditions, we found that Adam was more likely to be

561 described as "lazy" and less likely to be described as "caring", "thoughtful", and "genuine" in

562 both outsourcing conditions compared to the control condition. Adam was also more likely to be

563 described as "romantic" and as someone who "loves" his partner when he used AI as a

564 collaborative tool, compared to when he fully outsourced the love letter to AI.

565 When we included all the variables in a single path model, we found that outsourcing

566 influenced character evaluations both directly and indirectly through our proposed mechanisms

567 (Figure 9). The indirect effects showed that people perceived outsourced work as less effortful,

568 and less effortful work was seen as less authentic and indicating less care about the task. In turn,

569 less authenticity and care were associated with more negative evaluations of the person. Effort

570 itself was not directly related to character evaluations, suggesting that effort works solely through

571 perceptions of authenticity and care.

572 General Discussion

573 The release of openly available generative AI LLMs has changed lives, promising to let

574 people do more tasks, more efficiently, and perhaps to do so better than they could alone. People

575 can — and *do* — use AI tools like ChatGPT to, for example, create dinner recipes, assist with

576 coding, and even write job applications ([Department for Science, Innovation & Technology, 2024](#)).

577 But it is not only such routine, everyday, and non-social tasks that AI now "assists" with.

578 People can use AI for a seemingly endless range of social tasks too, from crafting apology letters

579 to writing condolences to even writing wedding vows. In this paper, across five pre-registered

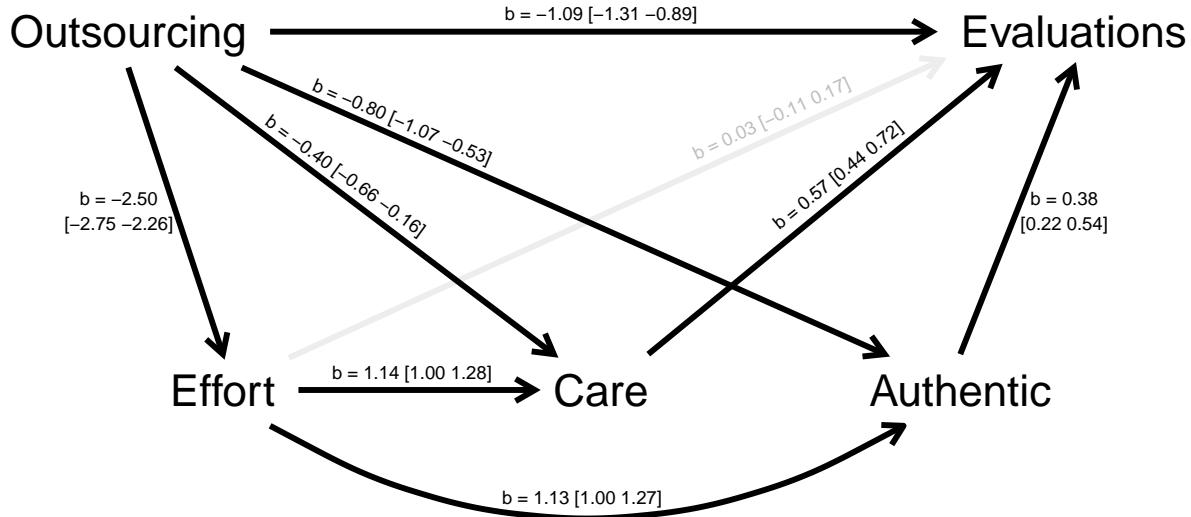
580 experiments, we show how — and why — AI-outsourcing shapes perceptions of others in a world

581 where outsourcing has never been easier and cheaper.

582 In Study 1, we showed that people who outsourced tasks to AI were perceived more

583 negatively than people who completed the tasks by themselves. These negative impressions were

584 particularly strong for people who used AI to complete socio-relational tasks, such as writing a

Figure 9*Path Model in Study 5*

Note. All predictors were modelled as monotonic effects, such that parameters can be interpreted as the expected average difference between two adjacent categories of the ordinal predictor on the logit scale. The “evaluations” outcome variable was modelled as a single latent variable with loadings from all five character evaluations (competence, warmth, morality, laziness, and trustworthiness).

585 love letter or writing wedding vows, and for people who copied the model’s first output verbatim
 586 without acknowledging their reliance on AI. Moreover, these negative perceptions were found
 587 even for the “best case” of openly acknowledging the use of AI as a collaborative tool. In Study 2,
 588 we showed that people perceive both the outsourcer and the outworked work more negatively,
 589 with outsourced work perceived as less meaningful, less authentic, and less reward-worthy than
 590 ostensibly human-generated writing. In Study 3, we showed that while it matters whether people
 591 spent time crafting the AI prompts or simply gave a rushed initial prompt, even expending effort
 592 into crafting the best prompts was still not enough to counteract the negative effects from using

593 AI. In Study 4, we explored the potential role of inferred importance and found that while
594 explicitly telling participants that the person used AI because they cared about the task reduced
595 negative perceptions for social tasks, it was still not enough to eliminate negative perceptions
596 completely. In Study 5, we showed that a perceived lack of effort is taken to signal both a lack of
597 authenticity and lack of importance attached to the task, and these independently influenced
598 character judgments above and beyond the effect of effort.

599 Our findings extend work on the moralisation of effort. Studies have shown that people
600 inherently value effort and perceive displays of effort as costly signals of one's moral character
601 and cooperative intent ([Celniker et al., 2023](#)). And yet it has remained unclear how we might view
602 others who outsource to AI; how these effects might vary based on how socio-relational the task
603 is; how different ways of outsourcing influence perceptions; how outsourcing has different effects
604 on different kinds of social perceptions; and why exactly effort has the effects that it does. Across
605 our studies, we provide new insight into all of these questions. In line with previous work on the
606 importance of effort, we show that people negatively judge those who outsource to AI. We show
607 that the type of task does matter, whereby outsourcing to AI for socio-relational tasks leads to
608 particularly negative perceptions. We show that different ways of outsourcing lead to differences
609 in the degree of negative perceptions but that, critically, even outsourcing to AI in the “best” way
610 (e.g., using it as a tool and finishing the work oneself while being honest about the AI use) is still
611 not enough to eliminate the negative consequences. We show that negative perceptions from
612 outsourcing tended to go together, even if outsourcing on social tasks led to particularly negative
613 effects on warmth and morality traits. And finally, we provide further insight into why effort
614 matters. The reduced effort from outsourcing socio-relational tasks to AI signals that the work is
615 less authentically one's own and that the person cares less about the task (and therefore, perhaps,
616 the relationship). The lack of a direct effect of perceived effort in our path model showed that it is
617 inferences of authenticity and care, rather than perceived effort per se, that are associated with
618 negative character evaluations. As a participant in our final study put it: “*If he really cared, he*
619 *would have just done it by himself from scratch*” (female, 25 years old).

620 Our findings cohere with the philosophical idea that there is value in *how* a task was done,
621 and not merely *whether* it was done (Aristotle, 2009; Goodman, 2010; Hursthouse & Pettigrove,
622 2023; Stohr, 2006). For many socio-relational tasks, it might seem that part of the constitutive
623 action is the *process* by which it occurs: an apology that does not contain a genuine reflection and
624 commitment to do better, rather than just the words “I am sorry”, might not seem to be an apology
625 at all. In contrast, for many of the non-social tasks, it is easier to distinguish the importance of the
626 process from the outcome. In this way, our work suggests that people rarely adopt a purely
627 utilitarian perspective in which outcomes are the sole determinant (Everett & Kahane, 2020;
628 Kahane et al., 2018). Instead, their judgments cohere more with ideas from virtue ethics about the
629 importance of *doing* (Hursthouse & Pettigrove, 2023; Stohr, 2006). Outsourcing to AI –
630 especially for social tasks — may allow us to produce similar outputs, but by severing the
631 outcome from the practice of doing, it may risk the development and maintenance of our human
632 virtues (Vallor, 2015, 2024).

633 AI is often being marketed as being able to help us to do more and more tasks, promising
634 gains of efficiency that align with societal incentives for “hacks” that encourage people to do
635 more-and-more with less energy and effort. Our work, however, highlights that when it comes to
636 our psychology, efficiency is not the only currency. Instead, *inefficiency* can sometimes pay off
637 more, especially for social tasks. By expending effort themselves instead of outsourcing to AI,
638 people are able to signal authenticity and care for the task, and this can lead to better reputations
639 (see also Celniker et al., 2023). Correspondingly, expending effort, even “unnecessarily”, is not as
640 irrational, biased, or suboptimal as we might think from a utilitarian perspective in which
641 outcomes are the only things that matter. Instead, it is precisely this inefficiency that helps people
642 signal things that they care about and connect with others, thereby arguably reflecting a deeply
643 rational reflection of virtues and the importance of social ties (Everett et al., 2016).

644 Most speculatively, our results on the negative effects of AI-outsourcing on character
645 judgments highlight potential risks in how increased use of AI could lead to negative
646 consequences for social ties, especially if people start to assume, by default, that others are using

647 AI for the kind of tasks that matter. Sociologists have highlighted concerns about the negative
648 effects that outsourcing to AI can have on our “connective labour”, arguing that while AI can
649 enhance certain tasks, it cannot replicate the depth of human relationships essential for effective
650 caregiving, education, and support (Pugh, 2024). Similar arguments have been made about the
651 risks of outsourcing empathy to AI (Landes & Everett, 2025). In this way, the rapid move towards
652 using AI for more and more tasks could have serious and unintended consequences on the way we
653 connect with one another, serving to further weaken the social ties that bind us into a community.

654 **Limitations and Directions for Future Research**

655 The studies in this paper are not without their limitations. While we included a range of
656 different socio-relational and professional tasks in an effort to improve the generalisability of our
657 findings across domains, it would be interesting for future work to additionally explore the
658 generalisability and variability of our findings across countries with different AI infrastructures
659 and readiness levels (Oxford Insights, 2024; Tortoise Media, 2024) and over time as AI use
660 becomes more commonplace. By focusing on generalisability across various real-world tasks in
661 which people outsource, it could also be argued that our design lacks the richness of information
662 in extended vignettes that might influence character evaluations. While we have advanced
663 previous research in highlighting the ways in which effort influences perceptions of authenticity
664 and care, it will be interesting for future research to delve deeper into these mechanisms, both
665 philosophically and psychologically: *why* is it that the perceived care for the task matters, and
666 what are the boundary conditions of these effects? Finally, while we have demonstrated negative
667 perceptions of outsourcing in this paper, it will be important for future research to explore when
668 people might deem outsourcing to AI as acceptable or even preferable. Several of the participants
669 in our final study expressed in their free responses that they would have been okay with Adam
670 using AI to write the love letter if he was not a confident writer or had a learning difficulty that
671 made writing challenging, such as dyslexia. In line with this, some research has found that people
672 are more accepting of cognition-enhancing technologies and drugs when they are used to repair
673 cognitive functions, rather than to enhance cognitive functions beyond “normal” levels (Medaglia

674 et al., 2019; Rudski, 2014). Future research should explore whether negative perceptions of
675 outsourcing persist when AI is used in a reparative way.

676 **Conclusions**

677 To conclude, across five pre-registered studies, we have demonstrated negative perceptions
678 of outsourcing to AI. Our participants perceived individuals who outsource tasks to AI more
679 negatively across a range of character dimensions and perceived outsourced work as less
680 meaningful and authentic. Negative perceptions were particularly strong for socio-relational tasks,
681 such as writing wedding vows, and were compounded when the outsourcer copied the AI's output
682 verbatim and did not honestly acknowledge their use of AI. These findings connect with broader
683 debates about the importance of *doing* in social relationships, and highlight that for many tasks –
684 especially those that are more socio-relational – it might be better to move away from a focus on
685 making things more efficient at all costs and instead bring back a recognition of the power of
686 inefficiency. Doing something oneself, even if AI could do it quicker and easier, signals one that is
687 authentic and cares about the task and therefore can help bind us together. In a world of
688 algorithm-mediated interactions, AI is no substitute for investing effort into our interpersonal
689 relationships.

Acknowledgements

This work was generously supported by funding from REDACTED.

Data and Code Availability

All data and original code can be found here:

https://osf.io/ac9g3/?view_only=912d9b57023d49baa87eea999574f0ce.

Statement of Interests

The authors have no conflicts of interest to disclose.

References

- Abele, A. E., Ellemers, N., Fiske, S. T., Koch, A., & Yzerbyt, V. (2021). Navigating the social world: Toward an integrated framework for evaluating self, individuals, and groups. *Psychological Review*, 128(2), 290–314. <https://doi.org/10.1037/rev0000262>
- Allaire, J. J., Teague, C., Scheidegger, C., Xie, Y., Dervieux, C., & Woodhull, G. (2024). *Quarto* (Version 1.7.23) [Computer software]. <https://doi.org/10.5281/zenodo.5960048>
- Anderson, L. (2025). *How are students using ChatGPT? For therapy, breakups, and even texting friends*. Teen Vogue.
<https://www.teenvogue.com/story/how-students-using-chatgpt-therapy-breakups>
- Aristotle. (2009). *The Nicomachean ethics* (D. Ross & L. Brown, Translators). Oxford University Press.
- Bürkner, P.-C. (2017). brms: An R package for Bayesian multilevel models using Stan. *Journal of Statistical Software*, 80(1), 1–28. <https://doi.org/10.18637/jss.v080.i01>
- Celniker, J. B., Gregory, A., Koo, H. J., Piff, P. K., Ditto, P. H., & Shariff, A. F. (2023). The moralization of effort. *Journal of Experimental Psychology: General*, 152(1), 60–79.
<https://doi.org/10.1037/xge0001259>
- Cubitt, R. P., Drouvelis, M., Gächter, S., & Kabalin, R. (2011). Moral judgments in social dilemmas: How bad is free riding? *Journal of Public Economics*, 95(3), 253–264.
<https://doi.org/10.1016/j.jpubeco.2010.10.011>
- Department for Science, Innovation & Technology. (2024). *Public attitudes to data and AI: Tracker survey (Wave 4) report*. <https://www.gov.uk/government/publications/public-attitudes-to-data-and-ai-tracker-survey-wave-4/public-attitudes-to-data-and-ai-tracker-survey-wave-4-report>
- Earp, B. D., Porsdam Mann, S., Liu, P., Hannikainen, I., Khan, M. A., Chu, Y., & Savulescu, J. (2024). Credit and blame for AI-generated content: Effects of personalization in four countries. *Annals of the New York Academy of Sciences*, 1542(1), 51–57.
<https://doi.org/10.1111/nyas.15258>

- Everett, J. A., & Kahane, G. (2020). Switching tracks? Towards a multidimensional model of utilitarian psychology. *Trends in Cognitive Sciences*, 24(2), 124–134.
<https://doi.org/10.1016/j.tics.2019.11.012>
- Everett, J. A., Pizarro, D. A., & Crockett, M. J. (2016). Inference of trustworthiness from intuitive moral judgments. *Journal of Experimental Psychology: General*, 145(6), 772–787.
<https://doi.org/10.1037/xge0000165>
- Fiske, A. P. (1992). The four elementary forms of sociality: Framework for a unified theory of social relations. *Psychological Review*, 99(4), 689–723.
<https://doi.org/10.1037/0033-295X.99.4.689>
- Fiske, S. T., Cuddy, A. J., & Glick, P. (2007). Universal dimensions of social cognition: Warmth and competence. *Trends in Cognitive Sciences*, 11(2), 77–83.
<https://doi.org/10.1016/j.tics.2006.11.005>
- Goodman, R. (2010). Cognitive enhancement, cheating, and accomplishment. *Kennedy Institute of Ethics Journal*, 20(2), 145–160. <https://doi.org/10.1353/ken.0.0309>
- Goodwin, G. P., Piazza, J., & Rozin, P. (2014). Moral character predominates in person perception and evaluation. *Journal of Personality and Social Psychology*, 106(1), 148–168.
<https://doi.org/10.1037/a0034726>
- Heider, F. (1958). *The psychology of interpersonal relations*. John Wiley & Sons Inc.
<https://doi.org/10.1037/10628-000>
- Hursthouse, R., & Pettigrove, G. (2023). Virtue ethics. In E. N. Zalta & U. Nodelman (Eds.), *The Stanford encyclopedia of philosophy*.
<https://plato.stanford.edu/archives/fall2023/entries/ethics-virtue/>
- Kahane, G., Everett, J. A., Earp, B. D., Caviola, L., Faber, N. S., Crockett, M. J., & Savulescu, J. (2018). Beyond sacrificial harm: A two-dimensional model of utilitarian psychology. *Psychological Review*, 125(2), 131–164. <https://doi.org/10.1037/rev0000093>
- Kerr, N. L. (1983). Motivation losses in small groups: A social dilemma analysis. *Journal of Personality and Social Psychology*, 45(4), 819–828.

- <https://doi.org/10.1037/0022-3514.45.4.819>
- Kruger, J., Wirtz, D., Van Boven, L., & Altermatt, T. W. (2004). The effort heuristic. *Journal of Experimental Social Psychology*, 40(1), 91–98.
- [https://doi.org/10.1016/S0022-1031\(03\)00065-9](https://doi.org/10.1016/S0022-1031(03)00065-9)
- Landau, W. M. (2021). The targets R package: A dynamic Make-like function-oriented pipeline toolkit for reproducibility and high-performance computing. *Journal of Open Source Software*, 6(57), 2959. <https://doi.org/10.21105/joss.02959>
- Landes, E., & Everett, J. A. C. (2025). *AI should develop human empathy, not replace it*. PsyArXiv. https://doi.org/10.31234/osf.io/y3qzu_v1
- Liu, B., Kang, J., & Wei, L. (2024). Artificial intelligence and perceived effort in relationship maintenance: Effects on relationship satisfaction and uncertainty. *Journal of Social and Personal Relationships*, 41(5), 1232–1252. <https://doi.org/10.1177/02654075231189899>
- Malle, B. F. (2022). Attribution theories. In *Theories in social psychology, second edition* (pp. 93–120). John Wiley & Sons, Ltd. <https://doi.org/10.1002/9781394266616.ch4>
- Medaglia, J. D., Yaden, D. B., Helion, C., & Haslam, M. (2019). Moral attitudes and willingness to enhance and repair cognition with brain stimulation. *Brain Stimulation*, 12(1), 44–53. <https://doi.org/10.1016/j.brs.2018.09.014>
- miramar0. (2024). *My (30F) husband (36M) used AI to write his vows. How do I approach him about this?* [Online forum post]. Reddit. https://www.reddit.com/r/relationship_advice/comments/1d9tw9j/my_30f_husband_36m_used_ai_to_write_his_vows_how/
- Newman, G. E. (2019). The psychology of authenticity. *Review of General Psychology*, 23(1), 8–18. <https://doi.org/10.1037/gpr0000158>
- Noy, S., & Zhang, W. (2023). Experimental evidence on the productivity effects of generative artificial intelligence. *Science*, 381(6654), 187–192. <https://doi.org/10.1126/science.adh2586>
- Oxford Insights. (2024). *Government AI readiness index*. <https://oxfordinsights.com/ai-readiness/ai-readiness-index/>
- Pedersen, T. L. (2025). *patchwork: The composer of plots*. <https://patchwork.data-imaginist.com>

- Pugh, A. J. (2024). *The last human job: The work of connecting in a disconnected world.* Princeton University Press.
- R Core Team. (2022). *R: A language and environment for statistical computing.* R Foundation for Statistical Computing. <https://www.R-project.org/>
- Roth, L. H. O., & Tissot, T. (2025). *A costly co-pilot? AI usage signals reduced morality via effort perceptions.* OSF Preprints. https://doi.org/10.31219/osf.io/vwxgz_v1
- Rudski, J. M. (2014). A comparison of attitudes toward pharmacological treatment versus enhancement under competitive and noncompetitive conditions. *AJOB Empirical Bioethics*, 5(2), 80–90. <https://doi.org/10.1080/23294515.2013.850122>
- Stohr, K. (2006). Contemporary virtue ethics. *Philosophy Compass*, 1(1), 22–27. <https://doi.org/10.1111/j.1747-9991.2006.00004.x>
- Tait, A. (2024). *My boyfriend sent me a heartfelt message - then I learned he'd used ChatGPT.* The i Paper. <https://inews.co.uk/inews-lifestyle/boyfriend-sent-heartfelt-message-learned-chatgpt-3413336>
- Tortoise Media. (2024). *The global AI index.* <https://www.tortoisemedia.com/data/global-ai>
- Vallor, S. (2015). Moral deskilling and upskilling in a new machine age: Reflections on the ambiguous future of character. *Philosophy & Technology*, 28, 107–124. <https://doi.org/10.1007/s13347-014-0156-9>
- Vallor, S. (2024). *The AI mirror: How to reclaim our humanity in an age of machine thinking.* Oxford University Press. <https://doi.org/10.1093/oso/9780197759066.001.0001>
- Wickham, H. (2016). *ggplot2: Elegant graphics for data analysis.* Springer-Verlag New York. <https://ggplot2.tidyverse.org>

Supplementary Materials

Negative Perceptions of Outsourcing to Artificial Intelligence

AUTHORS REDACTED

AFFILIATION REDACTED

Table of contents

Significance Statement	2
Pilot Study 1	2
Pilot Study 2	4
Vignette Wording in Study 5	7
Methods for Text Analysis in Study 5	8
Supplementary Figures	10
Supplementary Tables	29
Supplementary References	37

Pilot Study 1

Methods

Participants

We recruited a convenience sample of 200 participants from the United Kingdom through Prolific. After excluding participants who failed our pre-treatment attention check, we were left with a final sample of 186 participants (118 female; 67 male; 1 non-binary / third gender; 0 undisclosed gender; mean age = 38.99 years).

Procedure

We presented participants with six different tasks “that people might perform in their daily lives”. The six tasks were randomly drawn from a larger set of 20 tasks (see Supplementary Table 1 for the full list of tasks). For each task, we asked participants the following questions on 7-point Likert scales:

- Is this a social task?
- Does this task require social skills?
- Does this task impact other people?
- How important are the consequences of this task?
- How important is it that effort goes into this task?
- How important is it that others see the effort that goes into this task?

Statistical Analysis

We fitted a Bayesian multivariate multilevel cumulative-link ordinal model to the data using the *brms* R package. We modelled each task evaluation as a separate response variable and included correlated varying intercepts for participants and tasks. We used regularising priors for all parameters to impose conservatism on parameter estimates (see Supplementary Materials for full model specification). The model converged normally ($\hat{R} \leq 1.01$).

Results

We found that participants' responses to all six questions tended to be positively correlated. For example, tasks rated as more social were also rated as requiring more social skills (see Supplementary Figure 9). Estimated averages and rankings for the 20 tasks across each of the questions can be found in Supplementary Figures 10 – 15.

Pilot Study 2

Methods

Participants

We conducted a power simulation to determine our target sample size. The simulation suggested that a sample size of 150 participants per condition (overall $n = 450$ for three conditions) would be required to detect a small difference between conditions (Cohen's $d \approx 0.20$) with above 80% power.

We recruited a convenience sample of 500 participants from the United Kingdom through Prolific. After excluding participants who failed our pre-treatment attention check, we were left with a final sample of 466 participants (292 female; 169 male; 4 non-binary / third gender; 1 undisclosed gender; mean age = 42.32 years). 73% of these participants reported having used ChatGPT before.

Design

We randomly allocated participants into one of three conditions in a between-subjects design: (i) the control condition, (ii) the AI outsourcing condition, or (iii) the human outsourcing condition. These conditions determined how scenarios were presented to participants.

Procedure

We presented participants with six scenarios. Each scenario described a person completing a task, such as writing computer code or writing a love letter. The six tasks were randomly drawn from a larger set of 20 tasks (see Supplementary Table 1 for the full list of tasks). For each scenario, we told participants:

- *Control condition*: “In order to complete this task, [the person] works on it by themselves from start to finish.”
- *AI outsourcing condition*: “In order to complete this task, [the person] gets the AI tool ChatGPT to do it for them.”

- *Human outsourcing condition:* “In order to complete this task, [the person] gets someone else to do it for them.”

We then asked participants how well each of the following words described the person in the scenario: competent, warm, moral, lazy, and trustworthy. Participants answered these questions on 7-point Likert scales, ranging from “does not describe [the person] well” to “describes [the person] extremely well”.

After the six scenarios, we asked participants several questions about the AI tool ChatGPT, including their familiarity with ChatGPT, whether they had used ChatGPT before, how frequently they used ChatGPT, and how trustworthy they thought ChatGPT was.

Pre-registration

We pre-registered the study on the Open Science Framework (https://osf.io/xhmzk/?view_only=a4da193574d7410ba4d2aa3945a28b05).

Statistical Analysis

We fitted Bayesian multivariate multilevel cumulative-link ordinal models to the data using the *brms* R package. We modelled each character evaluation – competence, warmth, morality, laziness, and trustworthiness – as a separate response variable and included fixed effects for conditions, varying intercepts for participants, and varying intercepts and slopes for tasks. We used regularising priors for all parameters to impose conservatism on parameter estimates (see Supplementary Materials for full model specifications). All models converged normally ($\hat{R} \leq 1.01$).

Results

We found that people who outsourced tasks to AI or other humans were perceived more negatively than people who completed the tasks themselves (Supplementary Figure 16). In particular, people who outsourced were perceived as lazier and less competent, with smaller yet detectable differences for perceptions of warmth, morality, and trustworthiness (Supplementary Table 8). Across all measures, outsourcing to other humans was perceived more negatively than

outsourcing to AI.

We found that the effects of outsourcing varied across the different tasks, especially for perceptions of warmth and morality (Supplementary Figure 17). For example, people were perceived as less warm if they outsourced writing a love letter, but not if they outsourced writing computer code. Similarly, people were perceived as less moral if they outsourced writing an apology letter to a friend, but not if they outsourced writing a dinner recipe. By contrast, the effects of outsourcing on competence, laziness, and trustworthiness were more consistent across tasks.

To determine the factors that predict variation across tasks, we incorporated ratings of tasks from the first pilot study. Participants were asked to rate the 20 tasks on several features: whether the task is social, requires social skills, impacts others, has important consequences, and requires effort. All of these features predicted stronger causal effects of outsourcing compared to control (Supplementary Figures 18 and 19). In other words, outsourcing to AI or other humans is perceived more negatively for tasks that have these features, compared to tasks without these features.

Vignette Wording in Study 5

We presented participants in Study 5 with the following vignette text:

Adam has been dating his partner for almost a year, and Valentine's Day is coming up. He knows that many people exchange a card on Valentine's Day containing a love letter to their partner, and he decides to send a love letter to his partner too.

This year, Adam has been closely following developments in technology and has read of people using AI tools like ChatGPT for things like this, either using it to help with writing or getting AI to do the task completely.

This was followed by the manipulation text (see main text). Participants were then presented with the love letter that Adam ostensibly wrote, which was held constant across conditions:

Happy Valentine's Day, my love.

I don't think I tell you enough just how much you mean to me. Being with you feels like breathing a little easier, like the world is a bit softer just because you're in it. You make the everyday feel special, and somehow you always know how to calm my nerves or make me laugh at just the right moment. I feel like myself with you – maybe even a better version of myself – and that's such a rare and beautiful thing.

I'm so grateful for you – for the way you listen, the way you love, the way you show up, even in the small ways. I hope you know that no matter what, I'm always in your corner. I can't wait to keep making memories together, whether we're off on some adventure or just curled up on the couch. I love you more than I can really put into words, but I promise I'll spend every day trying.

Yours,

Adam

Methods for Text Analysis in Study 5

To generate frequency lists for each experimental condition in Study 5, we created three documents containing the raw text submissions to the open-ended question “In your own words, describe how you feel about Adam and why”. Each raw text submission was paired with a numbered text ID column. The number of submissions was roughly equivalent across conditions: the control condition ($N = 196$), the tool outsourcing condition ($N = 215$), and the full outsourcing condition ($N = 202$).

All text processing was conducted using the Basic Unit-Transposable Text Experimentation Resource (BUTTER; Version 0.9.4.1; Boyd, 2019). To prepare the data, each CSV file was converted into a folder containing individual text files – one per submission – using two plugins: *Read Text from CSV* (Version 1.0.2) and *Save .txt Files to Folder* (Version 1.0.6). The settings for *Read Text from CSV* were as follows: file encoding = UTF-8, row identifier = ID, text column = Text, CSV delimiter = , and CSV quote = “.

To generate frequency lists, we first loaded the .txt files using the *Load .txt Files from Folder* plugin (Version 1.0.4). Tokenization was performed using the Twitter-Aware Tokenizer (Version 1.0.2), with the options *convert text to lowercase* and *reduce elongation* enabled to minimize superficial variation in tokens. We removed filler and function words using the *Remove Stop Words* plugin (Version 1.0.31), applying the default English stop word list.

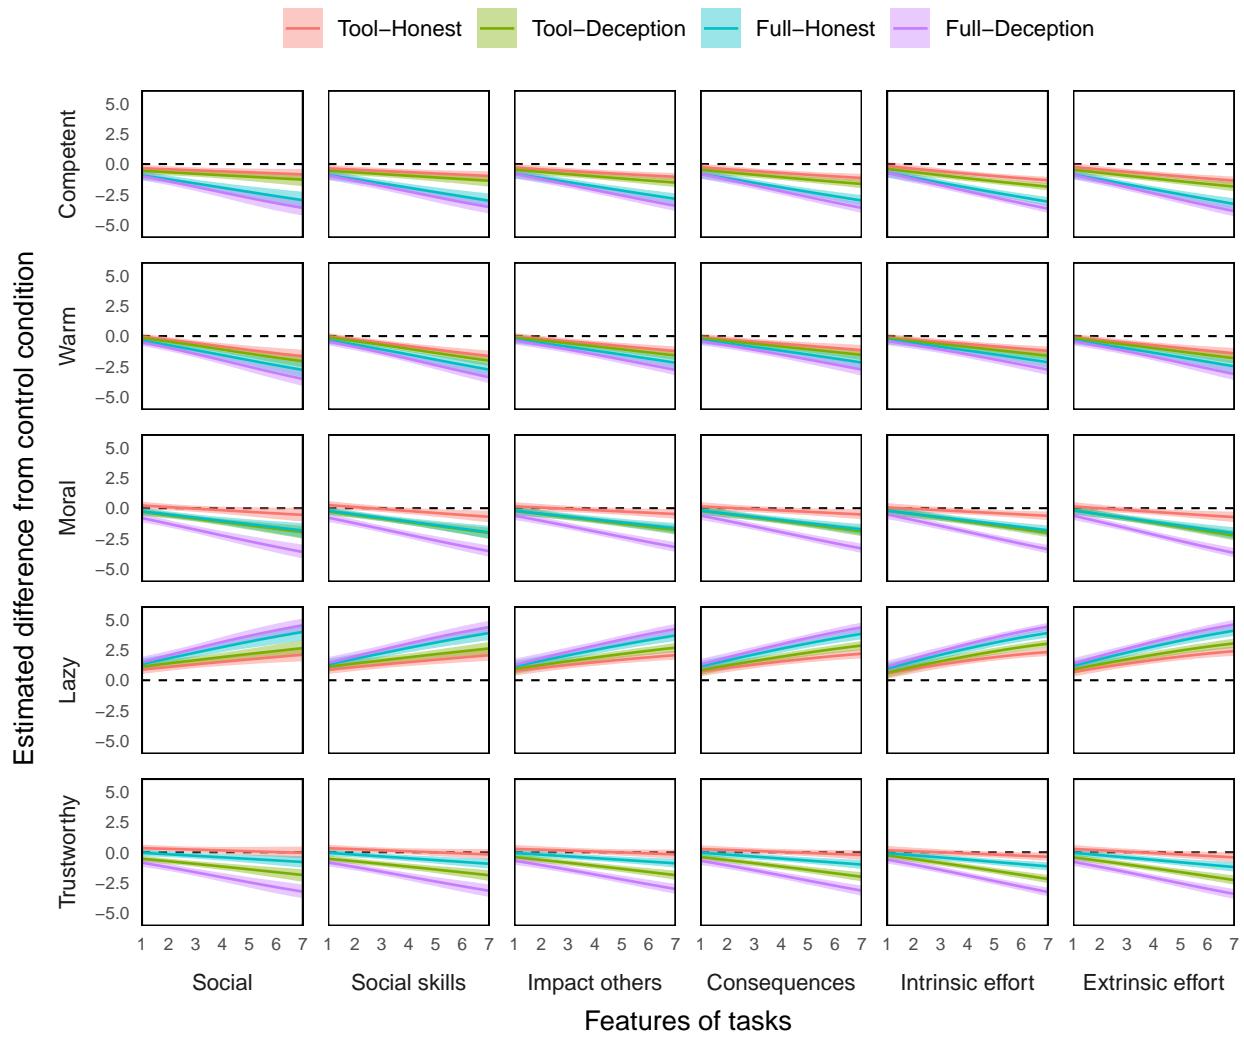
Frequency lists were created with the *Frequency List* plugin (Version 1.0.11). Settings included: unigram analysis ($N = 1$), omission of n-grams with frequency < 5 , exclusion of n-grams appearing in fewer than 0.1% of documents, filtering collocates by Normalized Pointwise Mutual Information (NPMI), and removal of collocates with metric values < 0.5 . Outputs were saved using the *Save Output to CSV* plugin (Version 1.0.5). This process was repeated separately for each condition folder.

For cross-condition comparison, we used the *Compare Frequencies* plugin (Version 1.1.02), retaining most default settings. The only modification was disabling the *Skip comparisons with 0 frequency values* option. This plugin calculates a range of comparative

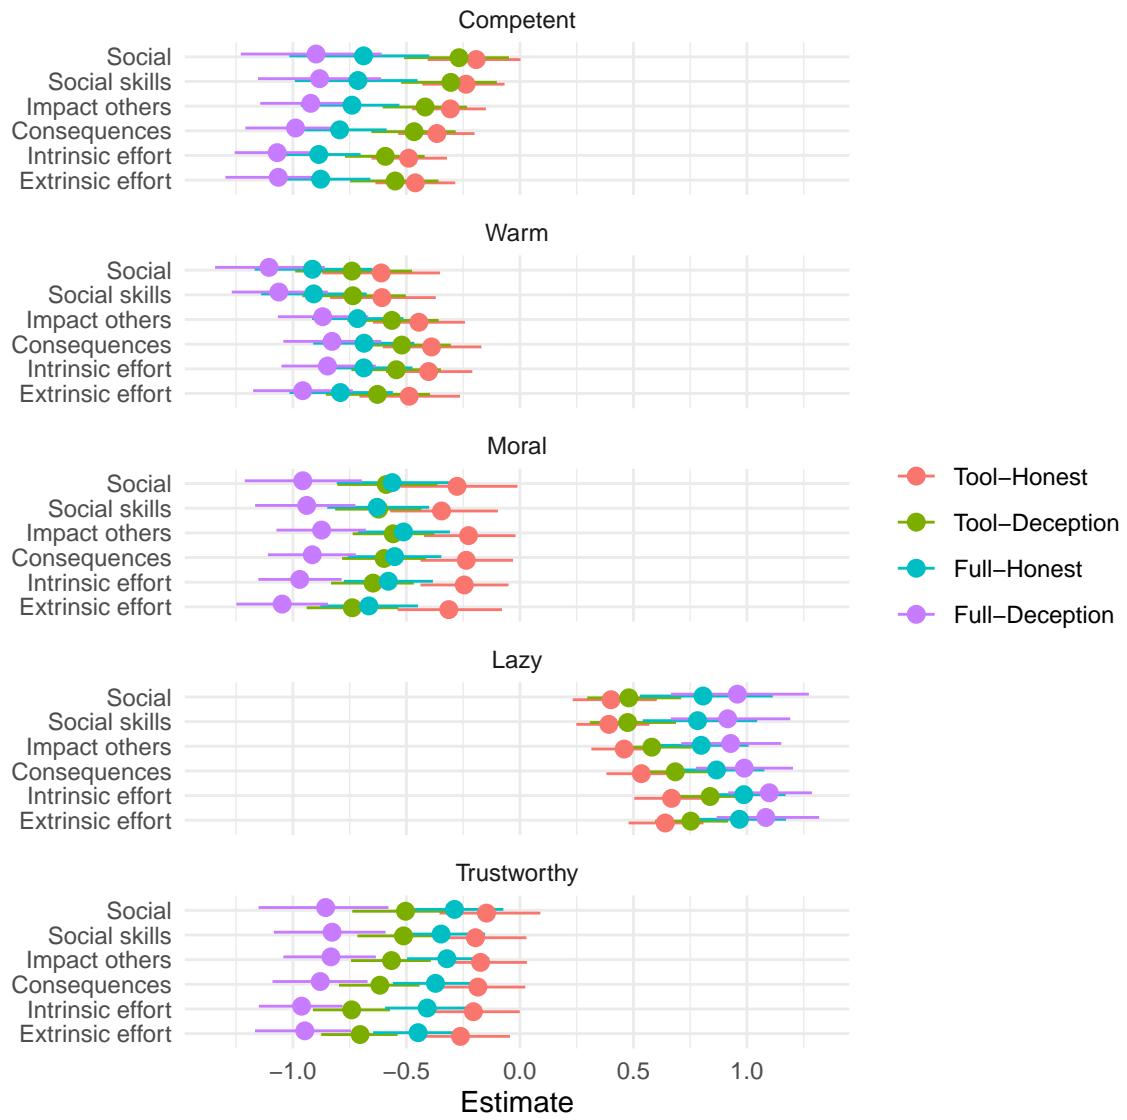
metrics, including log likelihood (LL), %DIFF, Bayes Information Criterion (BIC), relative risk (RRisk), log ratio, and odds ratio.

Following previous work (e.g., Rayson & Garside, 2000; Gregson et al., 2022), we interpret %DIFF as an indicator of effect size and direction. Frequentist statistical significance was determined using log likelihood values, with the following thresholds: $LL \geq 3.84$ ($p < .05$), $LL \geq 6.63$ ($p < .01$), $LL \geq 10.83$ ($p < .001$), and $LL \geq 15.13$ ($p < .0001$).

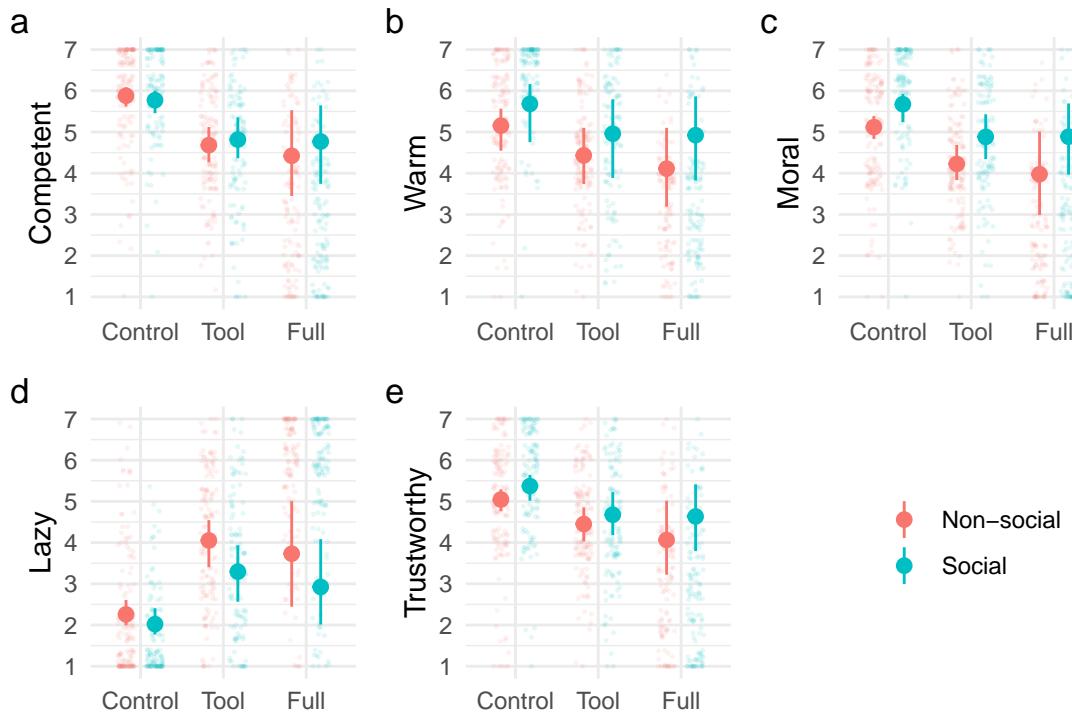
Supplementary Figures



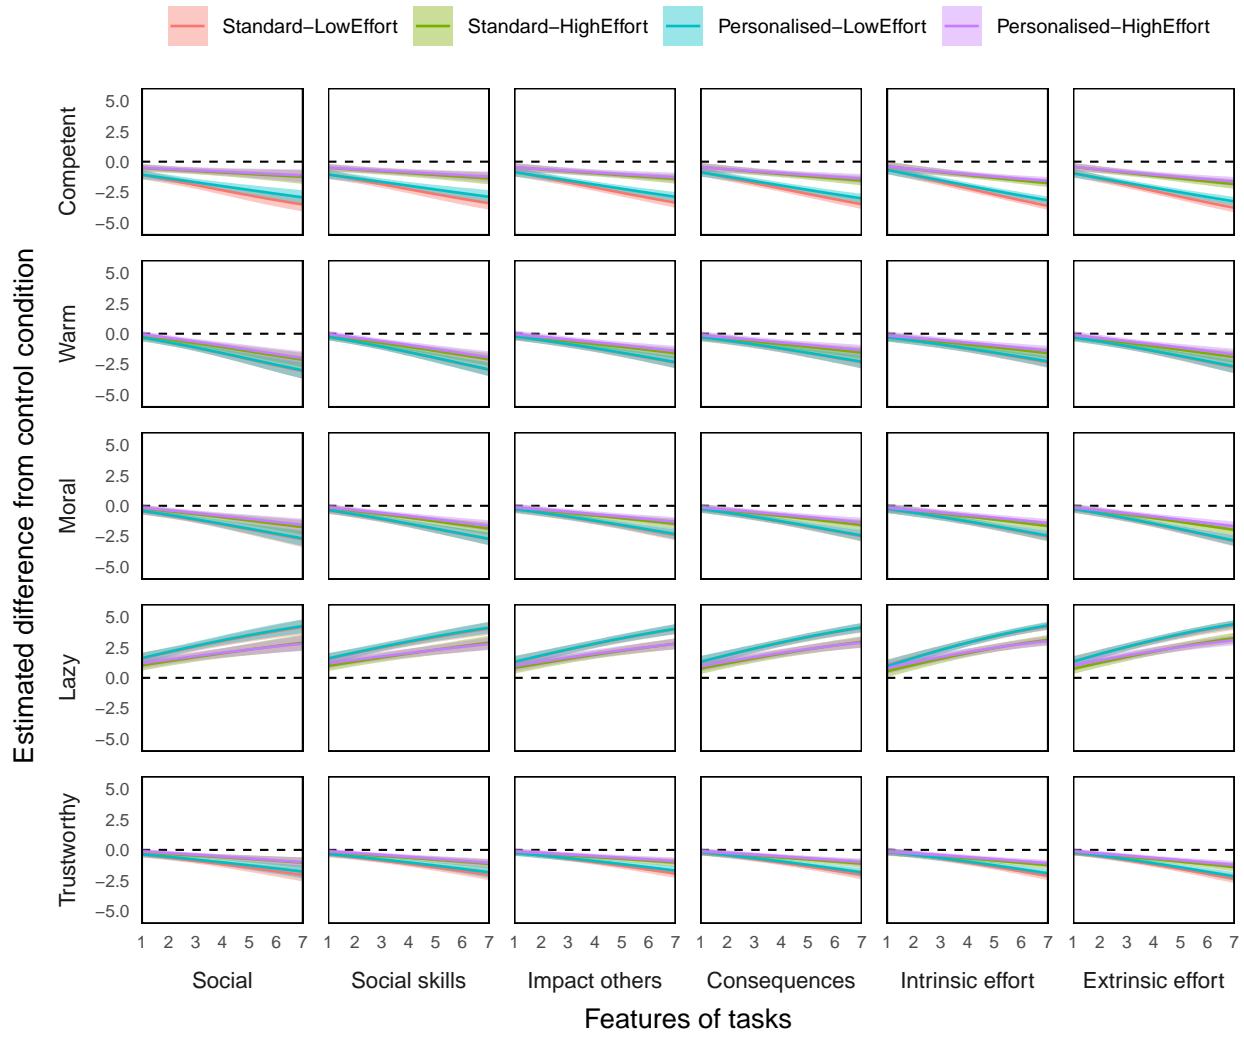
Supplementary Figure 1: The impact of task-specific features (e.g., being a social task) on the causal effects of outsourcing to AI compared to the control condition in Study 1. The y-axis reflects the estimated differences between the experimental conditions and the control condition (dashed line) on a 7-point Likert scale. Lines and shaded areas represent posterior medians and 95% credible intervals, respectively. The patterns indicate, for example, more negative effects of outsourcing on character evaluations for tasks that are rated as more social.



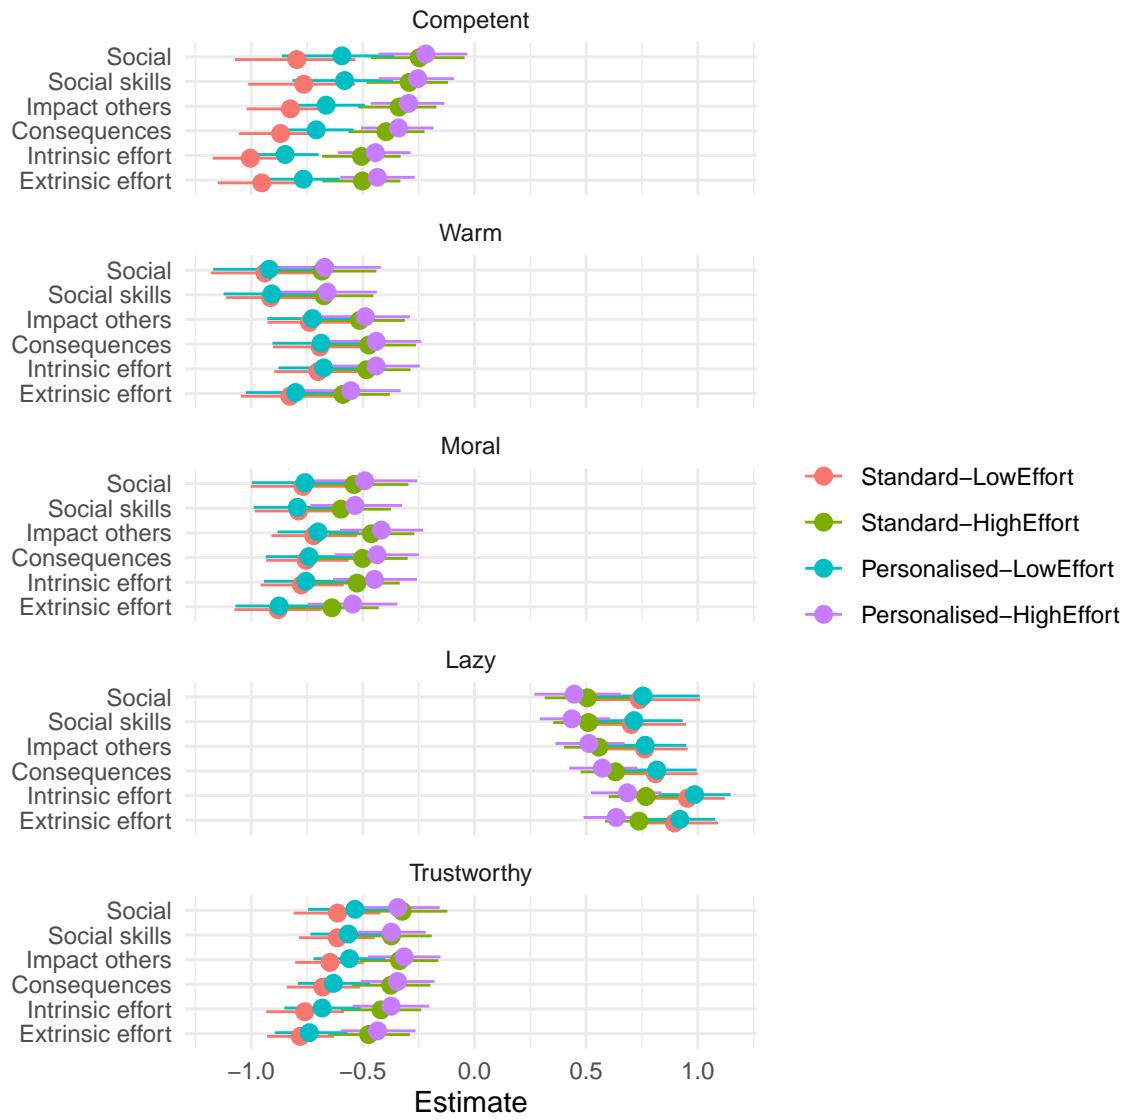
Supplementary Figure 2: Interaction parameters from models including task-specific features as moderators of the causal effects of AI outsourcing compared to the control condition in Study 1. Points and line ranges represent posterior medians and 95% credible intervals, respectively.



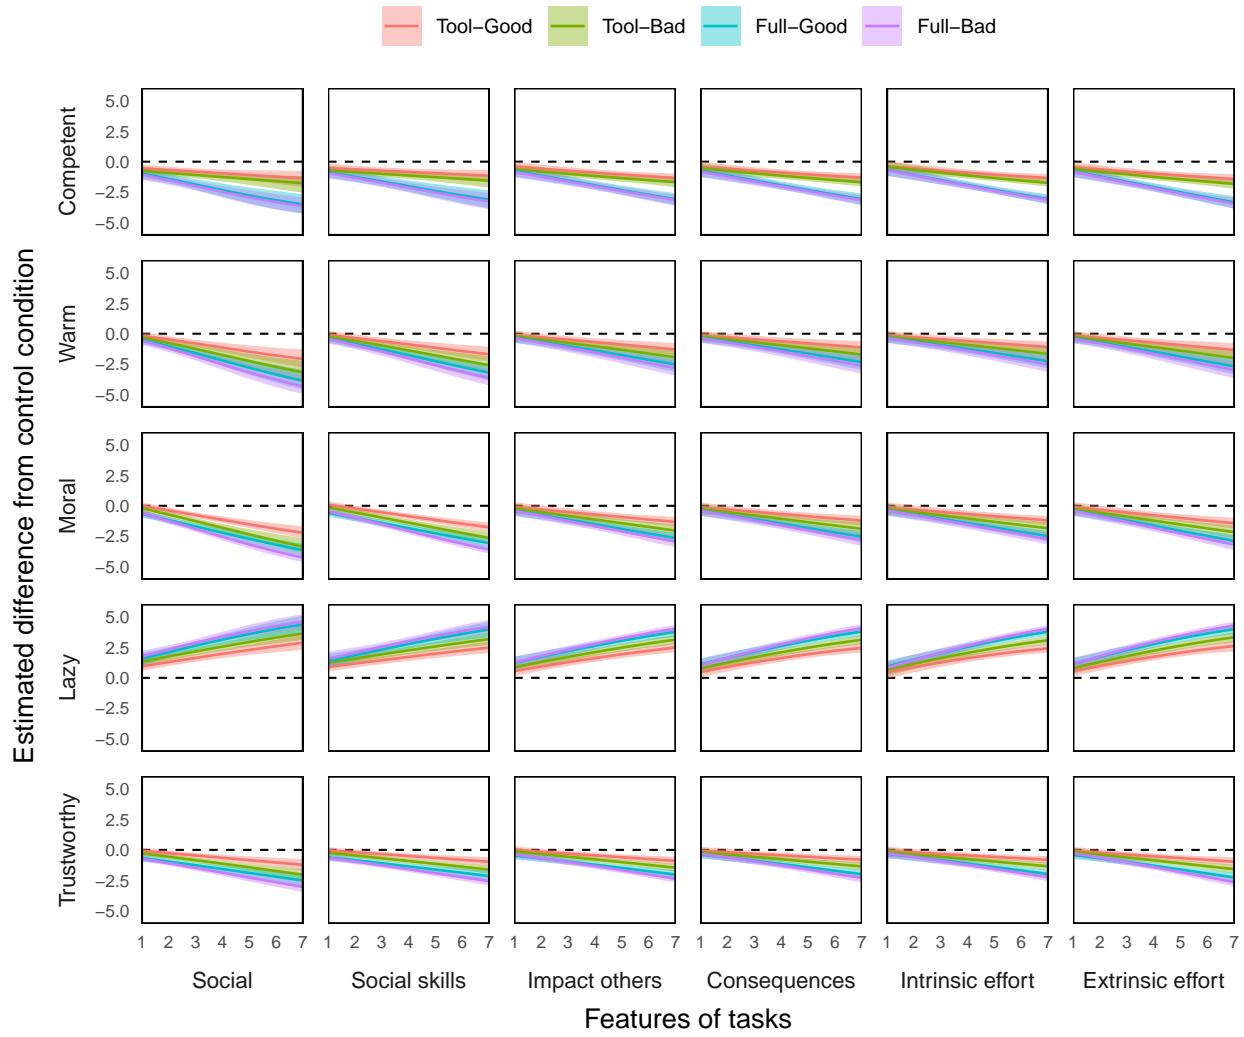
Supplementary Figure 3: Character evaluations in Study 2. Participants in the control condition, the tool outsourcing condition, and the full outsourcing condition evaluated the “other participant” on (a) competence, (b) warmth, (c) morality, (d) laziness, and (e) trustworthiness. Jittered points represent participant responses to the questions, split by whether the writing task was a non-social task (red) or a social task (blue). Point ranges are estimated marginal means from the fitted model, pooling over essay answers. Points and line ranges represent posterior medians and 95% credible intervals, respectively.



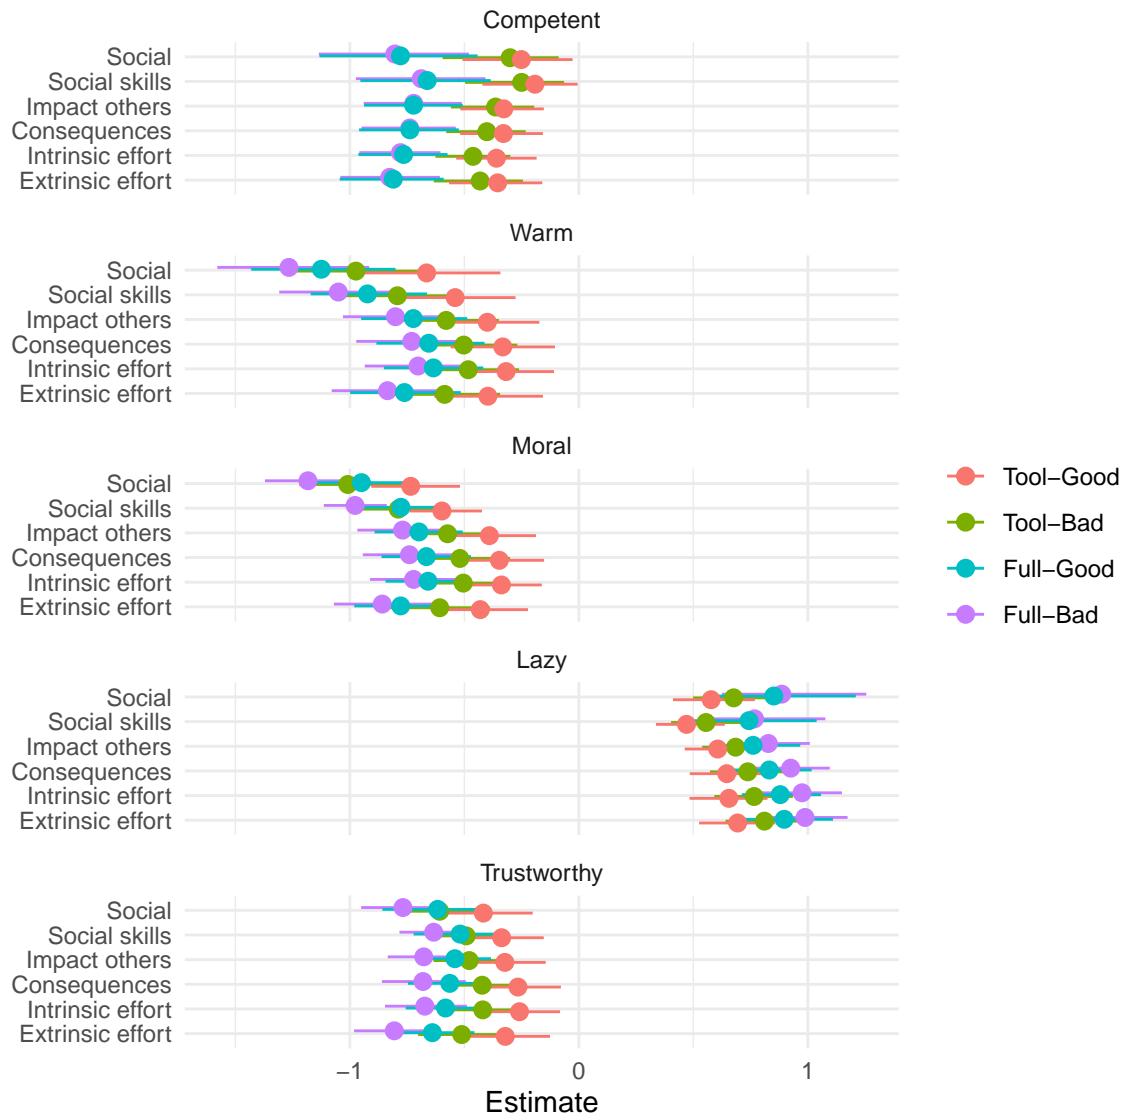
Supplementary Figure 4: The impact of task-specific features (e.g., being a social task) on the causal effects of outsourcing to AI compared to the control condition in Study 3. The y-axis reflects the estimated differences between the experimental conditions and the control condition (dashed line) on a 7-point Likert scale. Lines and shaded areas represent posterior medians and 95% credible intervals, respectively. The patterns indicate, for example, more negative effects of outsourcing on character evaluations for tasks that are rated as more social.



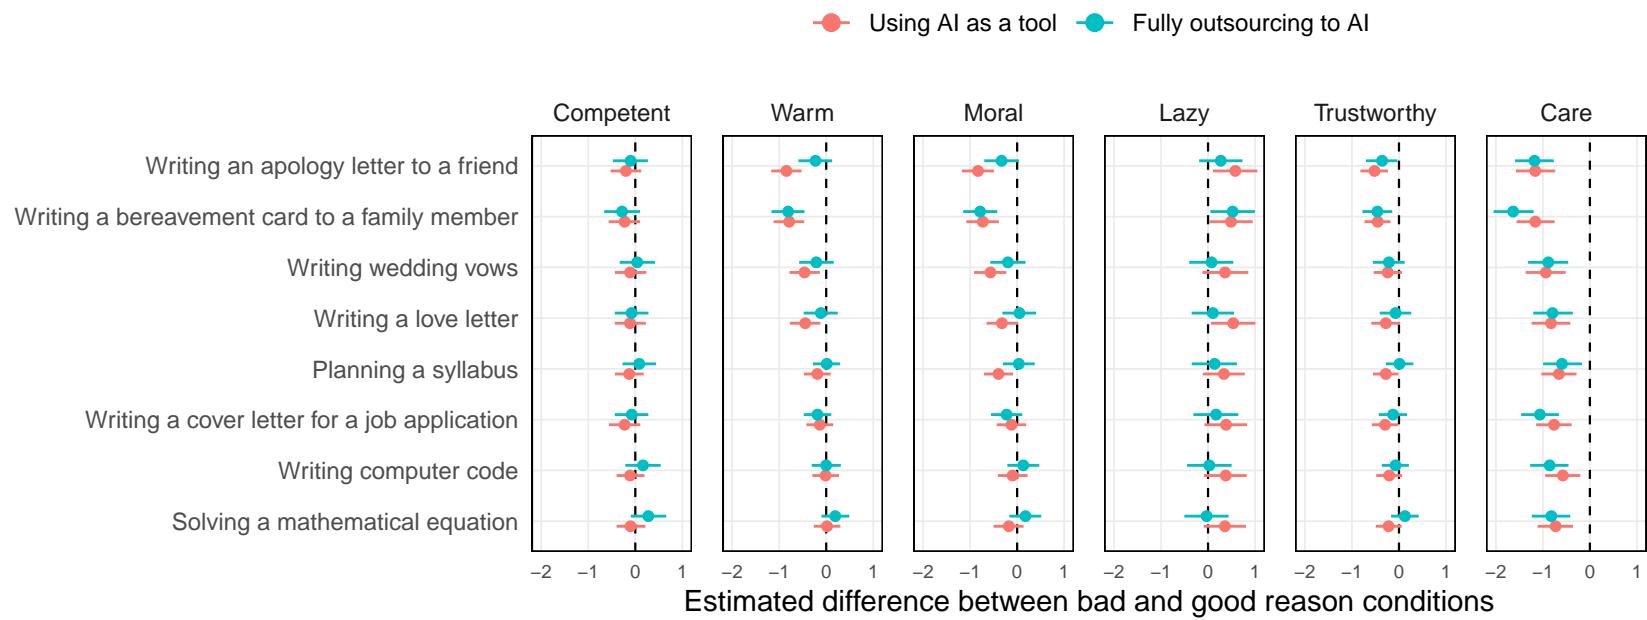
Supplementary Figure 5: Interaction parameters from models including task-specific features as moderators of the causal effects of AI outsourcing compared to the control condition in Study 3. Points and line ranges represent posterior medians and 95% credible intervals, respectively.



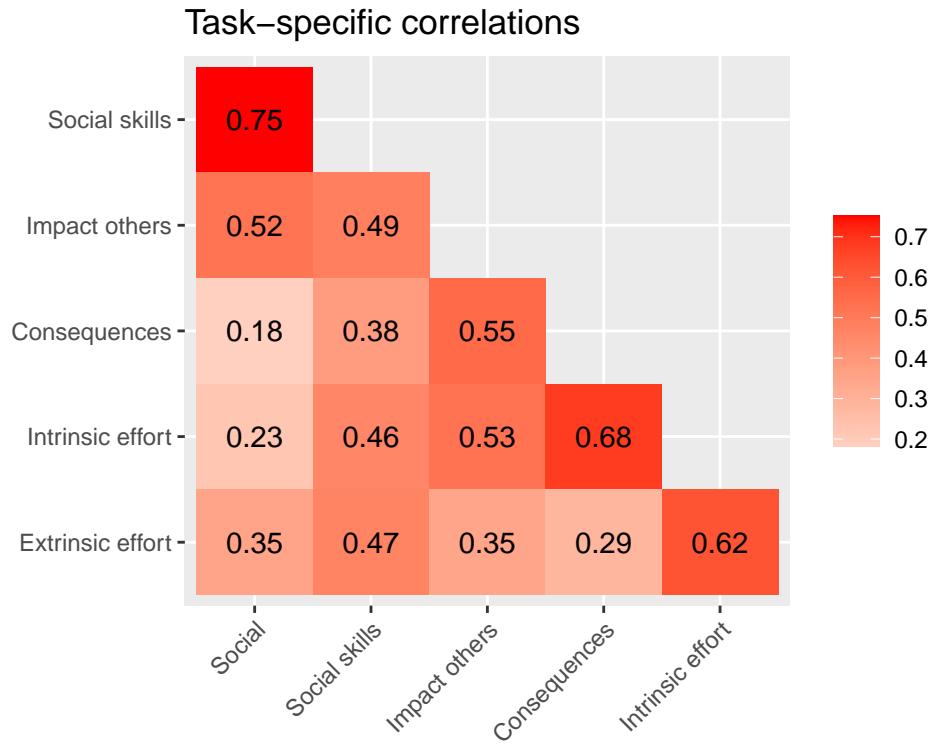
Supplementary Figure 6: The impact of task-specific features (e.g., being a social task) on the causal effects of outsourcing to AI compared to the control condition in Study 4. The y-axis reflects the estimated differences between the experimental conditions and the control condition (dashed line) on a 7-point Likert scale. Lines and shaded areas represent posterior medians and 95% credible intervals, respectively. The patterns indicate, for example, more negative effects of outsourcing on character evaluations for tasks that are rated as more social.



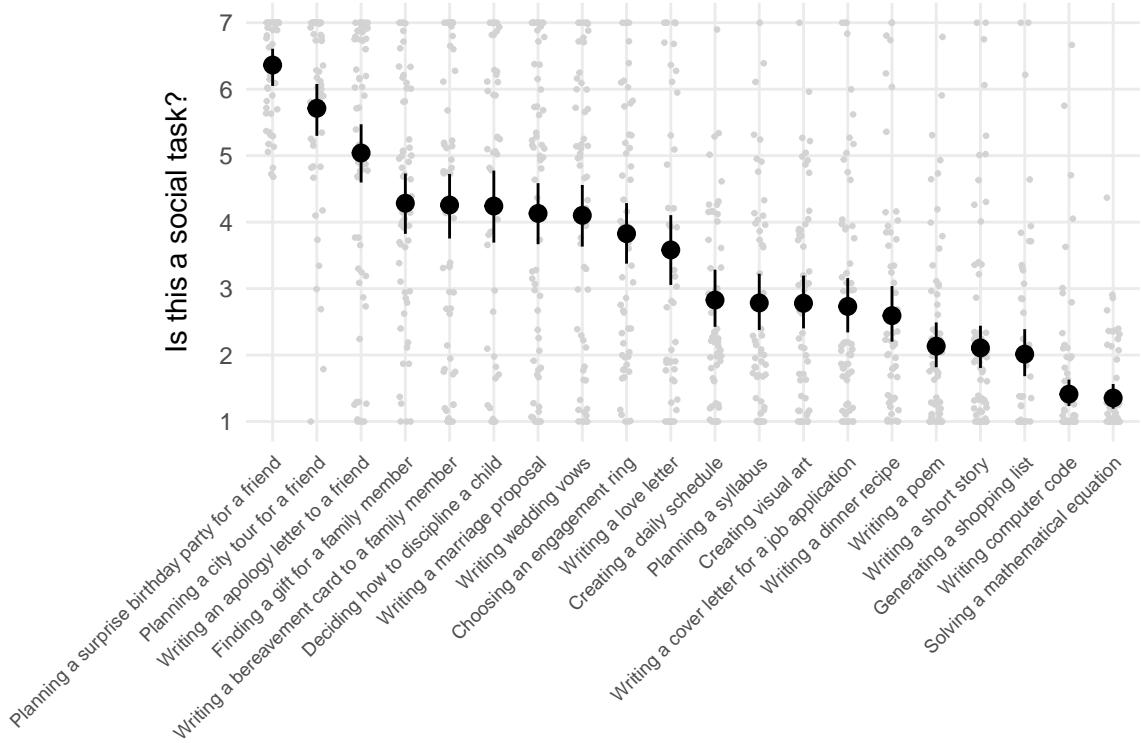
Supplementary Figure 7: Interaction parameters from models including task-specific features as moderators of the causal effects of AI outsourcing compared to the control condition in Study 4. Points and line ranges represent posterior medians and 95% credible intervals, respectively.



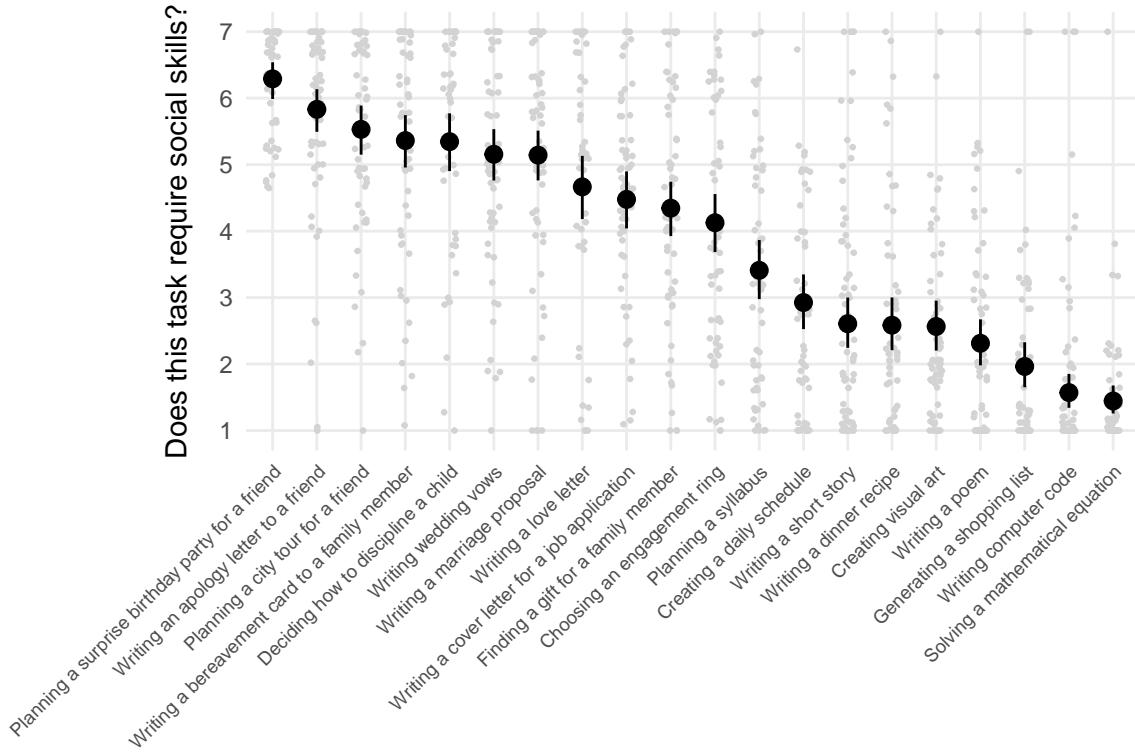
Supplementary Figure 8: Variation in the effect of reasons across tasks in Study 4. Tasks are ordered from most social (top) to least social (bottom) according to ratings from a pilot study. Point ranges are differences in marginal means on a 7-point Likert scale between the “bad reason” and “good reason” conditions, split by outsourcing type. Points and ranges represent posterior medians and 95% credible intervals, respectively.



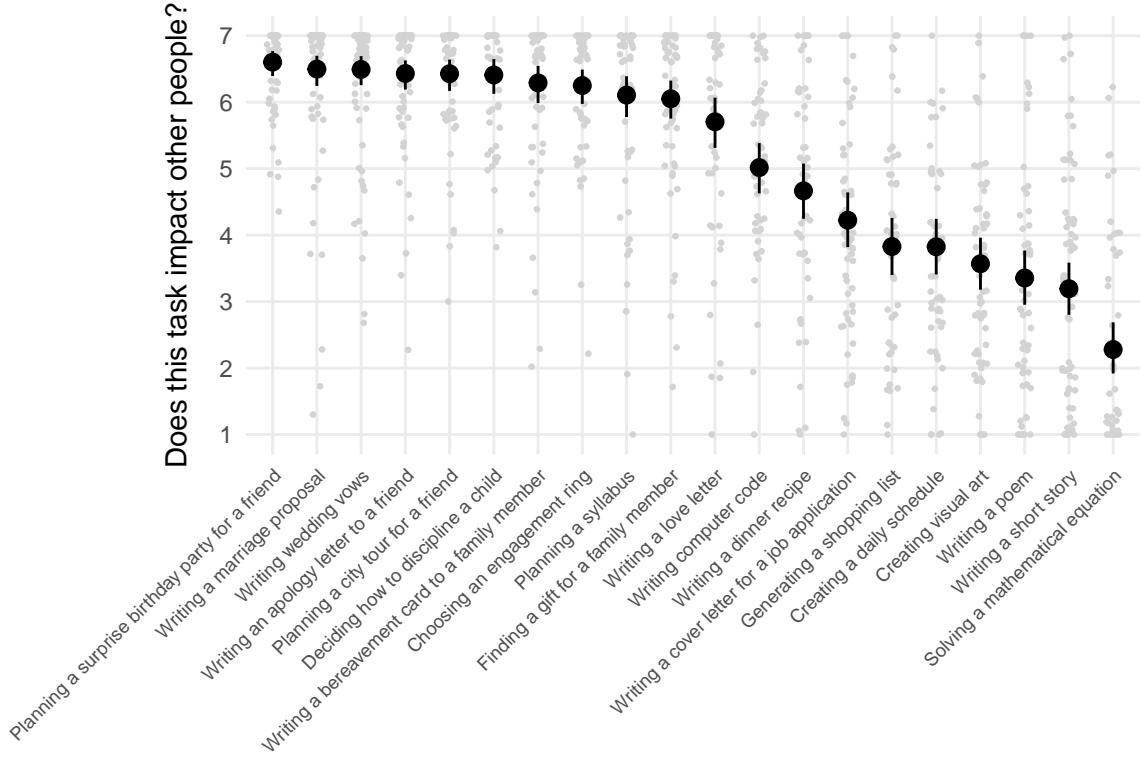
Supplementary Figure 9: Model-estimated task-specific correlations between all six questions in the first pilot study. Values are posterior median correlations. A positive correlation indicates that tasks that are rated highly on one question tend to be rated highly on another question.



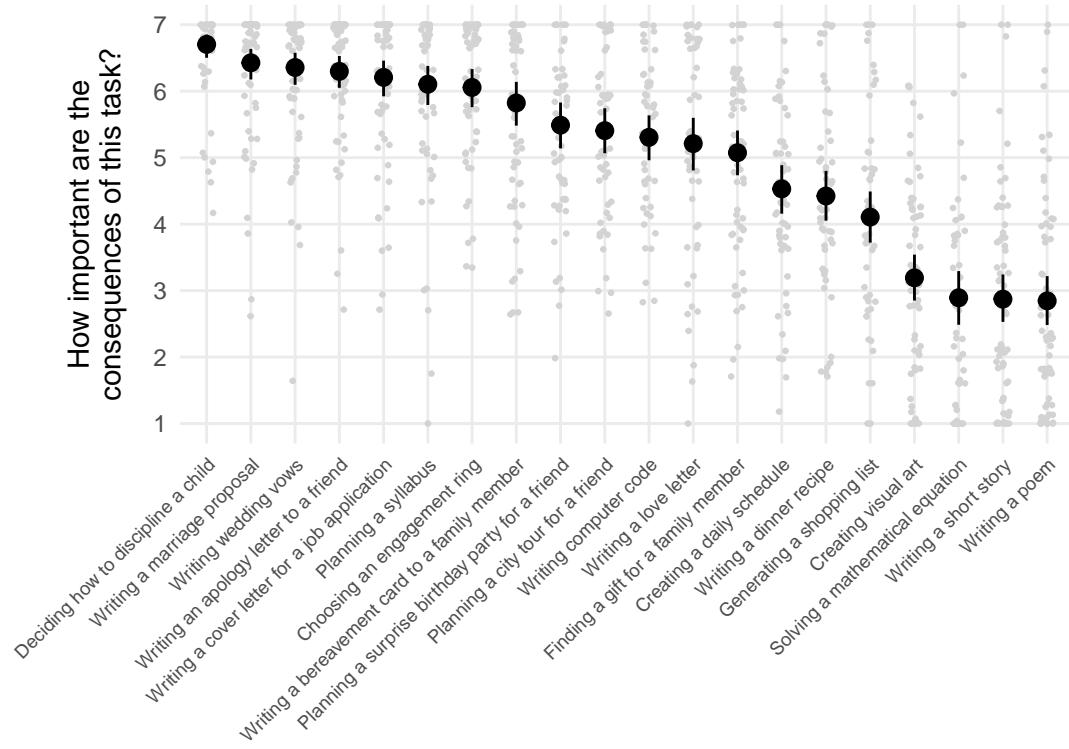
Supplementary Figure 10: Model-estimated means for the question “Is this a social task?” across all 20 tasks in the first pilot study. Grey points represent participant responses to the question, jittered for easier viewing. Black points are estimated means from the fitted model, pooling over participants. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.



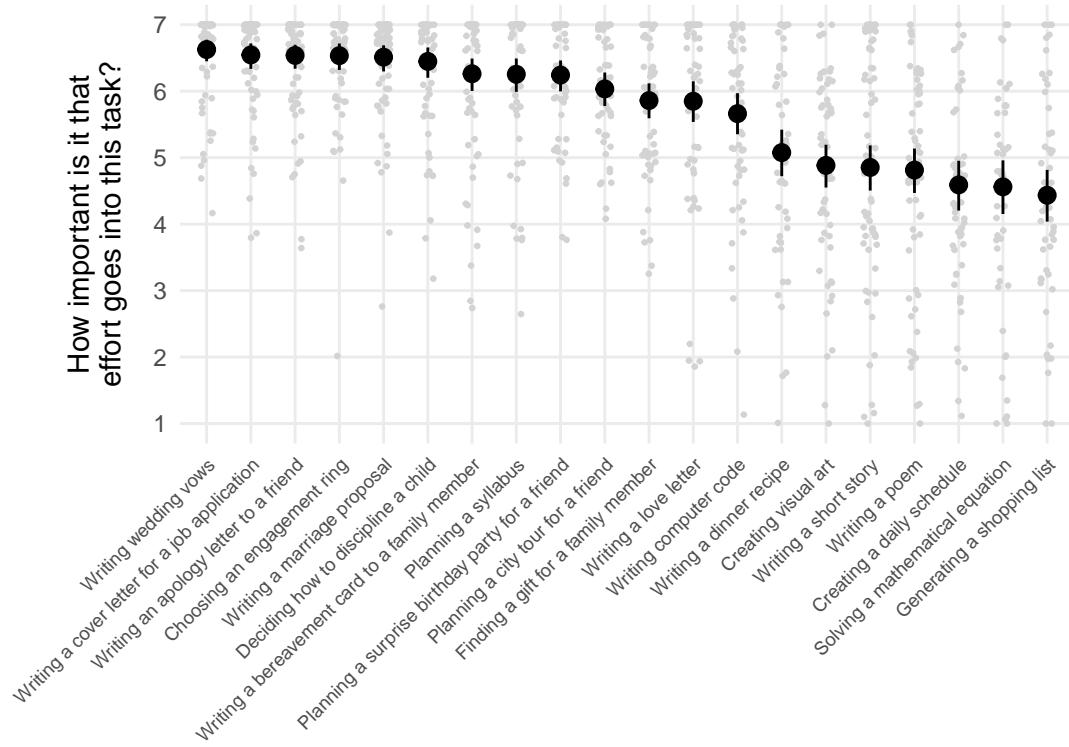
Supplementary Figure 11: Model-estimated means for the question “Does this task require social skills?” across all 20 tasks in the first pilot study. Grey points represent participant responses to the question, jittered for easier viewing. Black points are estimated means from the fitted model, pooling over participants. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.



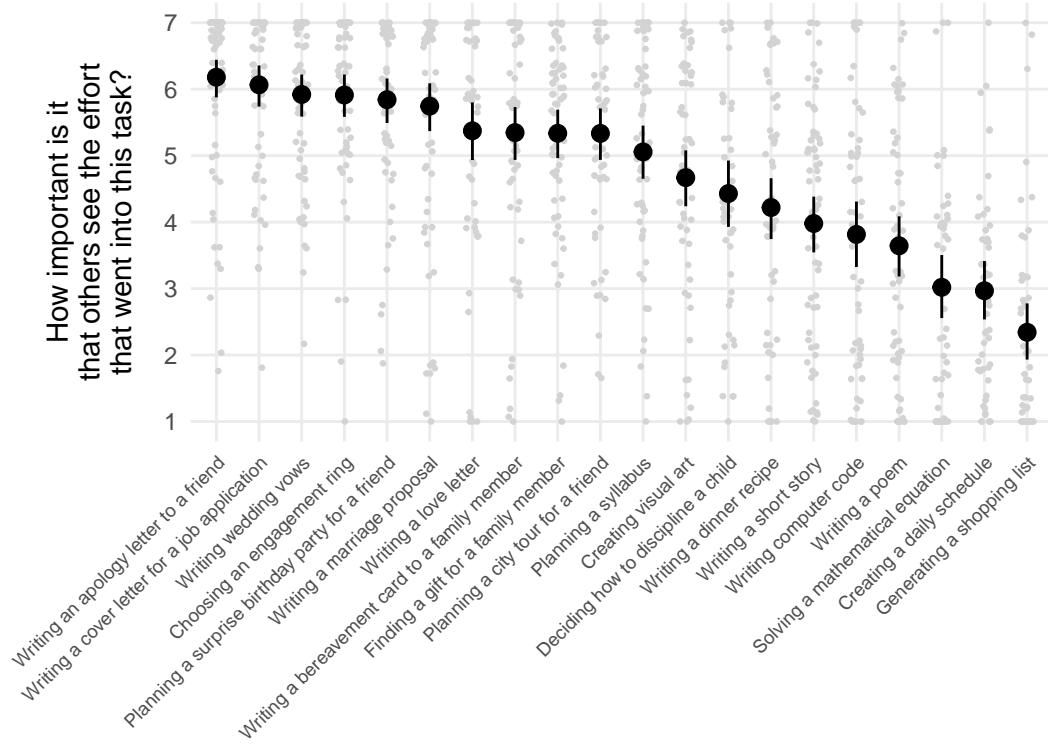
Supplementary Figure 12: Model-estimated means for the question “Does this task impact other people?” across all 20 tasks in the first pilot study. Grey points represent participant responses to the question, jittered for easier viewing. Black points are estimated means from the fitted model, pooling over participants. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.



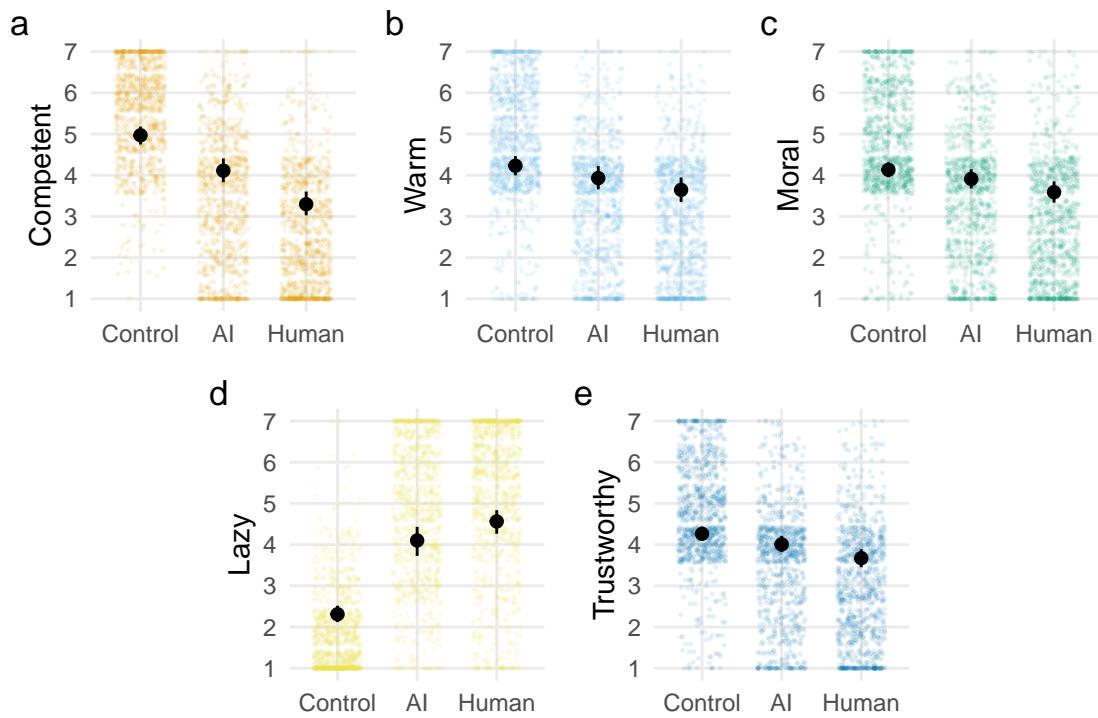
Supplementary Figure 13: Model-estimated means for the question “How important are the consequences of this task?” across all 20 tasks in the first pilot study. Grey points represent participant responses to the question, jittered for easier viewing. Black points are estimated means from the fitted model, pooling over participants. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.



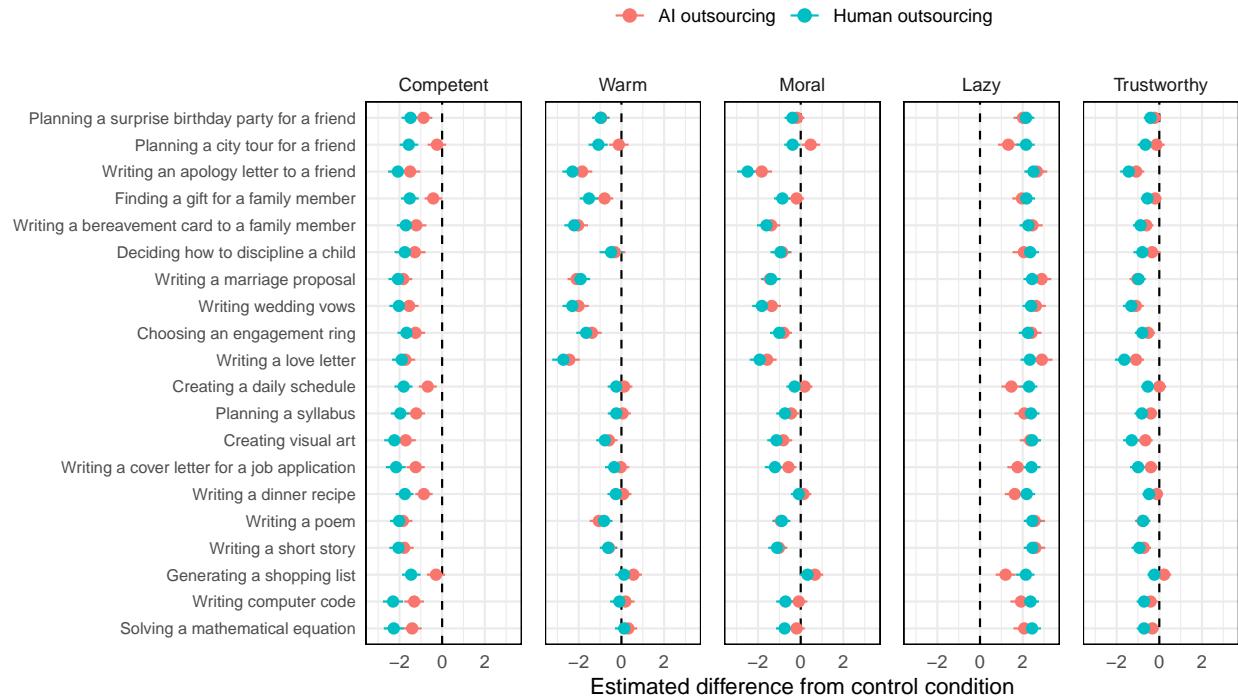
Supplementary Figure 14: Model-estimated means for the question “How important is it that effort goes into this task?” across all 20 tasks in the first pilot study. Grey points represent participant responses to the question, jittered for easier viewing. Black points are estimated means from the fitted model, pooling over participants. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.



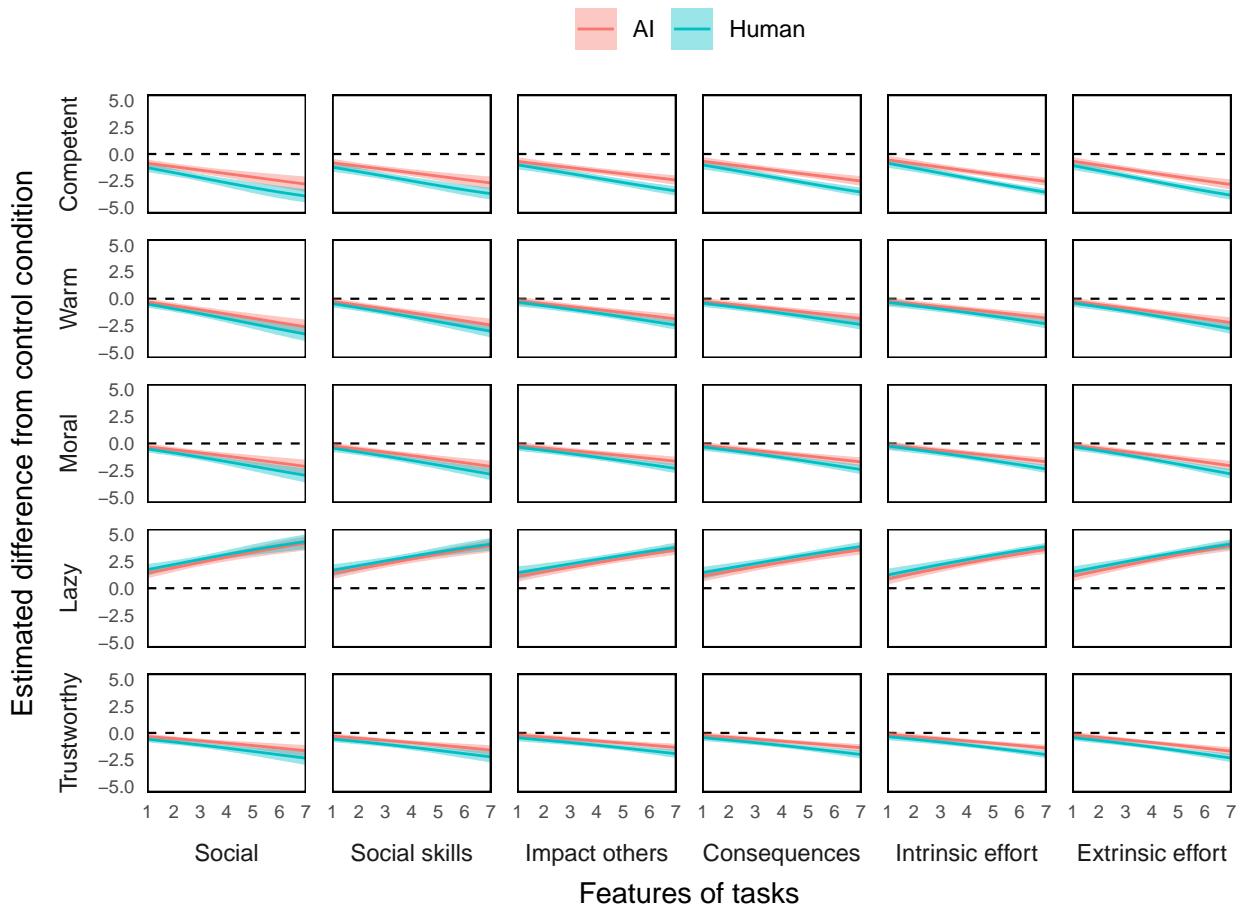
Supplementary Figure 15: Model-estimated means for the question “How important is it that others see the effort that goes into this task?” across all 20 tasks in the first pilot study. Grey points represent participant responses to the question, jittered for easier viewing. Black points are estimated means from the fitted model, pooling over participants. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.



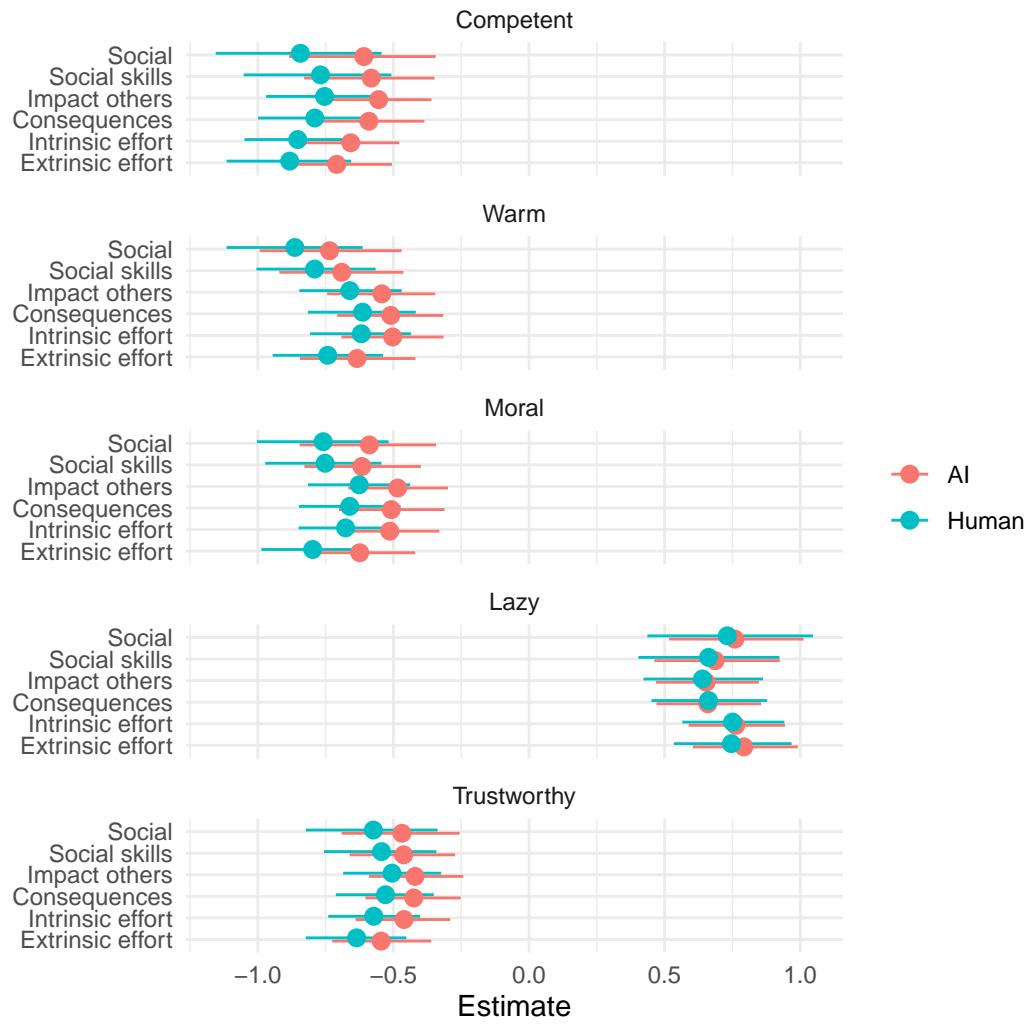
Supplementary Figure 16: Character evaluations in the second pilot study. Participants in the control condition, the AI outsourcing condition, and the human outsourcing condition evaluated people in the scenarios on (a) competence, (b) warmth, (c) morality, (d) laziness, and (e) trustworthiness. Coloured points represent participant responses to the questions, jittered for easier viewing. Black points are estimated marginal means from the fitted model, pooling over participants and tasks. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.



Supplementary Figure 17: Variation in the effects of outsourcing across tasks in the second pilot study. Tasks are ordered from most social (top) to least social (bottom) according to ratings from the first pilot study. Point ranges are differences in marginal means on a 7-point Likert scale for the AI outsourcing condition (red) and the human outsourcing condition (blue) compared to the control condition. Points and ranges represent posterior medians and 95% credible intervals, respectively.



Supplementary Figure 18: The impact of task-specific features (e.g., being a social task) on the causal effects of outsourcing to AI (red) and humans (blue) compared to the control condition in the second pilot study. The y-axis reflects the estimated differences between the experimental conditions and the control condition (dashed line) on a 7-point Likert scale. Lines and shaded areas represent posterior medians and 95% credible intervals, respectively. The patterns indicate, for example, more negative effects of outsourcing on character evaluations for tasks that are rated as more social.



Supplementary Figure 19: Interaction parameters from models including task-specific features as moderators of the causal effects of AI outsourcing (red) and human outsourcing (blue) compared to the control condition in the second pilot study. Points and line ranges represent posterior medians and 95% credible intervals, respectively.

Supplementary Tables

Supplementary Table 1: Tasks included in the studies.

Task	Pilot Study 1	Pilot Study 2	Study 1	Study 2	Study 4
Writing wedding vows	✓	✓	✓	✓	✓
Writing a love letter	✓	✓	✓	✓	✓
Writing a marriage proposal	✓	✓	✓	✓	
Choosing an engagement ring	✓	✓	✓		
Finding a gift for a family member	✓	✓	✓		
Deciding how to discipline a child	✓	✓	✓		
Writing a bereavement card to a family member	✓	✓	✓	✓	✓
Writing an apology letter to a friend	✓	✓	✓	✓	✓
Planning a city tour for a friend	✓	✓	✓	✓	
Planning a surprise birthday party for a friend	✓	✓	✓	✓	
Writing a cover letter for a job application	✓	✓	✓	✓	✓
Writing computer code	✓	✓	✓	✓	✓
Solving a mathematical equation	✓	✓	✓	✓	✓
Planning a syllabus	✓	✓	✓	✓	✓
Writing a short story	✓	✓	✓	✓	
Writing a poem	✓	✓	✓	✓	
Creating visual art	✓	✓	✓		
Creating a daily schedule	✓	✓	✓	✓	
Generating a shopping list	✓	✓	✓	✓	
Writing a dinner recipe	✓	✓	✓	✓	

Supplementary Table 2: Pre-generated essay answers to the social prompt in Study 2.

Answer	Text
Father	<p>My dad is one of the most important people in my life. He's always been someone I look up to and rely on. Throughout my whole life, he's been there to guide me, teach me, and support me in everything I do. What makes my dad special is how much he cares about our family. He works hard every day to make sure we have what we need, but no matter how busy he is, he always makes time for us. My dad is emotionally strong. Even though he doesn't show his emotions a lot, I can tell how much he cares by how much he does for us. When things get hard, he stays calm and steady, and that helps me feel better. One of my favorite things about my dad is how much he loves to teach. He knows so much and is always happy to share what he knows. He explains things in a way that makes sense and is easy to understand. I also love my dad's sense of humor. He always knows how to make me laugh with a joke or a funny story. His laughter makes everything feel lighter and happier. My dad has taught me so much about working hard, being kind, and staying strong when life is tough. I'm so thankful for everything he's done for me, and I'm proud to have him as my dad!</p>
Sister	<p>My sister is one of the most important people in my life. She is special because she always supports me. She has a way of making me feel confident, even when I'm unsure of myself. Whenever I'm scared to try something new, she's the first to remind me of what I can do. Her belief in me helps me believe in myself. My sister also has a really kind heart. She always thinks about others and does her best to help. She's always putting others first, whether it's being there for a friend or helping out with family. Her kindness is something I look up to and try to follow. Another thing I love about my sister is how funny she is. She has a great sense of humor and always knows how to make people laugh, even in serious moments. If I'm ever feeling down, she can cheer me up with a joke or a funny story. Her laughter makes everything feel lighter and happier. What I admire most about my sister is how strong she is. She's faced tough times but never lets them hold her back. Her strength gives me courage to keep going when life gets hard. My sister is more than just a family member — she's my role model and my rock!</p>
Friend	<p>My best friend is one of the most amazing people I know. She's someone I can count on no matter what. What makes her so special is her kindness. She always makes people feel important and cared for. Whether it's helping someone she just met or being there for her friends, she's the first to offer support. She never hesitates to help me, whether I'm upset or just having a bad day. She also has a great sense of humor that can cheer anyone up. She finds ways to laugh about even the smallest things, and her laugh is so contagious! Her laughter makes everything feel lighter and happier. What I admire most about her is how strong she is. Life hasn't always been easy for her, but she never gives up. She stays calm and keeps going, no matter what happens. Watching her face challenges in adulthood has taught me to be brave and not let hard times hold me back. My best friend has shown me what it means to be loyal, caring, and strong. I feel so lucky to have her in my life. I try to be as good of a friend to her as she is to me. She inspires me to be a better person!</p>

Supplementary Table 3: Pre-generated essay answers to the non-social prompt in Study 2.

Answer	Text
The Hobbit	I will focus on describing the book “The Hobbit” by Tolkien. The Hobbit is a fantasy adventure story about Bilbo Baggins. Bilbo is a quiet hobbit who lives in the Shire. His life changes when Gandalf the wizard and a company of dwarves ask him to join their quest to take back treasure stolen by a dragon. At the beginning of the journey, Bilbo and the dwarves are nearly eaten by trolls, but Gandalf saves them. Then later, in the Misty Mountains, Bilbo meets a creature called Gollum and finds a magical ring that makes him invisible. This ring later becomes very important in “The Lord of the Rings”. As they travel, the group fights goblins, giant spiders in Mirkwood forest, and they get captured by Wood-elves. Bilbo shows his bravery by saving the group several times. Finally, they reach the Lonely Mountain where the dragon Smaug lives. Bilbo sneaks into the dragon’s lair and finds a weak spot in Smaug’s armor. The dragon gets angry and attacks the nearby town by a lake. Eventually, Smaug is killed. With the dragon dead, humans, elves, and dwarves all want the treasure. This leads to the “Battle of the Five Armies”. Tolkien doesn’t describe the battle in too much detail, but we later learn that the leader of the dwarves Thorin has fought bravely and died from his wounds. At the end of the story, Bilbo returns home to the Shire, richer and wiser from his adventure. He is happy to be back in his quiet life, and sets out to write a book of his adventures - which sets the stage for the sequel, The Lord of the Rings.
Buffy the Vampire Slayer	I will focus on describing the TV show “Buffy the Vampire Slayer”. Buffy is a TV show that completely flips the script on traditional high school dramas and supernatural horror. It’s about a teenager, Buffy Summers, who’s tasked with being the Slayer – basically a chosen one who hunts vampires and other demons. But what sets the show apart is how Buffy struggles to balance her responsibility with the regular teenage experience. She’s not just fighting creatures of the night, she’s also balancing school and friendships at the same time. One of the most striking things about Buffy is how layered the characters are. Buffy is tough and witty, but she’s also vulnerable. She’s faced with loss, guilt, and trying to make sense of her life outside of the supernatural chaos. And then there’s her team. Willow is the nerdy, sweet heart of the group, Xander is the funny loyal friend, and Giles (Buffy’s Watcher) is the stern mentor who’s also loving. Each character feels real, with their own flaws and growth arcs. The show has this incredible ability to mix humor, heart, and horror seamlessly. The dialogue is sharp and full of clever pop culture references. Yet, the writing isn’t afraid to get serious, exploring themes like trauma and growing up. The monsters Buffy faces often mirror real-life challenges, making the stakes feel personal. I love Buffy. It’s a show that’s smart and emotional, blending witty banter with moments of real depth. It’s got a cult following for a reason!
Titanic	I will focus on describing the film “Titanic”. The genre is a mix of romance, disaster, and historical tragedy. The film tells the love story of Jack and Rose, two passengers from different social classes aboard the passenger ship Titanic. Jack is a poor artist, but he manages to win a ticket to the ship’s maiden voyage. Rose is a young upper-class woman who is feeling trapped in her engagement to her fiance. Jack and Rose cross paths on the ship, and they fall in love. The film balances the spectacle of the ship’s design and atmosphere with the tension that gradually builds as the audience knows what fate awaits. The Titanic sails into the icy waters of the Atlantic and strikes an iceberg. Chaos immediately erupts. The film allows viewers to experience the terror, confusion, and heartbreak of the tragedy, showcasing both personal stories and the broader catastrophe. At its core, the film is a romance. But Titanic also touches on themes of class and fate. It highlights the disparity between the elite and the working-class passengers who are doomed to different fates. The film also explores the sense of inevitability that comes with knowing the ship’s doom. The most iconic scene from the film is arguably the scene where Jack and Rose stand together at the bow, arms outstretched. They seem free, but the scene also foreshadows the devastating crash to come. The film is truly heartbreakingly tragic!

Supplementary Table 4: Reading times and comprehension rates for the essay answers in Study 2. Expected reading times were calculated based on an estimated reading speed of 275 words per minute. Comprehension rates are the percentage of participants who answered the comprehension question correctly.

Prompt	Answer	Number of words	Expected reading time (secs)	Average reading time (secs)	Comprehension (%)
Social	Father	234	51.05	47.25	100.00
	Friend	211	46.04	49.91	98.50
	Sister	218	47.56	50.25	99.21
Non-social	Buffy	251	54.76	62.46	100.00
	Hobbit	278	60.65	65.97	99.28
	Titanic	239	52.15	63.90	100.00

Supplementary Table 5: Percentage of participants in Study 2 who passed the manipulation check and reported that they believed the manipulation, split by condition.

Condition	Pass manipulation check (%)	Believe manipulation (%)
Control	98.10	71.10
Tool outsourcing	96.11	77.99
Full outsourcing	100.00	86.07

Supplementary Table 6: Pairwise contrasts for character evaluations in Study 2. Numbers reflect differences in marginal means on a 7-point Likert scale. Estimates are pooled over essay answers. The bottom rows represent the interactions between outsourcing type and task type. Main numbers are posterior medians, numbers in the square brackets are 95% credible intervals.

	Response				
	Competent	Warm	Moral	Lazy	Trustworthy
Effect of outsourcing type					
Task type = Social					
Tool Social - Control Social	-0.97 [-1.33 -0.45]	-0.74 [-1.25 -0.12]	-0.81 [-1.16 -0.30]	1.31 [0.59 1.76]	-0.71 [-1.07 -0.21]
Full Social - Control Social	-0.99 [-1.96 -0.17]	-0.75 [-1.55 -0.03]	-0.78 [-1.63 -0.04]	0.87 [0.05 1.96]	-0.73 [-1.50 0.00]
Full Social - Tool Social	-0.03 [-1.13 0.92]	-0.03 [-1.06 0.92]	0.01 [-0.95 0.87]	-0.41 [-1.38 0.84]	-0.03 [-0.91 0.77]
Task type = Non-social					
Tool Non-social - Control Non-social	-1.21 [-1.55 -0.76]	-0.74 [-1.23 -0.10]	-0.91 [-1.28 -0.42]	1.82 [1.13 2.26]	-0.60 [-0.98 -0.17]
Full Non-social - Control Non-social	-1.48 [-2.41 -0.37]	-1.07 [-1.89 -0.11]	-1.15 [-2.11 -0.11]	1.48 [0.21 2.75]	-1.00 [-1.81 -0.04]
Full Non-social - Tool Non-social	-0.28 [-1.30 0.91]	-0.33 [-1.31 0.72]	-0.25 [-1.28 0.86]	-0.31 [-1.73 1.03]	-0.40 [-1.28 0.62]
Effect of task type					
Control Social - Control Non-social	-0.11 [-0.39 0.17]	0.57 [-0.16 0.99]	0.57 [0.15 0.88]	-0.24 [-0.55 0.11]	0.33 [-0.02 0.65]
Tool Social - Tool Non-social	0.13 [-0.30 0.61]	0.55 [-0.37 1.32]	0.66 [0.11 1.15]	-0.76 [-1.35 -0.15]	0.22 [-0.28 0.76]
Full Social - Full Non-social	0.34 [-0.37 1.07]	0.83 [-0.16 1.68]	0.92 [0.13 1.68]	-0.82 [-1.72 0.12]	0.56 [-0.14 1.27]
Interaction effect					
Interaction: Tool - Control	0.24 [-0.18 0.69]	-0.01 [-0.52 0.59]	0.10 [-0.32 0.57]	-0.52 [-1.06 0.00]	-0.11 [-0.53 0.38]
Interaction: Full - Control	0.46 [-0.21 1.13]	0.29 [-0.37 0.94]	0.35 [-0.33 1.05]	-0.57 [-1.41 0.28]	0.24 [-0.41 0.90]
Interaction: Full - Tool	0.21 [-0.60 1.02]	0.29 [-0.58 1.14]	0.25 [-0.53 1.09]	-0.06 [-1.03 0.99]	0.35 [-0.42 1.14]

Supplementary Table 7: Pairwise comparisons of word frequencies between conditions. LL = log likelihood.

Word	Control Freq.	Tool Freq.	Full Freq.	%DIFF Full vs Control	LL Full vs Control	%DIFF Tool vs Control	LL Tool vs Control	%DIFF Full vs Tool	LL Full vs Tool
Lazy	0	46	82	14138.18	97.16	6061.29	42.55	131.09	21.76
Genuine	36	11	12	-71.06	16.19	-79.54	25.91	41.42	0.69
Loves	9	9	0	-95.18	10.50	-33.03	0.72	-92.80	7.21
Romantic	0	7	0	-13.18	0.00	837.59	4.42	-90.74	5.16
Thoughtful	13	0	0	-96.66	16.27	-97.42	19.99	29.64	0.02
Caring	35	12	0	-98.76	49.01	-77.04	22.85	-94.60	10.36

Supplementary Table 8: Pairwise contrasts in the second pilot study. Numbers reflect differences in marginal means on a 7-point Likert scale, pooling over participants and tasks. Main numbers are posterior medians, numbers in the square brackets are 95% credible intervals.

	Response				
	Competent	Warm	Moral	Lazy	Trustworthy
AI - Control	-0.86 [-1.16 -0.55]	-0.30 [-0.57 -0.01]	-0.23 [-0.47 0.02]	1.80 [1.42 2.14]	-0.26 [-0.44 -0.05]
Human - Control	-1.68 [-1.98 -1.35]	-0.59 [-0.87 -0.28]	-0.54 [-0.80 -0.28]	2.26 [1.90 2.58]	-0.59 [-0.81 -0.39]
Human - AI	-0.81 [-1.15 -0.46]	-0.29 [-0.61 0.04]	-0.32 [-0.63 -0.02]	0.46 [0.04 0.90]	-0.33 [-0.58 -0.10]

Supplementary References

Boyd, R. L. (2019). BUTTER: Basic unit-transposable text experimentation resource. Available from <https://www.butter.tools/>

Gregson, R., Piazza, J., & Boyd, R. L. (2022). ‘Against the cult of veganism’: Unpacking the social psychology and ideology of anti-vegans. *Appetite*, 178, 106143.
doi:[10.1016/j.appet.2022.106143](https://doi.org/10.1016/j.appet.2022.106143)

Rayson, P., & Garside, R. (2000). Comparing corpora using frequency profiling. *Proceedings of the Workshop on Comparing Corpora - Volume 9*, 1–6. Presented in Hong Kong.
doi:[10.3115/1117729.1117730](https://doi.org/10.3115/1117729.1117730)