

Negative Perceptions of Outsourcing to Artificial Intelligence

Scott Claessens, Pierce Veitch, and Jim A.C. Everett

School of Psychology, University of Kent

Author Note

Scott Claessens  <https://orcid.org/0000-0002-3562-6981>

Pierce Veitch  <https://orcid.org/0009-0005-3364-7470>

Jim A.C. Everett  <https://orcid.org/0000-0003-2801-5426>

This pre-print is currently not yet peer-reviewed, and may differ from the final version.

Word count (excl. Methods and Results) = 4619 words. Author roles were classified using the Contributor Role Taxonomy (CRediT; <https://credit.niso.org/>) as follows: Scott Claessens: conceptualization, data curation, formal analysis, investigation, methodology, visualization, writing, editing; Pierce Veitch: formal analysis, methodology, editing; Jim A.C. Everett: conceptualization, funding acquisition, methodology, supervision, editing

Correspondence concerning this article should be addressed to Jim A.C. Everett, School of Psychology, University of Kent, Keynes College, Canterbury CT2 7NP, UK, Email: j.a.c.everett@kent.ac.uk

Abstract

As artificial intelligence (AI) tools become increasingly integrated into daily life, people are beginning to outsource not only professional tasks but also socio-relational ones. Large language models like ChatGPT can generate wedding vows, speeches, and personal messages, raising questions about how individuals who use AI for such tasks are perceived by others. In this paper, we conduct five pre-registered studies with British participants ($N = 3,649$) to understand how people view those who outsource tasks to AI, and how this depends on how socio-relational the task is, whether AI is used as a tool or fully delegated to, and the acknowledgment of the AI use. We find negative perceptions of outsourcing, particularly for socio-relational tasks. We show that outsourcing makes us think more negatively about not only the person and their motivations, but also the outsourced work itself. Moreover, we provide insight into why this occurs: the reduced effort from outsourcing socio-relational tasks to AI signals that the output is less authentically one's own and that the person cares less about the task. Our research highlights the way that AI use shapes our perceptions of people, raising key philosophical questions about efficiency, authenticity, and social ties in a world filled with AI-mediated interactions.

Keywords: artificial intelligence, person perception, outsourcing, effort, trust

Negative Perceptions of Outsourcing to Artificial Intelligence

1 The widespread release of generative AI language models has transformed daily life,
2 offering the potential to perform a variety of tasks more efficiently and, in some cases, with
3 greater effectiveness than by doing them oneself. But as AI becomes more widely available,
4 people are not only using it to assist them with things like preparing dinner recipes, writing data
5 analysis code, and planning daily schedules. Increasingly, AI might be used beyond routine or
6 technical domains to instead assist in tasks that are more socio-relational in nature, like writing
7 wedding vows, apology notes, and love letters. Anecdotal evidence suggests that not only is
8 AI-outsourcing of this kind already happening, but that it potentially has serious effects on how
9 we judge others. In a recent Reddit post, a disgruntled newlywed tells the story of her husband
10 using ChatGPT to write his wedding vows, expressing her discomfort with outsourcing something
11 to AI that, to her, is deeply meaningful and a reflection of their love for one another ([miramar0](#),
12 [2024](#)). Outsourcing tasks – especially socio-relational ones – to AI tools may be efficient, but
13 could have negative consequences for person perception.

14 There is nothing new, in principle, about outsourcing tasks. For hundreds of years,
15 personal assistants have organised daily schedules, recipe-books have provided meal plans, and
16 guidebooks have created travel itineraries. In the socio-relational domain, ghostwriters have
17 long-existed, and the internet is abound with professional paid services for writing wedding vows
18 and personal speeches. AI merely supercharges what is an ancient human impulse: the push to
19 reduce mental energy by outsourcing parts of our work onto people, books, tools, or systems. But
20 even if outsourcing is an old phenomenon, the rapid shift in availability and use of AI models has
21 fundamentally changed the ease with which people can outsource work, what kinds of tasks they
22 can outsource, and the way in which they can outsource. These new developments in society
23 mean that even as an old phenomenon in new clothes, there is much we still need to know about
24 outsourcing.

25 First, we need to know how people who outsource tasks to AI are perceived. We know that
26 people are increasingly using large language models (LLMs) for a wide variety of tasks

27 (Department for Science, Innovation & Technology, 2024). Due to their ubiquity, perhaps
28 outsourcing to LLMs might not lead to negative perceptions? We are sceptical. We know that
29 people dislike it when others “free ride” or reduce effort while benefiting from collective resources
30 (e.g. Cubitt et al., 2011; Kerr, 1983) and that people’s outputs are perceived as more valuable the
31 more effort was ostensibly put into them (Kruger et al., 2004). Moreover, exertion of effort is
32 deemed morally admirable and is rewarded, even in situations where effort does not directly
33 generate additional product, quality, or economic value, suggesting that effort itself is moralised
34 (Celniker et al., 2023). Given this, even if AI tools are widely available and pitched as improving
35 efficiency, the core social psychological processes are likely to remain: someone is expending less
36 effort to achieve a task, and people value effort. Indeed, some work shows that describing
37 someone as using AI for a relational task led to the perception they expended less effort and were
38 less satisfied with their relationship (Liu et al., 2024) and other unpublished work looking at
39 perceptions of people using AI to complete academic assignments finds that using AI leads to
40 more negative perceptions of moral character and suitability as a partner (Roth & Tissot, 2025).

41 Second, we need to know whether the *type* of task that people are outsourcing matters.
42 One might expect outsourcing to be perceived negatively regardless of the type of task being
43 outsourced – if effort is generally moralised, then the domain in which it is expended (or not)
44 should have little impact. However, there are reasons to expect differences between social tasks
45 like writing vows and non-social tasks like writing computer code. We know that different norms,
46 standards, and expectations can be applied to social and non-social tasks and exchanges (e.g. A. P.
47 Fiske, 1992; Heider, 1958; Malle, 2022). Moreover, from a philosophical perspective, it often
48 matters not only *whether* something is done, but *how* it is done (Aristotle, 2009; Hursthouse &
49 Pettigrove, 2023; Stohr, 2006). An apology is not just about hearing someone say “I am sorry”,
50 but seeing genuine regret; a love letter is not just about hearing someone say “I love you”, but
51 seeing depth of emotion; and a bereavement letter is not just about hearing someone say “I am
52 sorry for your loss”, but seeing an understanding for the powerful human experience of loss.
53 There is, perhaps especially for social tasks, value not only in the outcome of doing something,

54 but the process too (Goodman, 2010). To understand any potential negative effects of outsourcing
55 to AI, we must therefore look at a broad range of non-social and social tasks, rather than draw
56 broad conclusions based on a few use cases.

57 Third, we need to know how different ways of outsourcing to AI influence negative
58 perceptions. Someone who “fully” outsources a task to AI by simply giving it a prompt and
59 copying the output word-for-word might be perceived very differently to someone who gives the
60 AI a prompt, revises the work accordingly, and finishes it themselves – using AI as a *collaborative*
61 *tool*, rather than as a replacement. Similarly, someone could deceive others about their use of AI
62 or be perfectly honest about it. While it seems reasonable to assume that “fully” outsourcing
63 would be perceived worse than using AI as a tool, and that not acknowledging AI use would be
64 perceived worse than being honest about it, it remains unclear how much this reduces negative
65 perceptions: if someone uses AI in the “best” way, by using it as a collaborative tool and being
66 open about this use, would they still suffer negative social consequences from doing so?

67 Fourth, we need to know how outsourcing to AI, in different kinds of tasks and in different
68 ways, may shape different *kinds* of social perceptions. People can judge others on separate
69 dimensions of warmth and competence (e.g. Abele et al., 2021; S. T. Fiske et al., 2007) as well as
70 on dimensions of morality and trustworthiness (Goodwin et al., 2014). It remains unclear how
71 outsourcing to AI might lead to differential character judgments across these different dimensions.

72 Fifth, we need to understand *why* outsourcing to AI, and therefore expending less effort,
73 might have these effects. Previous work has focused on how expending less effort leads to
74 negative perceptions of others (Celniker et al., 2023). But this raises the question of *why* effort is
75 seen as important and what exactly it is signalling to others, beyond one’s general cooperative
76 intent. It is possible that outsourcing leads to negative perceptions because the lack of effort spent
77 on the task signals something more fundamental about how authentic one is and how much one
78 cares about the task: when someone chooses to outsource a love letter to an AI, they might be
79 seen as valuing that love letter and what it represents less. It could be this second-step order of
80 perceptions that is the key driver of negative perceptions, especially for socio-relational tasks.

81 **Present Research**

82 In this paper, we build on classic social psychological work on character inferences from
83 reduced effort to understand how people view others who outsource different kinds of tasks, in
84 different ways, for different reasons, to AI. Across five pre-registered experiments with British
85 participants, we seek not only to understand how reduced effort through AI-outsourcing might
86 shape perceptions of others, but also to understand in more depth *why* it is that reduced effort has
87 the effect that it does.

88 In our initial pilot studies to motivate this work, we found that people who outsource a
89 range of tasks to AI or another person are perceived more negatively than people who complete
90 the tasks by themselves (see Supplementary Materials). In Study 1, we look at the effects of task
91 type, AI use, and honesty. We explore how people perceive others who outsource different kinds
92 of tasks with different levels of social relevance (e.g., from daily schedules, computer code and
93 dinner recipes to wedding vows, apology letters, or bereavement cards), manipulating whether
94 people use AI as a collaborative tool or “fully” outsource to AI and whether they are honest or
95 deceptive about their use of AI. After turning to look at perceptions of both outsourcers and the
96 outsourced work in Study 2, in Studies 3-5 we probe why outsourcing may have negative effects
97 on how we evaluate others. In Study 3, we test potential mechanisms of perceived effort and
98 authenticity by looking at how people evaluate others who either spend a lot or little time crafting
99 the AI prompts, and who either outsource to a generic or personalised AI. In Study 4, we test the
100 potential mechanism of perceived importance in the task by manipulating people’s reasons for
101 using AI – either because they wanted to save time or because they cared about the task and
102 thought that AI would improve their work. Finally, in Study 5, we bring these different potential
103 mechanisms together to explore the different pathways that influence the relationship between
104 outsourcing and negative perceptions, focusing on perceived effort, authenticity, and care in the
105 task.

Study 1**107 Methods****108 Ethical Approval**

109 Ethical approval was granted for all studies in this paper by the REDACTED Psychology
110 Research Ethics Panel. Participants in all studies provided informed consent and were debriefed
111 after the study.

112 Participants

113 We conducted a power simulation to determine our target sample size. The simulation
114 suggested that a sample size of 150 participants per condition (overall $n = 750$ for five conditions)
115 would be required to detect a small difference between conditions (Cohen's $d \approx 0.20$) with above
116 80% power.

117 We recruited a convenience sample of 800 participants from the United Kingdom through
118 the online platform Prolific (<https://www.prolific.com/>). After excluding participants who failed
119 our pre-treatment attention check, we were left with a final sample of 762 participants (438
120 female; 316 male; 4 non-binary / third gender; 4 undisclosed gender; mean age = 42.16 years).
121 78% of these participants reported having used ChatGPT before.

122 Design

123 We used a “control plus 2x2” between-subjects design. Participants were randomly
124 allocated to either the control condition, in which people in the scenarios complete the tasks
125 themselves, or one of four experimental conditions, in which people in the scenarios use AI to
126 complete the tasks. In the experimental conditions, we manipulated whether people in the
127 scenarios used AI as a collaborative tool or “fully” outsourced to AI, and whether people were
128 honest or deceptive about their use of AI. This resulted in five conditions overall: (i) the control
129 condition, (ii) the tool-honest condition, (iii) the tool-deception condition, (iv) the full-honest
130 condition, and (v) the full-deception condition.

131 **Procedure**

132 We presented participants with six scenarios. Each scenario described a person
133 completing a task, such as writing computer code or writing a love letter. The six tasks were
134 randomly drawn from a larger set of 16 tasks (see Supplementary Table 1 for the full list of tasks).

135 For each scenario, we first told participants:

- 136 • *Control condition*: “In order to complete this task, [the person] works on it by themselves
137 from start to finish.”
- 138 • *Tool outsourcing conditions*: “In order to complete this task, [the person] uses the AI tool
139 ChatGPT. They ask ChatGPT to provide ideas, inspiration, and feedback, but they edit and
140 rewrite the suggestions and finish the task themselves.”
- 141 • *Full outsourcing conditions*: “In order to complete this task, [the person] uses the AI tool
142 ChatGPT. They copy ChatGPT’s output word-for-word, rather than doing it themselves.”

143 We then told participants in the experimental conditions:

- 144 • *Honest conditions*: “After completing the task, [the person] is asked how they came up with
145 their ideas. [The person] acknowledges that they used ChatGPT as a tool / got ChatGPT to
146 do the task for them.”
- 147 • *Deception conditions*: “After completing the task, [the person] is asked how they came up
148 with their ideas. [The person] does not acknowledge that they used ChatGPT as a tool / got
149 ChatGPT to do the task for them.”

150 We then asked participants how well each of the following words described the person in

151 the scenario: competent, warm, moral, lazy, and trustworthy. Participants answered these
152 questions on 7-point Likert scales, ranging from “does not describe [the person] well” to
153 “describes [the person] extremely well”.

154 After the six scenarios, we asked participants several questions about the AI tool

155 ChatGPT, including their familiarity with ChatGPT, whether they had used ChatGPT before, how
156 frequently they used ChatGPT, and how trustworthy they thought ChatGPT was.

157 ***Pre-registration***

158 We pre-registered the study on the Open Science Framework
159 (https://osf.io/xhmzk/?view_only=a4da193574d7410ba4d2aa3945a28b05).

160 ***Statistical Analysis***

161 We fitted Bayesian multivariate multilevel cumulative-link ordinal models to the data
162 using the *brms* R package (Bürkner, 2017). We modelled each character evaluation – competence,
163 warmth, morality, laziness, and trustworthiness – as a separate response variable and included
164 fixed effects for conditions, varying intercepts for participants, and varying intercepts and slopes
165 for tasks. We used regularising priors for all parameters to impose conservatism on parameter
166 estimates. All models converged normally ($\hat{R} \leq 1.01$).

167 ***Transparency and Openness***


168 For all studies in this paper, we report how we determined our sample size, all data
169 exclusions, all manipulations, and all measures in the studies. All studies were pre-registered.
170 Analyses for all studies were conducted in R v4.4.2 (R Core Team, 2022). Visualisations were
171 produced using the *ggplot2* and *patchwork* packages (Pedersen, 2025; Wickham, 2016). The
172 manuscript was reproducibly generated using the *targets* package (Landau, 2021) and *quarto*
173 (Allaire et al., 2024). All data and code to reproduce the analyses and figures in this paper can be
174 found here: https://osf.io/xhmzk/?view_only=a4da193574d7410ba4d2aa3945a28b05

175 **Results**

176 We first looked at the overall results averaging over tasks. Across all five character
177 evaluations, we found that fully outsourcing to AI (i.e., copying the AI output verbatim) was
178 perceived more negatively than using AI as a collaborative tool (Figure 1; Table 1). By contrast,
179 we found that deception about AI usage had specific negative effects on perceptions of morality
180 and trustworthiness: people who did not acknowledge their use of AI were perceived as less moral
181 and less trustworthy. We did not find any interaction effects between full outsourcing and
182 deception.

Figure 1

Overall Character Evaluations in Study 1



Note. Participants in the control condition and four AI outsourcing conditions evaluated people in the scenarios on (a) competence, (b) warmth, (c) morality, (d) laziness, and (e) trustworthiness. Coloured points represent participant responses to the questions, jittered for easier viewing. Black points are estimated marginal means from the fitted model, pooling over participants and tasks. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.

Table 1*Overall Pairwise Contrasts in Study 1*

	Response				
	Competent	Warm	Moral	Lazy	Trustworthy
Comparison to control					
Tool Honest - Control	-0.04 [-0.25 0.19]	0.01 [-0.22 0.25]	0.53 [0.28 0.76]	0.69 [0.37 1.01]	0.67 [0.45 0.88]
Tool Deception - Control	-0.18 [-0.42 0.07]	-0.09 [-0.31 0.15]	-0.19 [-0.41 0.05]	0.88 [0.55 1.20]	-0.42 [-0.62 -0.20]
Full Honest - Control	-0.98 [-1.25 -0.71]	-0.33 [-0.59 -0.07]	-0.01 [-0.27 0.27]	1.42 [1.03 1.77]	0.35 [0.13 0.58]
Full Deception - Control	-1.05 [-1.33 -0.75]	-0.43 [-0.69 -0.16]	-0.76 [-1.04 -0.47]	1.66 [1.28 2.00]	-0.86 [-1.11 -0.59]
Effect of full outsourcing					
Full Honest - Tool Honest	-0.94 [-1.26 -0.63]	-0.34 [-0.64 -0.03]	-0.53 [-0.87 -0.20]	0.72 [0.29 1.14]	-0.32 [-0.60 -0.03]
Full Deception - Tool Deception	-0.87 [-1.18 -0.52]	-0.36 [-0.66 -0.02]	-0.57 [-0.91 -0.24]	0.78 [0.35 1.17]	-0.44 [-0.73 -0.15]
Effect of deception					
Tool Deception - Tool Honest	-0.14 [-0.42 0.13]	-0.09 [-0.39 0.20]	-0.72 [-1.01 -0.40]	0.19 [-0.19 0.58]	-1.08 [-1.35 -0.81]
Full Deception - Full Honest	-0.06 [-0.41 0.31]	-0.11 [-0.43 0.24]	-0.76 [-1.12 -0.38]	0.24 [-0.21 0.69]	-1.21 [-1.52 -0.89]
Interaction effect					
Interaction effect	0.08 [-0.36 0.53]	-0.01 [-0.45 0.43]	-0.03 [-0.51 0.42]	0.06 [-0.55 0.64]	-0.12 [-0.53 0.28]

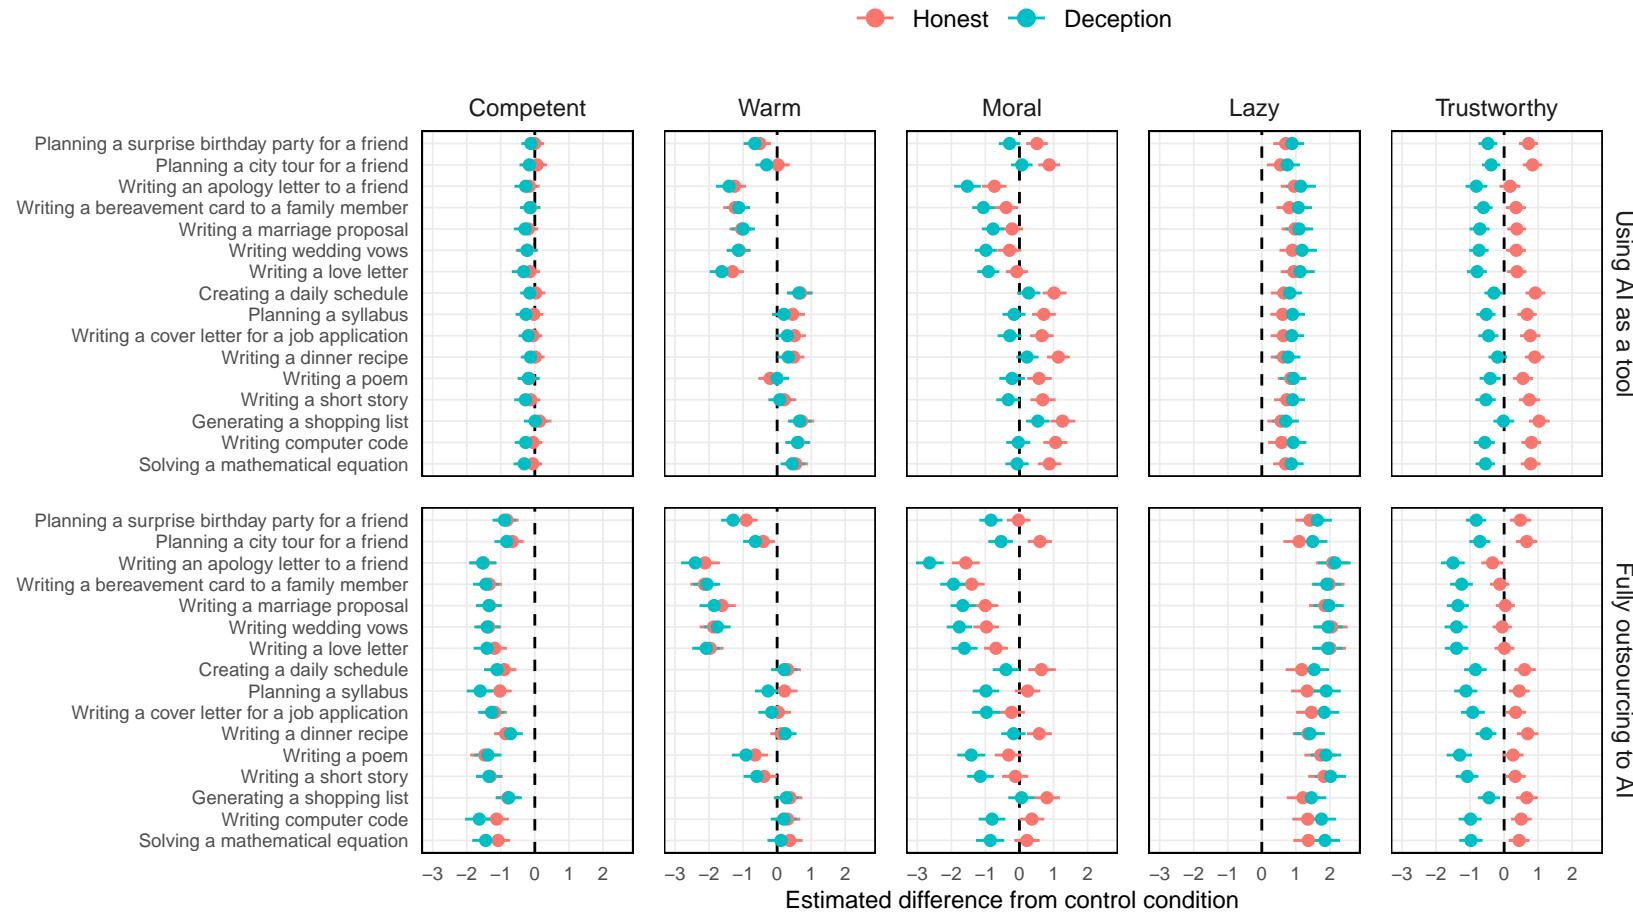
Note. Numbers reflect differences in marginal means on a 7-point Likert scale, pooling over participants and tasks. The bottom row represents the interaction between full outsourcing and deception (i.e., the difference between the differences in the rows above). Main numbers are posterior medians, numbers in the square brackets are 95% credible intervals.

The effects of outsourcing to AI varied across the different tasks, especially for

perceptions of warmth and morality (Figure 2). For example, people who used AI for social tasks, such as writing an apology letter or a bereavement card, were perceived as less warm, less moral, and lazier compared to people who completed the task themselves. This was true even if the person used AI as a tool and was honest about their usage of AI. By contrast, we observed weaker effects of outsourcing for non-social tasks like writing computer code or planning a syllabus.

Figure 2

Variation in the Effects of Outsourcing across Tasks in Study 1



Note. Tasks are ordered from most social (top) to least social (bottom) according to ratings from a pilot study. Point ranges are differences in marginal means on a 7-point Likert scale for the honest conditions (red) and deception conditions (blue) compared to the control condition. Upper panels refer to the tool outsourcing conditions, and lower panels refer to the full outsourcing conditions. Points and ranges represent posterior medians and 95% credible intervals, respectively.

189 To determine the factors that predict variation across tasks, we incorporated ratings of
190 tasks from a pilot study (see Supplementary Materials for details). Participants were asked to rate
191 the different tasks on several features: whether the task is social, requires social skills, impacts
192 others, has important consequences, and requires effort. All of these features predicted stronger
193 causal effects of outsourcing to AI compared to the control (Supplementary Figures 1 and 2). In
194 other words, outsourcing to AI is perceived more negatively for tasks that have these features,
195 compared to tasks without these features.

196 Discussion

197 In Study 1, we looked at how people who outsourced to AI in different ways were
198 perceived across a broad range of social and non-social tasks. In line with our predictions, we
199 found that “fully” outsourcing to AI was perceived more negatively than using AI as a
200 collaborative tool, particularly for socio-relational tasks. We also found, predictably, that people
201 were seen as less moral and less trustworthy if they did not acknowledge their use of AI.
202 Importantly, though, we show that even using AI in the “best” way – only as a tool and being
203 honest about one’s usage – still led to negative social perceptions for the more socio-relational
204 tasks like writing a love letter, an apology, or wedding vows.

205 In Study 2, we investigate whether these negative perceptions extend to the work itself and
206 remain after seeing the output. It could be, for example, that someone is perceived badly for using
207 ChatGPT to write their bereavement card, but the writing itself is seen as equally well-written and
208 authentic, if not more so, than if the person had written the card themselves. Indeed, evidence
209 suggests that text generated by ChatGPT is rated as higher quality than human-written text (Noy
210 & Zhang, 2023). Moreover, it is possible that seeing appropriate output could mitigate negative
211 perceptions by highlighting how the AI can in fact perform the task well. We explored these
212 possibilities in Study 2.

Study 2**214 Methods****215 Participants**

216 We conducted a power simulation to determine our target sample size. The simulation
217 suggested that a sample size of 125 participants per condition (overall $n = 750$ for six conditions)
218 would be required to detect a small-to-medium difference between conditions (Cohen's $d \approx 0.40$)
219 with above 80% power.

220 We recruited a convenience sample of 800 participants from the United Kingdom through
221 Prolific. After excluding participants who failed our pre-treatment attention check, we were left
222 with a final sample of 766 participants (425 female; 337 male; 3 non-binary / third gender; 1
223 undisclosed gender; mean age = 41.93 years). 72% of these participants reported having used
224 ChatGPT before.

225 Design

226 We randomly allocated participants to one of six conditions in a 3x2 between-subjects
227 design. We manipulated the type of outsourcing: (i) no outsourcing control, (ii) using AI as a tool,
228 and (iii) fully outsourcing to AI. Here, in contrast to Study 1, we also explicitly manipulated
229 whether the task prompt was social or non-social.

230 Procedure

231 We told participants that they would read and evaluate a short piece of writing from
232 "another participant". In reality, we generated the writing using ChatGPT version 4.0. We asked
233 ChatGPT to generate a 300 word response to the prompt and to write convincingly like a real
234 human. We then edited the text to appear more human-like by, for example, removing classic AI
235 markers like dashes and concluding sentences and ensuring that the information was not too
236 generic, such that the writing could reasonably be attributed to both a human and AI.

237 The prompt for the piece of writing varied between conditions:

- 238 • *Social conditions*: "Please write a description of a close family member or friend,

239 explaining why they are special to you.”

- 240 • *Non-social conditions*: “Please write a short description of a book, TV show, or film of
241 your choice.”

242 We explained that the “other participant” was asked several questions about how they

243 produced their answer, including whether or not they used an AI tool like ChatGPT. We explained

244 that the participant was encouraged to be honest and told that they would be paid regardless. The

245 response from the “other participant” varied between conditions:

- 246 • *Control conditions*: “The participant reported that they did not use any AI tool like
247 ChatGPT. Instead, they worked on the response themselves from start to finish.”
- 248 • *Tool outsourcing conditions*: “The participant reported using ChatGPT to provide ideas,
249 inspiration, and feedback. The participant told us that they edited and rewrote ChatGPT’s
250 suggestions and finished writing the response themselves.”
- 251 • *Full outsourcing conditions*: “The participant reported using ChatGPT to complete the task.
252 The participant told us that they copied ChatGPT’s output word-for-word, rather than
253 producing the response themselves.”

254 Next, we presented participants with a randomly-chosen pre-generated essay answer to the

255 prompt (see Supplementary Tables 2 and 3 for full essay answers). In the social conditions, the

256 answer either referred to the participant’s father, their sister, or their best friend. In the non-social

257 conditions, the answer either referred to the book The Hobbit, the TV show Buffy the Vampire

258 Slayer, or the film Titanic. Reading times and responses to a follow-up comprehension question

259 suggested that participants read the essay answers in sufficient detail (see Supplementary Table 4).

260 Finally, we asked participants about their perceptions of the essay answer and the “other

261 participant”. We asked how well-written, meaningful, and authentic they thought the answer was

262 (7-point Likert scales), what letter grade they would give the answer (A-E), and how much they

263 would hypothetically reward the other participant for their work (from £0.00 to £1.00). We also

264 asked how well each of the following words described the other participant: competent, warm,
265 moral, lazy, and trustworthy (7-point Likert scales).

266 At the end of the study, we gave participants a manipulation check and asked them
267 whether they believed the manipulation. Almost all participants correctly reported the condition
268 that they were in and most participants stated that they believed the essay response was written in
269 the way we described, suggesting that the manipulation was successful (see Supplementary
270 Table 5). We also asked participants several questions about ChatGPT.

271 ***Pre-registration***

272 We pre-registered the study on the Open Science Framework¹.

273 ***Statistical Analysis***

274 We fitted two Bayesian multilevel models to the data. The first model was a multivariate
275 cumulative-link ordinal model including all Likert scales as separate response variables. The
276 second model was a zero-one-inflated-beta model applied specifically to the reward question,
277 which was a slider scale varying between 0 and 1. For both models, we included fixed effects for
278 the interaction between outsourcing type and task type and varying intercepts and slopes for essay
279 answers. We used regularising priors for all parameters to impose conservatism on parameter
280 estimates. All models converged normally ($\hat{R} \leq 1.01$).

281 **Results**

282 We first looked at character evaluations. We found that even when provided with concrete
283 output, people were still perceived more negatively across all character evaluations if they
284 outsourced the writing task to AI, either by using ChatGPT as a collaborative tool or by copying
285 the text from ChatGPT verbatim (Supplementary Figure 3; Supplementary Table 6). In contrast to
286 Study 1, however, we did not find any differences in character evaluations between the tool

¹ Due to a technical error with archiving this pre-registration on the Open Science Framework, the timestamp for the registration was lost. However, on our OSF project

(https://osf.io/xhmzk/?view_only=a4da193574d7410ba4d2aa3945a28b05), it is possible to view our pre-registration document file and its timestamped upload date.


287 outsourcing and full outsourcing conditions. We did not find differences in character evaluations
288 between social and non-social tasks and did not find any interaction effects.

289 Turning to evaluations of the work itself, we found that the AI-outsourced work (either
290 outsourced by using AI as a collaborative tool or fully outsourced) was judged as being equally
291 well written to the work in the control condition (Figure 3; Table 2). This is in line with the
292 writing indeed being identical in all conditions. Interestingly, however, we found that essay
293 responses that were ostensibly generated using AI were perceived as less meaningful and less
294 authentic compared to essay responses ostensibly written by a human. Participants also marked
295 AI-generated essays with a lower grade and rewarded AI-generated essays with a lower
296 hypothetical monetary bonus. In contrast to Study 1, we did not find differences in perceptions of
297 the work between the tool outsourcing and full outsourcing conditions, except for the reward
298 question, where fully outsourced essays (i.e., essays copied verbatim from ChatGPT) were
299 rewarded £0.23 less than essays generated using AI as a collaborative tool. We did not find any
300 differences between social and non-social tasks and did not find any interaction effects.

301 **Discussion**

302 In Study 2, we turned to look at how people perceived both the outsourcer and the
303 outsourced work when given specific output in a social or non-social task that was described as
304 being produced independently by a person, produced by a person in collaboration with AI as a
305 tool, or outsourced in full to AI. We find that our results generalise from character judgments to
306 perceptions of the work itself: text purportedly generated using AI was perceived to be less
307 meaningful, less authentic, and less reward-worthy compared to the same text described as
308 human-generated.

309 Surprisingly, we found no differences in the effect of AI-outsourcing between social and
310 non-social tasks. This may be due to the particular tasks we chose. Writing *about* someone close
311 to you is not quite the same as writing something *for* someone close to you, as is the case with
312 wedding vows, love letters, and bereavement cards. We also found no differences between the tool
313 and full outsourcing conditions, aside from the lower rewards given to participants in the latter

Figure 3*Perceptions of the Work in Study 2*

Note. Participants in the control condition, the tool outsourcing condition, and the full outsourcing condition evaluated the essay response to the writing task. Participants rated whether the essay response was (a) well-written, (b) meaningful, and (c) authentic. Participants also (d) graded the work and (e) rewarded the work with a hypothetical monetary bonus. Jittered points represent participant responses to the questions, split by whether the writing task was a non-social task (red) or a social task (blue). Point ranges are estimated marginal means from the fitted model, pooling over essay answers. Points and line ranges represent posterior medians and 95% credible intervals, respectively.

Table 2

Pairwise Contrasts for Perceptions of the Work in Study 2

	Response				
	Well written	Meaningful	Authentic	Grade	Reward
Effect of outsourcing type					
Task type = Social					
Tool Social - Control Social	-0.42 [-0.84 0.03]	-0.64 [-1.11 -0.08]	-1.01 [-1.47 -0.35]	-0.33 [-0.60 -0.02]	-0.20 [-0.27 -0.13]
Full Social - Control Social	-0.28 [-0.63 0.09]	-0.84 [-1.55 -0.06]	-0.95 [-1.93 -0.07]	-0.55 [-1.07 -0.04]	-0.43 [-0.50 -0.36]
Full Social - Tool Social	0.14 [-0.39 0.67]	-0.21 [-1.13 0.74]	0.04 [-1.12 1.12]	-0.23 [-0.84 0.36]	-0.23 [-0.32 -0.14]
Task type = Non-social					
Tool Non-social - Control Non-social	-0.46 [-0.92 -0.02]	-0.62 [-1.09 -0.08]	-1.03 [-1.47 -0.49]	-0.42 [-0.71 -0.07]	-0.19 [-0.26 -0.12]
Full Non-social - Control Non-social	-0.32 [-0.70 0.03]	-0.90 [-1.56 -0.08]	-1.23 [-2.09 -0.12]	-0.71 [-1.16 -0.13]	-0.41 [-0.49 -0.32]
Full Non-social - Tool Non-social	0.14 [-0.39 0.71]	-0.28 [-1.06 0.62]	-0.21 [-1.19 0.97]	-0.29 [-0.86 0.34]	-0.22 [-0.31 -0.12]
Effect of task type					
Control Social - Control Non-social	-0.10 [-0.56 0.33]	0.28 [-0.36 0.73]	0.27 [-0.21 0.64]	-0.04 [-0.27 0.17]	0.01 [-0.05 0.06]
Tool Social - Tool Non-social	-0.06 [-0.67 0.56]	0.27 [-0.61 0.99]	0.29 [-0.41 0.95]	0.04 [-0.30 0.41]	0.00 [-0.09 0.09]
Full Social - Full Non-social	-0.07 [-0.64 0.51]	0.30 [-0.59 1.15]	0.48 [-0.37 1.38]	0.09 [-0.33 0.57]	-0.02 [-0.12 0.09]
Interaction effect					
Interaction: Tool - Control	0.05 [-0.39 0.51]	-0.02 [-0.52 0.56]	0.02 [-0.51 0.62]	0.09 [-0.22 0.39]	-0.01 [-0.10 0.08]
Interaction: Full - Control	0.04 [-0.34 0.46]	0.03 [-0.59 0.73]	0.23 [-0.50 1.02]	0.13 [-0.23 0.55]	-0.03 [-0.12 0.08]
Interaction: Full - Tool	0.00 [-0.58 0.55]	0.05 [-0.74 0.89]	0.21 [-0.66 1.12]	0.05 [-0.43 0.53]	-0.01 [-0.13 0.10]

Note. Numbers reflect differences in marginal means on either a 7-point Likert scale (well-written, meaningful, authentic), a 5-point ordinal grade scale (grade), or a 0-1 sliding scale (reward). Estimates are pooled over essay answers. The bottom rows represent the interactions between outsourcing type and task type. Main numbers are posterior medians, numbers in the square brackets are 95% credible intervals.

314 condition. It is possible that because the set-up described to participants was of another
 315 participant who was asked to produce work on Prolific and then admitted they used AI,
 316 participants saw any kind of AI use as violating an implicit contract between the survey requester
 317 and respondent and judged them negatively accordingly.

318 In Study 3, we turn to explore potential mechanisms driving our effects. We assume that
 319 effort may play a role, since perceived effort is often used as a signal of one's moral character

320 (Cubitt et al., 2011) and cooperative intent (Celniker et al., 2023). Study 2 also suggested a role of
321 authenticity: in line with work on the psychological importance of authenticity (Newman, 2019),
322 people who outsource to AI may be perceived as producing work that is less authentically their
323 own, leading to negative evaluations. To explore these potential mechanisms, we experimentally
324 manipulate (1) how much effort someone puts into the task and (2) whether they outsource the
325 task to a standard LLM like ChatGPT or a personalised LLM trained specifically on their own
326 prior writings (and so therefore producing work that is more authentically “theirs”). We expected
327 negative perceptions of outsourcing to be mitigated when the person uses a personalised LLM and
328 expends significant effort on formulating prompts for the AI.

329 **Study 3**

330 **Methods**

331 **Participants**

332 We used the same power estimate from Study 1 to determine our target sample size of $n =$
333 750 (150 participants in each of five conditions). We recruited a convenience sample of 802
334 participants from the United Kingdom through Prolific. After excluding participants who failed
335 our pre-treatment attention check, we were left with a final sample of 753 participants (462
336 female; 278 male; 9 non-binary / third gender; 4 undisclosed gender; mean age = 44.29 years).
337 74% of these participants reported having used ChatGPT before.

338 **Design**

339 We used the same “control plus 2x2” between-subjects design as in Study 1. In the
340 experimental conditions, we manipulated whether people in the scenarios used a standard or
341 personalised AI model, and whether people put more or less effort into the task. This resulted in
342 five conditions overall: (i) the control condition, (ii) the standard-low-effort condition, (iii) the
343 standard-high-effort condition, (iv) the personalised-low-effort condition, and (v) the
344 personalised-high-effort condition. Our authenticity manipulation was inspired by recent
345 psychological work looking at the credit-blame asymmetry in AI use (Earp et al., 2024), showing

346 that people receive more personal credit for their work when they use an AI model trained on their
347 own prior writings.

348 ***Procedure***

349 The procedure was mostly identical to Study 1 to allow us to explore effects across a range
350 of tasks, but we updated the study preamble and the presentation of the scenarios. For participants
351 in the personalised AI conditions, we expanded the study preamble to explain that personalised AI
352 models were trained on people's own prior writings and "tailored to each specific person and their
353 own thoughts, feelings, and values". Then in the scenarios, we told participants in the
354 experimental conditions:

- 355 • *Standard AI conditions*: "In order to complete this task, [the person] uses the AI tool
356 ChatGPT."
- 357 • *Personalised AI conditions*: "In order to complete this task, [the person] uses a personalised
358 AI tool."

359 We then told participants:

- 360 • *Low effort conditions*: "[The person] quickly gives the AI a rushed prompt and uses its first
361 output."
- 362 • *High effort conditions*: "[The person] carefully gives the AI several detailed prompts and,
363 after multiple rounds of changes, uses its resulting output."

364 At the end of the study, we asked participants to choose which of these was more authentic
365 and effortful, respectively. 94% of participants stated that the personalised AI was more authentic
366 and 99% of participants stated that giving the AI several detailed prompts was more effortful.
367 This suggests that even if participants might not have felt the output was meaningfully authentic in
368 the way that mattered (see Discussion), our participants agreed that using a personalised AI was at
369 least more authentic than using a generic one.

370 ***Pre-registration***

371 We pre-registered the study on the Open Science Framework

372 (https://osf.io/ac9g3/?view_only=912d9b57023d49baa87eea999574f0ce).

373 ***Statistical Analysis***

374 We fitted the same Bayesian multivariate multilevel cumulative-link ordinal models as in

375 Studies 1 and 2. All models converged normally ($\hat{R} \leq 1.01$).

376 **Results**

377 We first looked across all the tasks. On average, we found that people who outsourced to

378 AI in a low effort way were perceived as less competent, less moral, lazier, and less trustworthy

379 than people who put more effort into their use of AI (Figure 4; Table 3). By contrast, we found

380 that character evaluations did not differ between people who used a standard AI model rather than

381 a personalised AI model. We also found no interaction effects between effort and the type of AI

382 used.

383 As in Study 1, the effects of outsourcing to AI varied across the different tasks, especially

384 for perceptions of warmth and morality (Figure 5). We again found that the negative causal effects


385 of outsourcing to AI were particularly strong for tasks that are social, require social skills, impact

386 others, have important consequences, and require effort (Supplementary Figures 4 and 5). Indeed,

387 for tasks like writing wedding vows or writing a love letter, outsourcing to a personalised AI in a

388 high effort way was still perceived more negatively than the control condition for all five character

389 dimensions.

Figure 4*Overall Character Evaluations in Study 3*

Note. Participants in the control condition and four AI outsourcing conditions evaluated people in the scenarios on (a) competence, (b) warmth, (c) morality, (d) laziness, and (e) trustworthiness. Coloured points represent participant responses to the questions, jittered for easier viewing. Black points are estimated marginal means from the fitted model, pooling over participants and tasks. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.

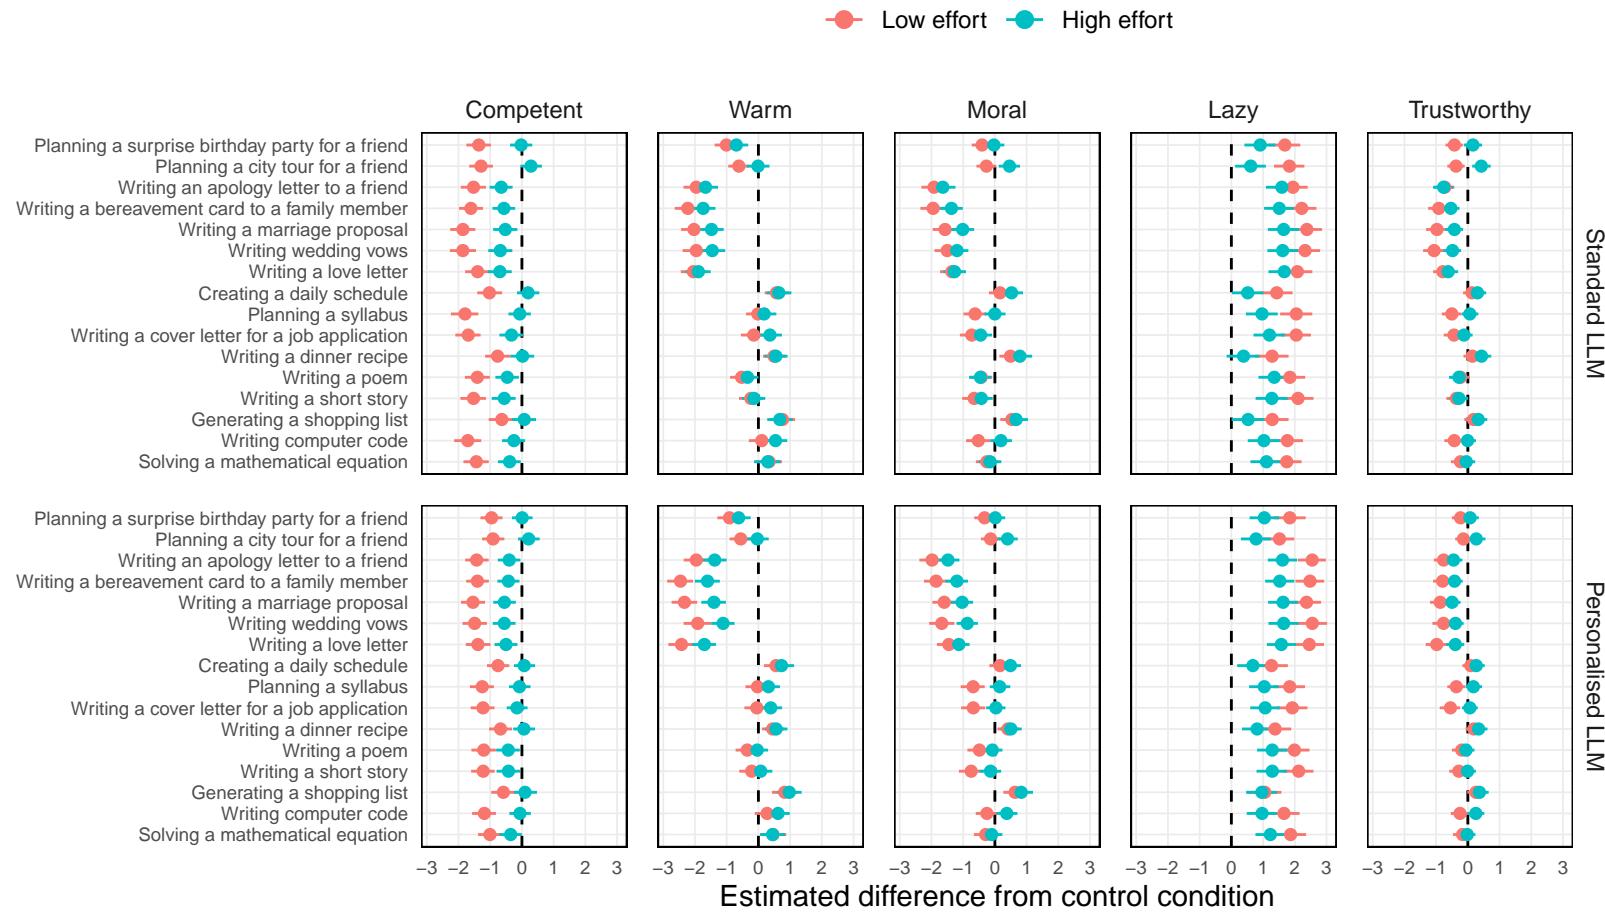
Table 3*Overall Pairwise Contrasts in Study 3*

	Response				
	Competent	Warm	Moral	Lazy	Trustworthy
Comparison to control					
Standard Low Effort - Control	-1.17 [-1.49 -0.85]	-0.26 [-0.53 0.00]	-0.34 [-0.58 -0.10]	1.67 [1.26 2.05]	-0.25 [-0.44 -0.06]
Standard High Effort - Control	-0.18 [-0.44 0.09]	-0.11 [-0.35 0.12]	-0.08 [-0.31 0.14]	0.94 [0.54 1.32]	0.00 [-0.19 0.19]
Personalised Low Effort - Control	-0.98 [-1.24 -0.70]	-0.24 [-0.50 0.02]	-0.33 [-0.56 -0.08]	1.71 [1.29 2.08]	-0.19 [-0.38 -0.01]
Personalised High Effort - Control	-0.13 [-0.38 0.13]	0.00 [-0.24 0.23]	0.00 [-0.22 0.23]	1.08 [0.69 1.44]	0.05 [-0.12 0.23]
Effect of AI type					
Standard Low Effort - Personalised Low Effort	-0.19 [-0.54 0.17]	-0.03 [-0.35 0.31]	-0.01 [-0.32 0.30]	-0.04 [-0.51 0.43]	-0.06 [-0.30 0.18]
Standard High Effort - Personalised High Effort	-0.04 [-0.36 0.27]	-0.11 [-0.41 0.20]	-0.07 [-0.36 0.22]	-0.14 [-0.61 0.34]	-0.05 [-0.28 0.17]
Effect of effort					
Standard Low Effort - Standard High Effort	-1.00 [-1.34 -0.62]	-0.15 [-0.47 0.17]	-0.26 [-0.56 0.04]	0.73 [0.25 1.21]	-0.25 [-0.50 -0.02]
Personalised Low Effort - Personalised High Effort	-0.84 [-1.16 -0.53]	-0.24 [-0.56 0.10]	-0.33 [-0.63 -0.03]	0.63 [0.15 1.08]	-0.25 [-0.48 -0.03]
Interaction effect					
Interaction effect	-0.15 [-0.61 0.32]	0.09 [-0.36 0.53]	0.06 [-0.36 0.49]	0.10 [-0.55 0.76]	0.00 [-0.32 0.32]

Note. Numbers reflect differences in marginal means on a 7-point Likert scale, pooling over participants and tasks. The bottom row represents the interaction between AI type and effort (i.e., the difference between the differences in the rows above). Main numbers are posterior medians, numbers in the square brackets are 95% credible intervals.

Figure 5

Variation in the Effects of Outsourcing across Tasks in Study 3



Note. Tasks are ordered from most social (top) to least social (bottom) according to ratings from a pilot study. Point ranges are differences in marginal means on a 7-point Likert scale for the low effort conditions (red) and high effort conditions (blue) compared to the control condition. Upper panels refer to the standard LLM conditions, and lower panels refer to the personalised LLM conditions. Points and ranges represent posterior medians and 95% credible intervals, respectively.

390 **Discussion**

391 In Study 3, we found that effort is an important mechanism by which outsourcing to AI
392 leads to negative character evaluations. People who engaged in effortless copying of the AI's first
393 output were perceived more negatively than people who spent time and effort crafting the AI's
394 output with multiple prompts. Nevertheless, for social tasks like writing wedding vows or love
395 letters, outsourcing to AI in a high effort way was still perceived more negatively than completing
396 the task oneself.

397 Interestingly, we found no effect of authenticity as proxied by the use of a personalised AI
398 that is trained on one's own prior writings compared to a standard AI like ChatGPT. This could
399 indicate that authenticity is not an important mechanism underlying the effect of outsourcing on
400 negative character evaluations. However, our specific manipulation may not have moved the
401 needle on authenticity enough to impact character evaluations. While previous work has found an
402 effect of personalised AI models on perceived credit (Earp et al., 2024), and the majority of
403 participants in our study stated that the personalised AI was more authentic than a standard model
404 like ChatGPT, it is possible that perceptions of *meaningful* authenticity in our study remained low
405 even with the personalised AI model. An AI could be perfectly trained on all apologies that a
406 person has ever written, but one might still think that a specific apology it then generates in a new
407 instance is not an *authentic* apology. Therefore, even if people were described as outsourcing to
408 an AI that was trained on their own writing and therefore personalised, participants still may not
409 have seen the specific output as being meaningfully authentic in the way that matters for character
410 judgments.

411 In Study 4, we turn to look at a third potential mechanism: a perceived lack of importance
412 attached to the task. When participants read about someone who outsources to AI in our studies,
413 they may be inferring that they simply did not care enough about the task – “If this was important
414 to them, they would do it themselves!”. To the extent that we especially want people to care about
415 their relationships with others – the kind of things demonstrated through love letters, apology
416 notes, and gift-giving – this could explain the particular negativity we see for social tasks

417 compared to tasks like writing daily schedules, recipes, or computer code. To test this, in Study 4,
418 we attempted to counteract inferences about care for the task by explicitly telling participants that
419 someone had a good reason for using AI: that they really cared about the task and used AI because
420 they wanted to get it right.

421 **Study 4**

422 **Methods**

423 **Participants**

424 We used the same power estimate from Study 1 to determine our target sample size of $n =$
425 750 (150 participants in each of five conditions). We recruited a convenience sample of 800
426 participants from the United Kingdom through Prolific. After excluding participants who failed
427 our pre-treatment attention check, we were left with a final sample of 758 participants (398
428 female; 346 male; 8 non-binary / third gender; 6 undisclosed gender; mean age = 41.72 years).
429 80% of these participants reported having used ChatGPT before.

430 **Design**

431 We used the same “control plus 2x2” between-subjects design as in Studies 1 and 3. In the
432 experimental conditions, we manipulated whether people in the scenarios used AI as a tool or
433 “fully” outsourced to AI, and whether people had bad or good reasons for using AI. This resulted
434 in five conditions overall: (i) the control condition, (ii) the tool-bad-reason condition, (iii) the
435 tool-good-reason condition, (iv) the full-bad-reason condition, and (v) the full-good-reason
436 condition.

437 **Procedure**

438 The procedure was mostly identical to Study 3, with two changes. First, we reduced the
439 number of tasks, focusing on eight tasks (four “social” tasks and four “non-social” tasks) that fit
440 with the manipulation of the updated design (since, for example, participants might find it difficult
441 to see how someone could deeply value a shopping list and want to get it right). Second, we
442 updated the presentation of the scenarios. We told participants in the experimental conditions:

- 443 • *Bad reason conditions*: “Because they are really short on time and want to complete the
444 task quickly, [the person] uses the AI tool ChatGPT.”
445 • *Good reason conditions*: “Because this task is really important to them and they want to
446 make sure they get it right, [the person] uses the AI tool ChatGPT.”

447 We then told participants:

- 448 • *Tool outsourcing conditions*: “[The person] asks ChatGPT to provide ideas, inspiration, and
449 feedback, but they edit and rewrite the suggestions and finish the task themselves.”
450 • *Full outsourcing conditions*: “[The person] copies ChatGPT’s output word-for-word, rather
451 than doing it themselves.”

452 In addition to the five character evaluations, on each page we also asked participants, on a
453 7-point Likert scale, how much they thought the person cared about the task.

454 ***Pre-registration***

455 We pre-registered the study on the Open Science Framework


456 (https://osf.io/ac9g3/?view_only=912d9b57023d49baa87eea999574f0ce).

457 ***Statistical Analysis***

458 We fitted the same Bayesian multivariate multilevel cumulative-link ordinal models as in
459 Studies 1 and 3. All models converged normally ($\hat{R} \leq 1.01$).

460 **Results**

461 We first looked across all the tasks. In line with our previous results, we found that people
462 who fully outsourced to AI by copying the output verbatim were perceived as less competent, less
463 moral, and less trustworthy than people who used AI as a collaborative tool (Figure 6; Table 4).
464 Perhaps surprisingly, though, people’s reasons for outsourcing to AI did not appear to influence
465 character evaluations when pooling across all the tasks. When looking at the tasks overall,
466 character evaluations did not differ between people who really cared about the task and wanted to
467 get it right and people who used AI because they were short on time and wanted to complete the
468 task quickly. This was true both when using the AI as a tool or outsourcing in full.

Figure 6*Overall Character Evaluations in Study 4*

Note. Participants in the control condition and four AI outsourcing conditions evaluated people in the scenarios on (a) competence, (b) warmth, (c) morality, (d) laziness, and (e) trustworthiness. Coloured points represent participant responses to the questions, jittered for easier viewing. Black points are estimated marginal means from the fitted model, pooling over participants and tasks. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.

Table 4*Overall Pairwise Contrasts in Study 4*

	Response					
	Competent	Warm	Moral	Lazy	Trustworthy	Care
Comparison to control						
Tool Bad Reason - Control	-0.27 [-0.53 -0.02]	-0.03 [-0.37 0.30]	-0.11 [-0.45 0.24]	1.16 [0.73 1.51]	-0.05 [-0.27 0.16]	-0.56 [-0.97 -0.11]
Tool Good Reason - Control	-0.12 [-0.38 0.13]	0.11 [-0.21 0.43]	0.11 [-0.19 0.39]	0.80 [0.43 1.17]	0.21 [-0.02 0.42]	-0.02 [-0.37 0.31]
Full Bad Reason - Control	-1.16 [-1.45 -0.83]	-0.22 [-0.59 0.14]	-0.39 [-0.75 0.00]	1.58 [0.75 2.01]	-0.42 [-0.68 -0.15]	-0.92 [-1.46 -0.29]
Full Good Reason - Control	-1.08 [-1.40 -0.62]	-0.19 [-0.54 0.16]	-0.40 [-0.73 -0.05]	1.50 [0.78 1.94]	-0.35 [-0.59 -0.09]	-0.50 [-0.91 -0.07]
Effect of outsourcing type						
Full Bad Reason - Tool Bad Reason	-0.89 [-1.19 -0.55]	-0.19 [-0.66 0.27]	-0.27 [-0.77 0.21]	0.42 [-0.53 0.98]	-0.37 [-0.68 -0.05]	-0.36 [-1.04 0.39]
Full Good Reason - Tool Good Reason	-0.96 [-1.32 -0.46]	-0.30 [-0.78 0.16]	-0.50 [-0.93 -0.04]	0.70 [-0.11 1.24]	-0.56 [-0.85 -0.23]	-0.48 [-0.99 0.07]
Effect of reasons for outsourcing						
Tool Bad Reason - Tool Good Reason	-0.15 [-0.44 0.14]	-0.14 [-0.59 0.28]	-0.22 [-0.63 0.23]	0.36 [-0.16 0.82]	-0.25 [-0.53 0.03]	-0.54 [-1.04 0.01]
Full Bad Reason - Full Good Reason	-0.07 [-0.60 0.31]	-0.03 [-0.54 0.45]	0.01 [-0.48 0.49]	0.08 [-0.85 0.85]	-0.07 [-0.41 0.26]	-0.42 [-1.10 0.34]
Interaction effect						
Interaction effect	0.07 [-0.50 0.55]	0.11 [-0.57 0.75]	0.22 [-0.41 0.86]	-0.28 [-1.28 0.64]	0.18 [-0.26 0.61]	0.13 [-0.73 1.01]

Note. Numbers reflect differences in marginal means on a 7-point Likert scale, pooling over participants and tasks. The bottom row represents the interaction between outsourcing type and the reasons for outsourcing (i.e., the difference between the differences in the rows above). Main numbers are posterior medians, numbers in the square brackets are 95% credible intervals.

469 Importantly, though, as in our previous studies, the type of task mattered (Figure 7).

470 Perceptions of outsourcing were particularly negative for tasks that are social, require social skills,

471 impact others, have important consequences, and require effort (Supplementary Figures 6 and 7).

472 Indeed, for socio-relational tasks like writing an apology letter and writing wedding vows, people

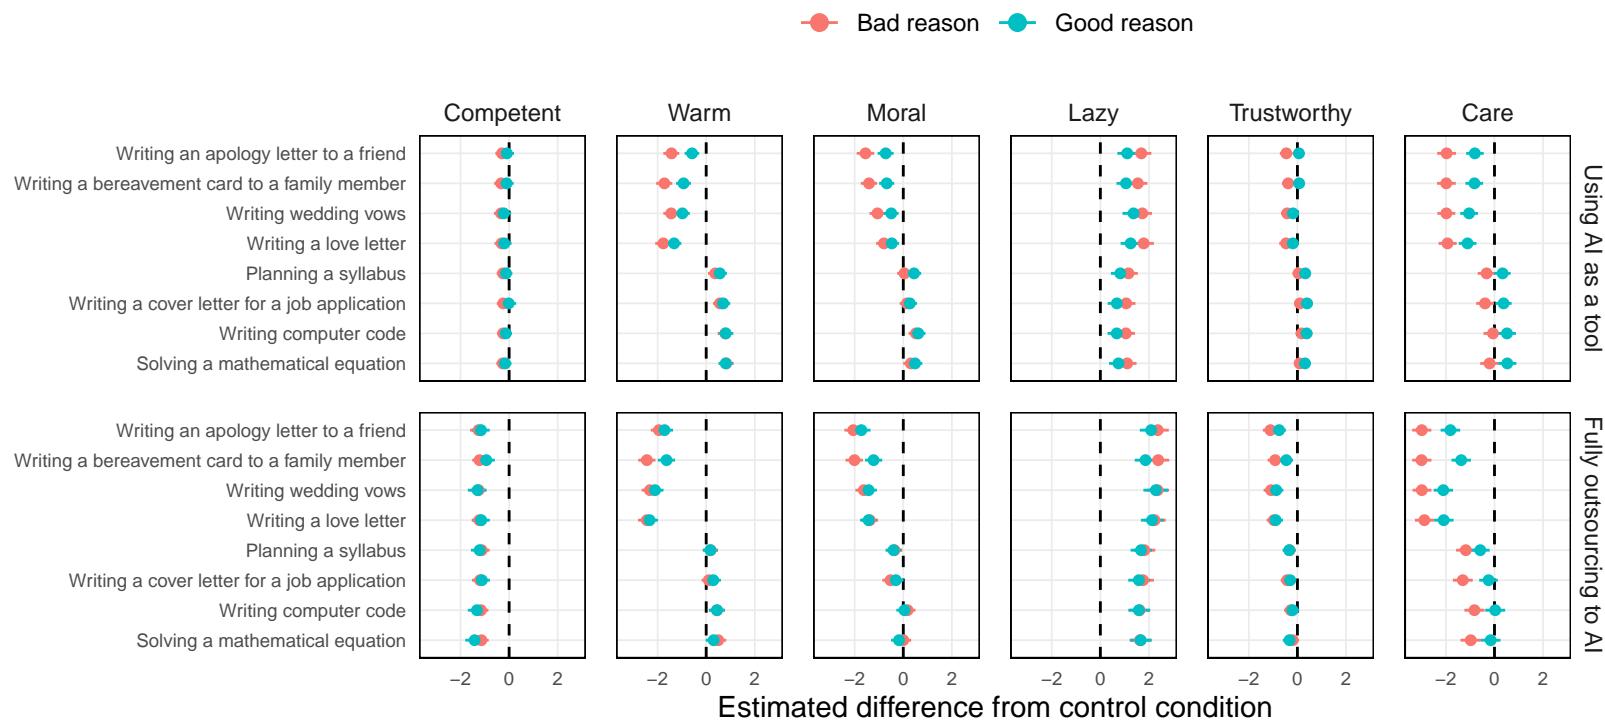
473 using AI as a tool for good reasons were still perceived more negatively than the control condition

474 on the dimensions of warmth, morality, laziness, and care, though not on the dimensions of

475 competence or trustworthiness.

Figure 7

Variation in the Effects of Outsourcing across Tasks in Study 4



Note. Tasks are ordered from most social (top) to least social (bottom) according to ratings from a pilot study. Point ranges are differences in marginal means on a 7-point Likert scale for the bad reason conditions (red) and good reason conditions (blue) compared to the control condition. Upper panels refer to the tool outsourcing conditions, and lower panels refer to the full outsourcing conditions. Points and ranges represent posterior medians and 95% credible intervals, respectively.

476 Moreover, when we delved further into the task-specific estimates, we found that the
477 reasons manipulation did indeed have an effect on character evaluations for social tasks – but not
478 non-social tasks (Supplementary Figure 8). When writing a bereavement card, for example,
479 people were perceived as less warm, less moral, lazier, and less trustworthy when they used AI to
480 save time compared to when they used it because they cared about doing the task well. The same
481 was not true for non-social tasks like writing computer code or solving a mathematical equation.

482 Discussion

483 In Study 4, we attempted to counteract the potential perception that outsourcing to AI
484 reflects caring less about the task by explicitly informing participants about the person's reason for
485 outsourcing: they outsourced to AI because they really cared about the task and wanted to get it
486 right. As well as replicating our finding that fully outsourcing to AI is perceived more negatively
487 than using AI as a tool, we also found an important effect of the reasons for outsourcing, but only
488 for socio-relational tasks. When writing a bereavement card or an apology letter, for example,
489 people were perceived more negatively if they used an AI tool to produce a quick output in a rush,
490 rather than to ensure they got it right. Nonetheless, for socio-relational tasks, the “best” use of AI
491 in this study – using AI as a tool because they cared about the task and wanted to get it right —
492 still led to targets being perceived more negatively than if they had completed the task themselves.

493 While we have so far shown varying evidence for three different mechanisms that might
494 underlie the negative perceptions of outsourcing to AI – effort, authenticity, and caring about the
495 task – it is likely that these mechanisms are related. For example, outsourcing to AI might indicate
496 a lack of effort, which then might signal a lack of authenticity and reduced care in the task,
497 leading to negative character evaluations. Our previous studies have been unable to test causal
498 models like these as we manipulated the mechanisms separately and independently. In Study 5,
499 therefore, we bring all three mechanisms together and test their combined associations with
500 character evaluations. To do this, we focus on a single socio-relational task — writing a love letter
501 — which we elaborate for participants with a more detailed vignette.

Study 5**503 Methods****504 Participants**

505 We conducted a power simulation to determine our target sample size. The simulation
506 suggested that a sample size of 200 participants per condition (overall $n = 600$ for three
507 conditions) would be required to detect a small-to-medium difference between conditions
508 (Cohen's $d \approx 0.30$) with above 80% power.

509 We recruited a convenience sample of 651 participants from the United Kingdom through
510 Prolific. After excluding participants who failed our pre-treatment attention check, we were left
511 with a final sample of 610 participants (371 female; 233 male; 4 non-binary / third gender; 2
512 undisclosed gender; mean age = 42.85 years). 82% of these participants reported having used AI
513 tools like ChatGPT before.

514 Design

515 We randomly allocated participants into one of three conditions in a between-subjects
516 design: (i) the control condition, (ii) the tool outsourcing condition, or (iii) the full outsourcing
517 condition. These conditions determined how the scenario was presented to participants.

518 Procedure

519 We presented participants with a vignette about a person, Adam, who is writing a love
520 letter in a Valentine's Day card to his partner (see Supplementary Materials for full vignette
521 wording). We told participants in each of the conditions:

- 522 • *Control condition*: “Adam decides to write the love letter in the card by himself.”
- 523 • *Tool outsourcing condition*: “Adam decides to use AI to help write the love letter in the
524 card. He asks ChatGPT to provide ideas, inspiration, and feedback, but he edits and
525 rewrites the suggestions and finishes writing the love letter himself.”
- 526 • *Full outsourcing condition*: “Adam decides to use AI to write the love letter in the card. He
527 asks ChatGPT to write the love letter and copies the output word-for-word, rather than

528 writing it himself.”

529 We then presented participants with the love letter that Adam wrote (in reality, this was
530 written by ChatGPT version 4o; see Supplementary Materials for wording). On the following
531 page, we asked participants what Adam wrote and whether he used AI to help. 95% of
532 participants answered both of these comprehension questions correctly.

533 Using 7-point Likert scales, we then asked participants how much effort they thought
534 Adam put into the love letter, how authentic they thought the love letter was, how much they
535 thought Adam cared about the love letter, and the same five character evaluations as in our
536 previous studies. In additional free response questions, we asked participants to explain how they
537 felt towards Adam and how they would feel if they were Adam’s partner. Finally, we asked
538 participants several questions about AI tools like ChatGPT.

539 ***Pre-registration***

540 We pre-registered the study on the Open Science Framework
541 (https://osf.io/ac9g3/?view_only=912d9b57023d49baa87eea999574f0ce).

542 ***Statistical Analysis***

543 We fitted two Bayesian regression models to the data. The first model was a multivariate
544 cumulative-link ordinal model including all Likert scales as separate response variables. The
545 second model was a path model capturing the effect of outsourcing on character evaluations, both
546 directly and indirectly through perceptions of effort, authenticity, and care. In this second model,
547 we included ordinal predictors as monotonic effects and modelled the five character evaluations as
548 a single latent variable. We used regularising priors for all parameters to impose conservatism on
549 parameter estimates. All models converged normally ($\hat{R} \leq 1.01$).


550 **Results**

551 Across all measures, we found that outsourcing the love letter to AI was perceived more
552 negatively compared to the control condition and that fully outsourcing to AI was perceived more
553 negatively than using AI as a collaborative tool (Figure 8; Table 5). Not only did outsourcing the

- 554 love letter lead to more negative character evaluations, but outsourcing to AI was also seen as less
 555 effortful, less authentic, and indicative of caring less about the task.

Figure 8

Perceptions of the Person and the Love Letter in Study 5



Note. Participants in the control, tool outsourcing, and full outsourcing conditions rated (a) the amount of effort put into the love letter, (b) how authentic the love letter was, (c) how much the person cared about the love letter, and (d-h) five character evaluation measures. Coloured points represent participant responses to the questions, jittered for easier viewing. Black points are estimated marginal means from the fitted model. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.

Table 5*Pairwise Contrasts in Study 5*

	Response							
	Effort	Authentic	Care	Competent	Warm	Moral	Lazy	Trustworthy
Tool - Control	-1.91 [-2.18 -1.64]	-1.94 [-2.23 -1.61]	-1.27 [-1.54 -1.00]	-1.15 [-1.40 -0.89]	-1.20 [-1.45 -0.95]	-1.14 [-1.39 -0.88]	1.77 [1.47 2.07]	-0.86 [-1.11 -0.62]
Full - Control	-3.76 [-3.98 -3.52]	-3.50 [-3.77 -3.22]	-2.90 [-3.16 -2.63]	-2.41 [-2.67 -2.14]	-2.62 [-2.88 -2.37]	-2.44 [-2.71 -2.17]	2.78 [2.47 3.09]	-2.07 [-2.30 -1.82]
Full - Tool	-1.85 [-2.11 -1.58]	-1.56 [-1.87 -1.26]	-1.63 [-1.93 -1.31]	-1.26 [-1.55 -0.97]	-1.42 [-1.72 -1.12]	-1.31 [-1.60 -1.02]	1.01 [0.65 1.35]	-1.20 [-1.45 -0.95]

Note. Numbers reflect differences in marginal means on a 7-point Likert scale. Main numbers are posterior medians, numbers in the square brackets are 95% credible intervals.

556 Exploratory text analysis of participants' free responses supported this quantitative pattern

557 (see Supplementary Materials for methodology and Supplementary Table 7 for results). When

558 comparing word frequencies between conditions, we found that Adam was more likely to be

559 described as "lazy" and less likely to be described as "caring", "thoughtful", and "genuine" in

560 both outsourcing conditions compared to the control condition. Adam was also more likely to be

561 described as "romantic" and as someone who "loves" his partner when he used AI as a

562 collaborative tool, compared to when he fully outsourced the love letter to AI.

563 When we included all the variables in a single path model, we found that outsourcing

564 influenced character evaluations both directly and indirectly through our proposed mechanisms

565 (Figure 9). The indirect effects showed that people perceived outsourced work as less effortful,

566 and less effortful work was seen as less authentic and indicating less care about the task. In turn,

567 less authenticity and care were associated with more negative evaluations of the person. Effort

568 itself was not directly related to character evaluations, suggesting that effort works solely through

569 perceptions of authenticity and care.

570 General Discussion

571 The release of openly available generative AI LLMs has changed lives, promising to let

572 people do more tasks, more efficiently, and perhaps to do so better than they could alone. People

573 can — and *do* — use AI tools like ChatGPT to, for example, create dinner recipes, assist with

574 coding, and even write job applications ([Department for Science, Innovation & Technology](#),

575 [2024](#)). But it is not only such routine, everyday, and non-social tasks that AI now "assists" with.

576 People can use AI for a seemingly endless range of social tasks too, from crafting apology letters

577 to writing condolences to even writing wedding vows. In this paper, across five pre-registered


578 experiments, we show how — and why — AI-outsourcing shapes perceptions of others in a world

579 where outsourcing has never been easier and cheaper.

580 In Study 1, we showed that people who outsourced tasks to AI were perceived more

581 negatively than people who completed the tasks by themselves. These negative impressions were

582 particularly strong for people who used AI to complete socio-relational tasks, such as writing a

Figure 9*Path Model in Study 5*

Note. All predictors were modelled as monotonic effects, such that parameters can be interpreted as the expected average difference between two adjacent categories of the ordinal predictor on the logit scale. The “evaluations” outcome variable was modelled as a single latent variable with loadings from all five character evaluations (competence, warmth, morality, laziness, and trustworthiness).

583 love letter or writing wedding vows, and for people who copied the model’s first output verbatim
 584 without acknowledging their reliance on AI. Moreover, these negative perceptions were found
 585 even for the “best case” of openly acknowledging the use of AI as a collaborative tool. In Study 2,
 586 we showed that people perceive both the outsourcer and the outworked work more negatively,
 587 with outsourced work perceived as less meaningful, less authentic, and less reward-worthy than
 588 ostensibly human-generated writing. In Study 3, we showed that while it matters whether people
 589 spent time crafting the AI prompts or simply gave a rushed initial prompt, even expending effort
 590 into crafting the best prompts was still not enough to counteract the negative effects from using

591 AI. In Study 4, we explored the potential role of inferred importance and found that while
592 explicitly telling participants that the person used AI because they cared about the task reduced
593 negative perceptions for social tasks, it was still not enough to eliminate negative perceptions
594 completely. In Study 5, we showed that a perceived lack of effort is taken to signal both a lack of
595 authenticity and lack of importance attached to the task, and these independently influenced
596 character judgments above and beyond the effect of effort.

597 Our findings extend work on the moralisation of effort. Studies have shown that people
598 inherently value effort and perceive displays of effort as costly signals of one's moral character
599 and cooperative intent ([Celniker et al., 2023](#)). And yet it has remained unclear how we might view
600 others who outsource to AI; how these effects might vary based on how socio-relational the task
601 is; how different ways of outsourcing influence perceptions; how outsourcing has different effects
602 on different kinds of social perceptions; and why exactly effort has the effects that it does. Across
603 our studies, we provide new insight into all of these questions. In line with previous work on the
604 importance of effort, we show that people negatively judge those who outsource to AI. We show
605 that the type of task does matter, whereby outsourcing to AI for socio-relational tasks leads to
606 particularly negative perceptions. We show that different ways of outsourcing lead to differences
607 in the degree of negative perceptions but that, critically, even outsourcing to AI in the “best” way
608 (e.g., using it as a tool and finishing the work oneself while being honest about the AI use) is still
609 not enough to eliminate the negative consequences. We show that negative perceptions from
610 outsourcing tended to go together, even if outsourcing on social tasks led to particularly negative
611 effects on warmth and morality traits. And finally, we provide further insight into why effort
612 matters. The reduced effort from outsourcing socio-relational tasks to AI signals that the work is
613 less authentically one's own and that the person cares less about the task (and therefore, perhaps,
614 the relationship). The lack of a direct effect of perceived effort in our path model showed that it is
615 inferences of authenticity and care, rather than perceived effort per se, that are associated with
616 negative character evaluations. As a participant in our final study put it: “*If he really cared, he*
617 *would have just done it by himself from scratch*” (female, 25 years old).

618 Our findings cohere with the philosophical idea that there is value in *how* a task was done,

619 and not merely *whether* it was done (Aristotle, 2009; Goodman, 2010; Hursthouse & Pettigrove,

620 2023; Stohr, 2006). For many socio-relational tasks, it might seem that part of the constitutive

621 action is the *process* by which it occurs: an apology that does not contain a genuine reflection and

622 commitment to do better, rather than just the words “I am sorry”, might not seem to be an apology

623 at all. In contrast, for many of the non-social tasks, it is easier to distinguish the importance of the

624 process from the outcome. In this way, our work suggests that people rarely adopt a purely

625 utilitarian perspective in which outcomes are the sole determinant (Everett & Kahane, 2020;

626 Kahane et al., 2018). Instead, their judgments cohere more with ideas from virtue ethics about the

627 importance of *doing* (Hursthouse & Pettigrove, 2023; Stohr, 2006). Outsourcing to AI –

628 especially for social tasks — may allow us to produce similar outputs, but by severing the

629 outcome from the practice of doing, it may risk the development and maintenance of our human

630 virtues (Vallor, 2015, 2024).

631 AI is often being marketed as being able to help us to do more and more tasks, promising

632 gains of efficiency that align with societal incentives for “hacks” that encourage people to do

633 more-and-more with less energy and effort. Our work, however, highlights that when it comes to

634 our psychology, efficiency is not the only currency. Instead, *inefficiency* can sometimes pay off

635 more, especially for social tasks. By expending effort themselves instead of outsourcing to AI,

636 people are able to signal authenticity and care for the task, and this can lead to better reputations

637 (see also Celniker et al., 2023). Correspondingly, expending effort, even “unnecessarily”, is not as

638 irrational, biased, or suboptimal as we might think from a utilitarian perspective in which

639 outcomes are the only things that matter. Instead, it is precisely this inefficiency that helps people

640 signal things that they care about and connect with others, thereby arguably reflecting a deeply

641 rational reflection of virtues and the importance of social ties (Everett et al., 2016).

642 Most speculatively, our results on the negative effects of AI-outsourcing on character

643 judgments highlight potential risks in how increased use of AI could lead to negative

644 consequences for social ties, especially if people start to assume, by default, that others are using

645 AI for the kind of tasks that matter. Sociologists have highlighted concerns about the negative
646 effects that outsourcing to AI can have on our “connective labour”, arguing that while AI can
647 enhance certain tasks, it cannot replicate the depth of human relationships essential for effective
648 caregiving, education, and support (Pugh, 2024). Similar arguments have been made about the
649 risks of outsourcing empathy to AI (Landes & Everett, 2025). In this way, the rapid move towards
650 using AI for more and more tasks could have serious and unintended consequences on the way we
651 connect with one another, serving to further weaken the social ties that bind us into a community.

652 **Limitations and Directions for Future Research**

653 The studies in this paper are not without their limitations. While we included a range of
654 different socio-relational and professional tasks in an effort to improve the generalisability of our
655 findings across domains, it would be interesting for future work to additionally explore the
656 generalisability and variability of our findings across countries with different AI infrastructures
657 and readiness levels (Oxford Insights, 2024; Tortoise Media, 2024) and over time as AI use
658 becomes more commonplace. By focusing on generalisability across various real-world tasks in
659 which people outsource, it could also be argued that our design lacks the richness of information
660 in extended vignettes that might influence character evaluations. While we have advanced
661 previous research in highlighting the ways in which effort influences perceptions of authenticity
662 and care, it will be interesting for future research to delve deeper into these mechanisms, both
663 philosophically and psychologically: *why* is it that the perceived care for the task matters, and
664 what are the boundary conditions of these effects? Finally, while we have demonstrated negative
665 perceptions of outsourcing in this paper, it will be important for future research to explore when
666 people might deem outsourcing to AI as acceptable or even preferable. Several of the participants
667 in our final study expressed in their free responses that they would have been okay with Adam
668 using AI to write the love letter if he was not a confident writer or had a learning difficulty that
669 made writing challenging, such as dyslexia. In line with this, some research has found that people
670 are more accepting of cognition-enhancing technologies and drugs when they are used to repair
671 cognitive functions, rather than to enhance cognitive functions beyond “normal” levels (Medaglia

672 et al., 2019; Rudski, 2014). Future research should explore whether negative perceptions of
673 outsourcing persist when AI is used in a reparative way.

674 **Conclusions**

675 To conclude, across five pre-registered studies, we have demonstrated negative perceptions
676 of outsourcing to AI. Our participants perceived individuals who outsource tasks to AI more
677 negatively across a range of character dimensions and perceived outsourced work as less
678 meaningful and authentic. Negative perceptions were particularly strong for socio-relational tasks,
679 such as writing wedding vows, and were compounded when the outsourcer copied the AI's output
680 verbatim and did not honestly acknowledge their use of AI. These findings connect with broader
681 debates about the importance of *doing* in social relationships, and highlight that for many tasks –
682 especially those that are more socio-relational – it might be better to move away from a focus on
683 making things more efficient at all costs and instead bring back a recognition of the power of
684 inefficiency. Doing something oneself, even if AI could do it quicker and easier, signals one that is
685 authentic and cares about the task and therefore can help bind us together. In a world of
686 algorithm-mediated interactions, AI is no substitute for investing effort into our interpersonal
687 relationships.

Acknowledgements

This work was generously supported by funding from REDACTED.

Data and Code Availability

All data and original code can be found here:

https://osf.io/ac9g3/?view_only=912d9b57023d49baa87eea999574f0ce.

Statement of Interests

The authors have no conflicts of interest to disclose.

References

- Abele, A. E., Ellemers, N., Fiske, S. T., Koch, A., & Yzerbyt, V. (2021). Navigating the social world: Toward an integrated framework for evaluating self, individuals, and groups. *Psychological Review*, 128(2), 290–314. <https://doi.org/10.1037/rev0000262>
- Allaire, J. J., Teague, C., Scheidegger, C., Xie, Y., Dervieux, C., & Woodhull, G. (2024). *Quarto* (Version 1.7.23) [Computer software]. <https://doi.org/10.5281/zenodo.5960048>
- Aristotle. (2009). *The Nicomachean ethics* (D. Ross & L. Brown, Translators). Oxford University Press.
- Bürkner, P.-C. (2017). brms: An R package for Bayesian multilevel models using Stan. *Journal of Statistical Software*, 80(1), 1–28. <https://doi.org/10.18637/jss.v080.i01>
- Celniker, J. B., Gregory, A., Koo, H. J., Piff, P. K., Ditto, P. H., & Shariff, A. F. (2023). The moralization of effort. *Journal of Experimental Psychology: General*, 152(1), 60–79. <https://doi.org/10.1037/xge0001259>
- Cubitt, R. P., Drouvelis, M., Gächter, S., & Kabalin, R. (2011). Moral judgments in social dilemmas: How bad is free riding? *Journal of Public Economics*, 95(3), 253–264. <https://doi.org/10.1016/j.jpubeco.2010.10.011>
- Department for Science, Innovation & Technology. (2024). *Public attitudes to data and AI: Tracker survey (Wave 4) report*. <https://www.gov.uk/government/publications/public-attitudes-to-data-and-ai-tracker-survey-wave-4/public-attitudes-to-data-and-ai-tracker-survey-wave-4-report>
- Earp, B. D., Porsdam Mann, S., Liu, P., Hannikainen, I., Khan, M. A., Chu, Y., & Savulescu, J. (2024). Credit and blame for AI-generated content: Effects of personalization in four countries. *Annals of the New York Academy of Sciences*, 1542(1), 51–57. <https://doi.org/10.1111/nyas.15258>
- Everett, J. A., & Kahane, G. (2020). Switching tracks? Towards a multidimensional model of utilitarian psychology. *Trends in Cognitive Sciences*, 24(2), 124–134. <https://doi.org/10.1016/j.tics.2019.11.012>

- Everett, J. A., Pizarro, D. A., & Crockett, M. J. (2016). Inference of trustworthiness from intuitive moral judgments. *Journal of Experimental Psychology: General*, 145(6), 772–787.
<https://doi.org/10.1037/xge0000165>
- Fiske, A. P. (1992). The four elementary forms of sociality: Framework for a unified theory of social relations. *Psychological Review*, 99(4), 689–723.
<https://doi.org/10.1037/0033-295X.99.4.689>
- Fiske, S. T., Cuddy, A. J., & Glick, P. (2007). Universal dimensions of social cognition: Warmth and competence. *Trends in Cognitive Sciences*, 11(2), 77–83.
<https://doi.org/10.1016/j.tics.2006.11.005>
- Goodman, R. (2010). Cognitive enhancement, cheating, and accomplishment. *Kennedy Institute of Ethics Journal*, 20(2), 145–160. <https://doi.org/10.1353/ken.0.0309>
- Goodwin, G. P., Piazza, J., & Rozin, P. (2014). Moral character predominates in person perception and evaluation. *Journal of Personality and Social Psychology*, 106(1), 148–168.
<https://doi.org/10.1037/a0034726>
- Heider, F. (1958). *The psychology of interpersonal relations*. John Wiley & Sons Inc.
<https://doi.org/10.1037/10628-000>
- Hursthouse, R., & Pettigrove, G. (2023). Virtue ethics. In E. N. Zalta & U. Nodelman (Eds.), *The Stanford encyclopedia of philosophy*.
<https://plato.stanford.edu/archives/fall2023/entries/ethics-virtue/>
- Kahane, G., Everett, J. A., Earp, B. D., Caviola, L., Faber, N. S., Crockett, M. J., & Savulescu, J. (2018). Beyond sacrificial harm: A two-dimensional model of utilitarian psychology. *Psychological Review*, 125(2), 131–164. <https://doi.org/10.1037/rev0000093>
- Kerr, N. L. (1983). Motivation losses in small groups: A social dilemma analysis. *Journal of Personality and Social Psychology*, 45(4), 819–828.
<https://doi.org/10.1037/0022-3514.45.4.819>
- Kruger, J., Wirtz, D., Van Boven, L., & Altermatt, T. W. (2004). The effort heuristic. *Journal of Experimental Social Psychology*, 40(1), 91–98.

- [https://doi.org/10.1016/S0022-1031\(03\)00065-9](https://doi.org/10.1016/S0022-1031(03)00065-9)
- Landau, W. M. (2021). The targets R package: A dynamic Make-like function-oriented pipeline toolkit for reproducibility and high-performance computing. *Journal of Open Source Software*, 6(57), 2959. <https://doi.org/10.21105/joss.02959>
- Landes, E., & Everett, J. A. C. (2025). *AI should develop human empathy, not replace it.* PsyArXiv. https://doi.org/10.31234/osf.io/y3qzu_v1
- Liu, B., Kang, J., & Wei, L. (2024). Artificial intelligence and perceived effort in relationship maintenance: Effects on relationship satisfaction and uncertainty. *Journal of Social and Personal Relationships*, 41(5), 1232–1252. <https://doi.org/10.1177/02654075231189899>
- Malle, B. F. (2022). Attribution theories. In *Theories in social psychology, second edition* (pp. 93–120). John Wiley & Sons, Ltd. <https://doi.org/10.1002/9781394266616.ch4>
- Medaglia, J. D., Yaden, D. B., Helion, C., & Haslam, M. (2019). Moral attitudes and willingness to enhance and repair cognition with brain stimulation. *Brain Stimulation*, 12(1), 44–53. <https://doi.org/10.1016/j.brs.2018.09.014>
- miramar0. (2024). *My (30F) husband (36M) used AI to write his vows. How do I approach him about this?* [Online forum post]. Reddit. https://www.reddit.com/r/relationship_advice/comments/1d9tw9j/my_30f_husband_36m_used_ai_to_write_his_vows_how/
- Newman, G. E. (2019). The psychology of authenticity. *Review of General Psychology*, 23(1), 8–18. <https://doi.org/10.1037/gpr0000158>
- Noy, S., & Zhang, W. (2023). Experimental evidence on the productivity effects of generative artificial intelligence. *Science*, 381(6654), 187–192. <https://doi.org/10.1126/science.adh2586>
- Oxford Insights. (2024). *Government AI readiness index.* <https://oxfordinsights.com/ai-readiness/ai-readiness-index/>
- Pedersen, T. L. (2025). *patchwork: The composer of plots.* <https://patchwork.data-imaginist.com>
- Pugh, A. J. (2024). *The last human job: The work of connecting in a disconnected world.* Princeton University Press.
- R Core Team. (2022). *R: A language and environment for statistical computing.* R Foundation for

- Statistical Computing. <https://www.R-project.org/>
- Roth, L. H. O., & Tissot, T. (2025). *A costly co-pilot? AI usage signals reduced morality via effort perceptions*. OSF Preprints. https://doi.org/10.31219/osf.io/vwxgz_v1
- Rudski, J. M. (2014). A comparison of attitudes toward pharmacological treatment versus enhancement under competitive and noncompetitive conditions. *AJOB Empirical Bioethics*, 5(2), 80–90. <https://doi.org/10.1080/23294515.2013.850122>
- Stohr, K. (2006). Contemporary virtue ethics. *Philosophy Compass*, 1(1), 22–27.
<https://doi.org/10.1111/j.1747-9991.2006.00004.x>
- Tortoise Media. (2024). *The global AI index*. <https://www.tortoisemedia.com/data/global-ai>
- Vallor, S. (2015). Moral deskilling and upskilling in a new machine age: Reflections on the ambiguous future of character. *Philosophy & Technology*, 28, 107–124.
<https://doi.org/10.1007/s13347-014-0156-9>
- Vallor, S. (2024). *The AI mirror: How to reclaim our humanity in an age of machine thinking*. Oxford University Press. <https://doi.org/10.1093/oso/9780197759066.001.0001>
- Wickham, H. (2016). *ggplot2: Elegant graphics for data analysis*. Springer-Verlag New York.
<https://ggplot2.tidyverse.org>

Supplementary Materials

Negative Perceptions of Outsourcing to Artificial Intelligence

AUTHORS REDACTED

AFFILIATION REDACTED

Table of contents

Pilot Study 1	2
Pilot Study 2	4
Vignette Wording in Study 5	7
Methods for Text Analysis in Study 5	8
Supplementary Figures	10
Supplementary Tables	29
Supplementary References	37

Pilot Study 1

Methods

Participants

We recruited a convenience sample of 200 participants from the United Kingdom through Prolific. After excluding participants who failed our pre-treatment attention check, we were left with a final sample of 186 participants (118 female; 67 male; 1 non-binary / third gender; 0 undisclosed gender; mean age = 38.99 years).

Procedure

We presented participants with six different tasks “that people might perform in their daily lives”. The six tasks were randomly drawn from a larger set of 20 tasks (see Supplementary Table 1 for the full list of tasks). For each task, we asked participants the following questions on 7-point Likert scales:

- Is this a social task?
- Does this task require social skills?
- Does this task impact other people?
- How important are the consequences of this task?
- How important is it that effort goes into this task?
- How important is it that others see the effort that goes into this task?

Statistical Analysis

We fitted a Bayesian multivariate multilevel cumulative-link ordinal model to the data using the *brms* R package. We modelled each task evaluation as a separate response variable and included correlated varying intercepts for participants and tasks. We used regularising priors for all parameters to impose conservatism on parameter estimates (see Supplementary Materials for full model specification). The model converged normally ($\hat{R} \leq 1.01$).

Results

We found that participants' responses to all six questions tended to be positively correlated. For example, tasks rated as more social were also rated as requiring more social skills (see Supplementary Figure 9). Estimated averages and rankings for the 20 tasks across each of the questions can be found in Supplementary Figures 10 – 15.

Pilot Study 2

Methods

Participants

We conducted a power simulation to determine our target sample size. The simulation suggested that a sample size of 150 participants per condition (overall $n = 450$ for three conditions) would be required to detect a small difference between conditions (Cohen's $d \approx 0.20$) with above 80% power.

We recruited a convenience sample of 500 participants from the United Kingdom through Prolific. After excluding participants who failed our pre-treatment attention check, we were left with a final sample of 466 participants (292 female; 169 male; 4 non-binary / third gender; 1 undisclosed gender; mean age = 42.32 years). 73% of these participants reported having used ChatGPT before.

Design

We randomly allocated participants into one of three conditions in a between-subjects design: (i) the control condition, (ii) the AI outsourcing condition, or (iii) the human outsourcing condition. These conditions determined how scenarios were presented to participants.

Procedure

We presented participants with six scenarios. Each scenario described a person completing a task, such as writing computer code or writing a love letter. The six tasks were randomly drawn from a larger set of 20 tasks (see Supplementary Table 1 for the full list of tasks). For each scenario, we told participants:

- *Control condition*: “In order to complete this task, [the person] works on it by themselves from start to finish.”
- *AI outsourcing condition*: “In order to complete this task, [the person] gets the AI tool ChatGPT to do it for them.”

- *Human outsourcing condition:* “In order to complete this task, [the person] gets someone else to do it for them.”

We then asked participants how well each of the following words described the person in the scenario: competent, warm, moral, lazy, and trustworthy. Participants answered these questions on 7-point Likert scales, ranging from “does not describe [the person] well” to “describes [the person] extremely well”.

After the six scenarios, we asked participants several questions about the AI tool ChatGPT, including their familiarity with ChatGPT, whether they had used ChatGPT before, how frequently they used ChatGPT, and how trustworthy they thought ChatGPT was.

Pre-registration

We pre-registered the study on the Open Science Framework (https://osf.io/xhmzk/?view_only=a4da193574d7410ba4d2aa3945a28b05).

Statistical Analysis

We fitted Bayesian multivariate multilevel cumulative-link ordinal models to the data using the *brms* R package. We modelled each character evaluation – competence, warmth, morality, laziness, and trustworthiness – as a separate response variable and included fixed effects for conditions, varying intercepts for participants, and varying intercepts and slopes for tasks. We used regularising priors for all parameters to impose conservatism on parameter estimates (see Supplementary Materials for full model specifications). All models converged normally ($\hat{R} \leq 1.01$).

Results

We found that people who outsourced tasks to AI or other humans were perceived more negatively than people who completed the tasks themselves (Supplementary Figure 16). In particular, people who outsourced were perceived as lazier and less competent, with smaller yet detectable differences for perceptions of warmth, morality, and trustworthiness (Supplementary Table 8). Across all measures, outsourcing to other humans was perceived more negatively than

outsourcing to AI.

We found that the effects of outsourcing varied across the different tasks, especially for perceptions of warmth and morality (Supplementary Figure 17). For example, people were perceived as less warm if they outsourced writing a love letter, but not if they outsourced writing computer code. Similarly, people were perceived as less moral if they outsourced writing an apology letter to a friend, but not if they outsourced writing a dinner recipe. By contrast, the effects of outsourcing on competence, laziness, and trustworthiness were more consistent across tasks.

To determine the factors that predict variation across tasks, we incorporated ratings of tasks from the first pilot study. Participants were asked to rate the 20 tasks on several features: whether the task is social, requires social skills, impacts others, has important consequences, and requires effort. All of these features predicted stronger causal effects of outsourcing compared to control (Supplementary Figures 18 and 19). In other words, outsourcing to AI or other humans is perceived more negatively for tasks that have these features, compared to tasks without these features.

Vignette Wording in Study 5

We presented participants in Study 5 with the following vignette text:

Adam has been dating his partner for almost a year, and Valentine's Day is coming up. He knows that many people exchange a card on Valentine's Day containing a love letter to their partner, and he decides to send a love letter to his partner too.

This year, Adam has been closely following developments in technology and has read of people using AI tools like ChatGPT for things like this, either using it to help with writing or getting AI to do the task completely.

This was followed by the manipulation text (see main text). Participants were then presented with the love letter that Adam ostensibly wrote, which was held constant across conditions:

Happy Valentine's Day, my love.

I don't think I tell you enough just how much you mean to me. Being with you feels like breathing a little easier, like the world is a bit softer just because you're in it. You make the everyday feel special, and somehow you always know how to calm my nerves or make me laugh at just the right moment. I feel like myself with you – maybe even a better version of myself – and that's such a rare and beautiful thing.

I'm so grateful for you – for the way you listen, the way you love, the way you show up, even in the small ways. I hope you know that no matter what, I'm always in your corner. I can't wait to keep making memories together, whether we're off on some adventure or just curled up on the couch. I love you more than I can really put into words, but I promise I'll spend every day trying.

Yours,

Adam

Methods for Text Analysis in Study 5

To generate frequency lists for each experimental condition in Study 5, we created three documents containing the raw text submissions to the open-ended question “In your own words, describe how you feel about Adam and why”. Each raw text submission was paired with a numbered text ID column. The number of submissions was roughly equivalent across conditions: the control condition ($N = 196$), the tool outsourcing condition ($N = 215$), and the full outsourcing condition ($N = 202$).

All text processing was conducted using the Basic Unit-Transposable Text Experimentation Resource (BUTTER; Version 0.9.4.1; Boyd, 2019). To prepare the data, each CSV file was converted into a folder containing individual text files – one per submission – using two plugins: *Read Text from CSV* (Version 1.0.2) and *Save .txt Files to Folder* (Version 1.0.6). The settings for *Read Text from CSV* were as follows: file encoding = UTF-8, row identifier = ID, text column = Text, CSV delimiter = , and CSV quote = “.

To generate frequency lists, we first loaded the .txt files using the *Load .txt Files from Folder* plugin (Version 1.0.4). Tokenization was performed using the Twitter-Aware Tokenizer (Version 1.0.2), with the options *convert text to lowercase* and *reduce elongation* enabled to minimize superficial variation in tokens. We removed filler and function words using the *Remove Stop Words* plugin (Version 1.0.31), applying the default English stop word list.

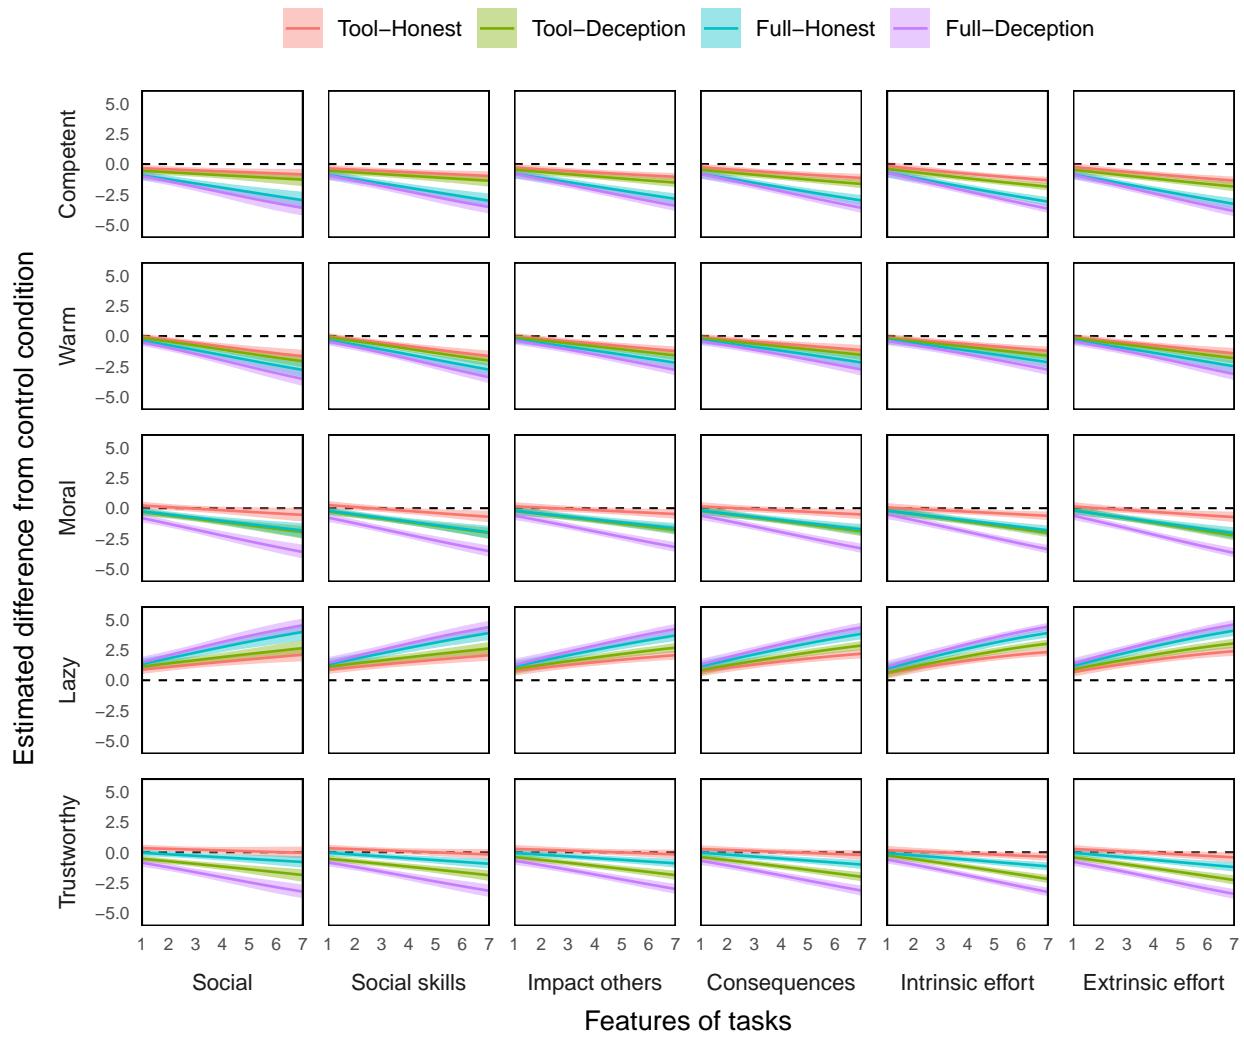
Frequency lists were created with the *Frequency List* plugin (Version 1.0.11). Settings included: unigram analysis ($N = 1$), omission of n-grams with frequency < 5 , exclusion of n-grams appearing in fewer than 0.1% of documents, filtering collocates by Normalized Pointwise Mutual Information (NPMI), and removal of collocates with metric values < 0.5 . Outputs were saved using the *Save Output to CSV* plugin (Version 1.0.5). This process was repeated separately for each condition folder.

For cross-condition comparison, we used the *Compare Frequencies* plugin (Version 1.1.02), retaining most default settings. The only modification was disabling the *Skip comparisons with 0 frequency values* option. This plugin calculates a range of comparative

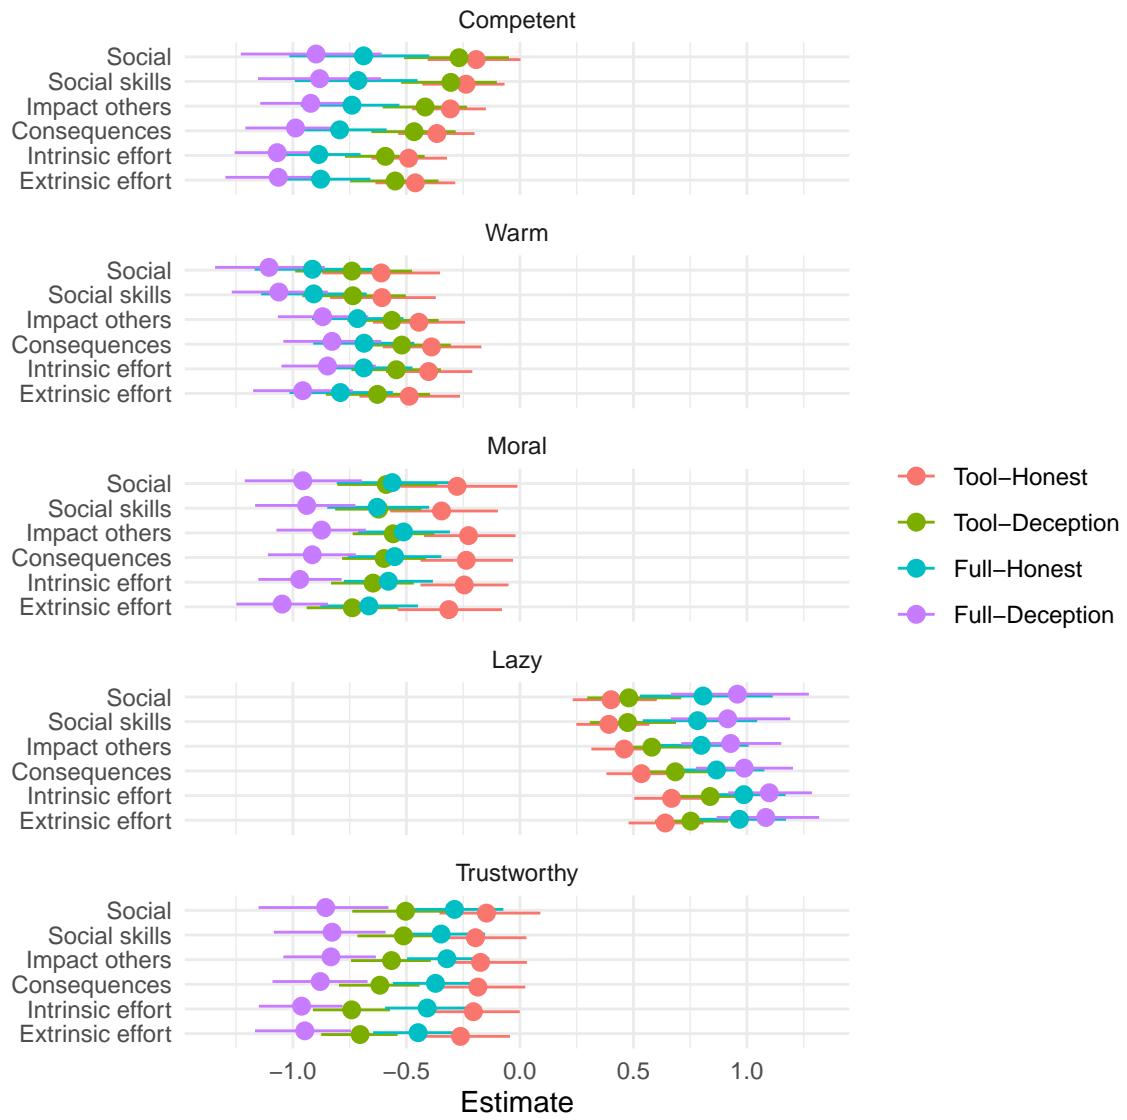
metrics, including log likelihood (LL), %DIFF, Bayes Information Criterion (BIC), relative risk (RRisk), log ratio, and odds ratio.

Following previous work (e.g., Rayson & Garside, 2000; Gregson et al., 2022), we interpret %DIFF as an indicator of effect size and direction. Frequentist statistical significance was determined using log likelihood values, with the following thresholds: $LL \geq 3.84$ ($p < .05$), $LL \geq 6.63$ ($p < .01$), $LL \geq 10.83$ ($p < .001$), and $LL \geq 15.13$ ($p < .0001$).

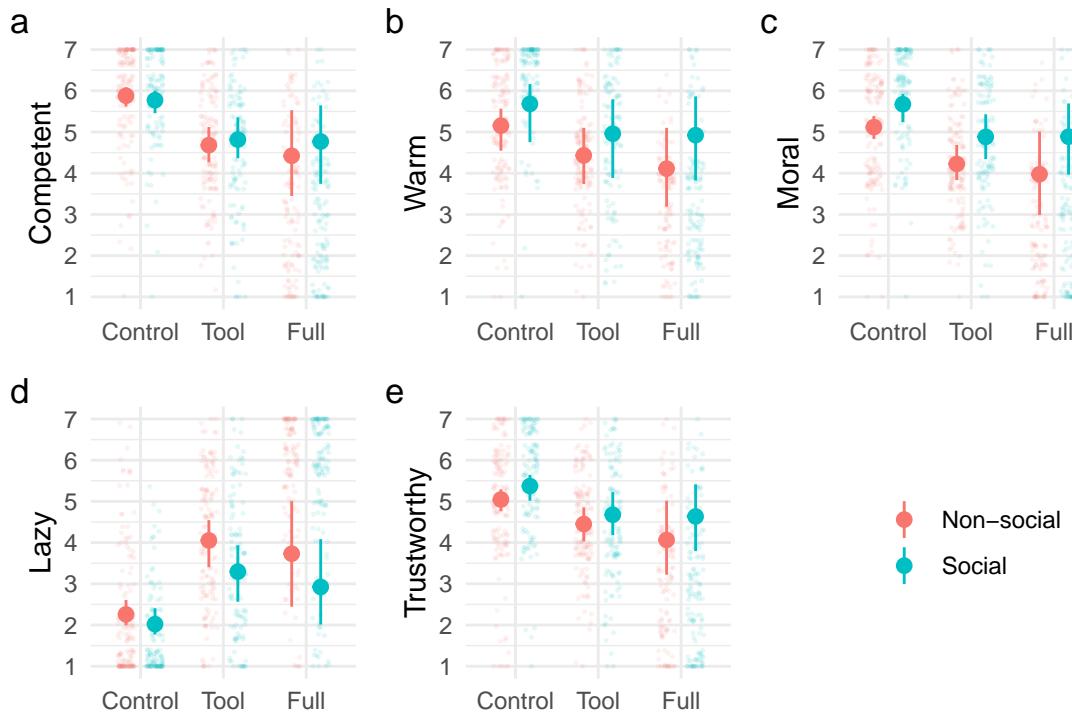
Supplementary Figures



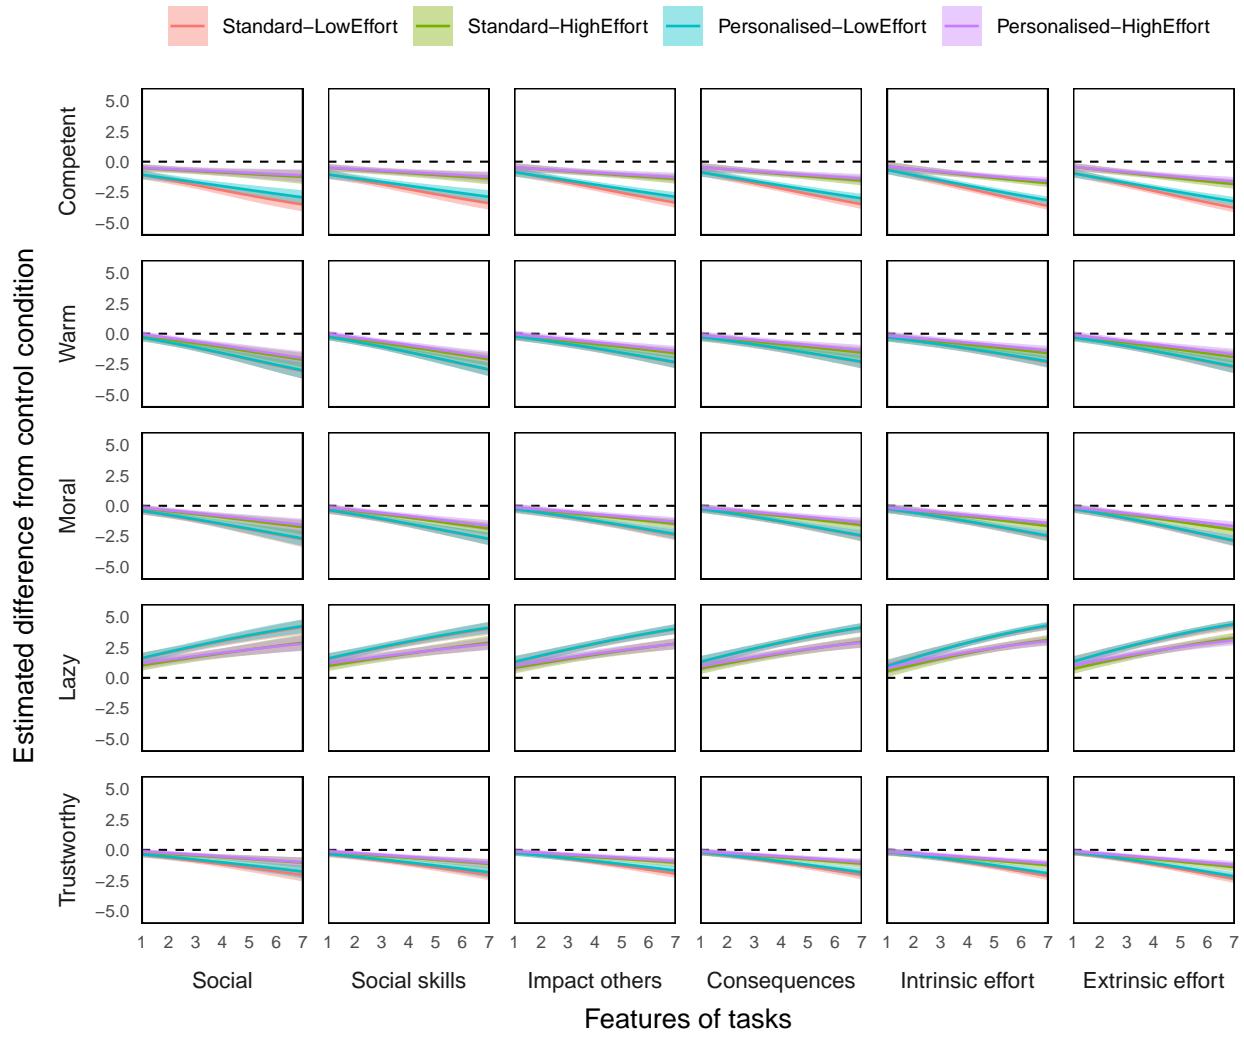
Supplementary Figure 1: The impact of task-specific features (e.g., being a social task) on the causal effects of outsourcing to AI compared to the control condition in Study 1. The y-axis reflects the estimated differences between the experimental conditions and the control condition (dashed line) on a 7-point Likert scale. Lines and shaded areas represent posterior medians and 95% credible intervals, respectively. The patterns indicate, for example, more negative effects of outsourcing on character evaluations for tasks that are rated as more social.



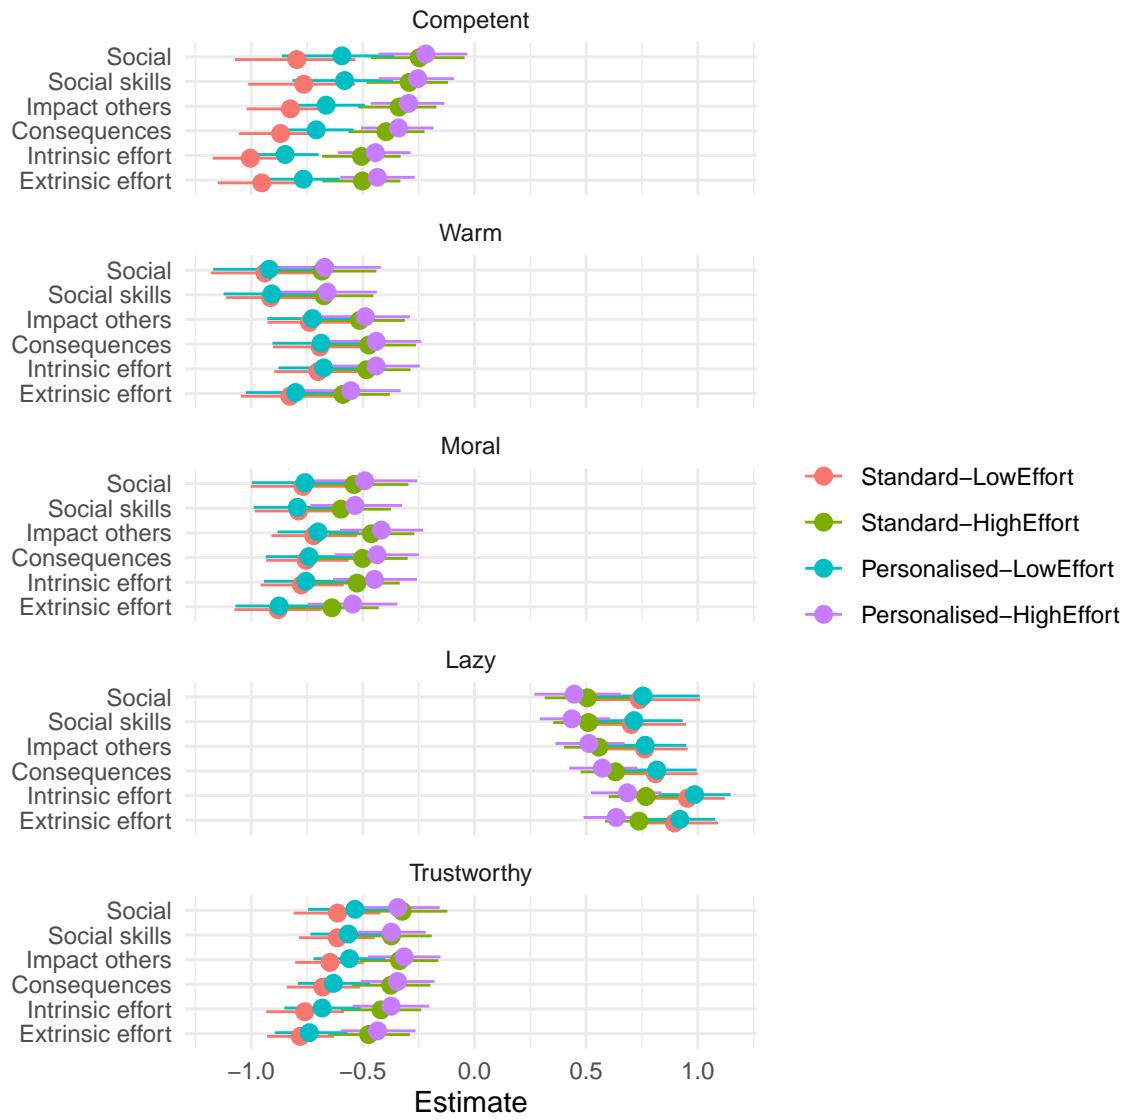
Supplementary Figure 2: Interaction parameters from models including task-specific features as moderators of the causal effects of AI outsourcing compared to the control condition in Study 1. Points and line ranges represent posterior medians and 95% credible intervals, respectively.



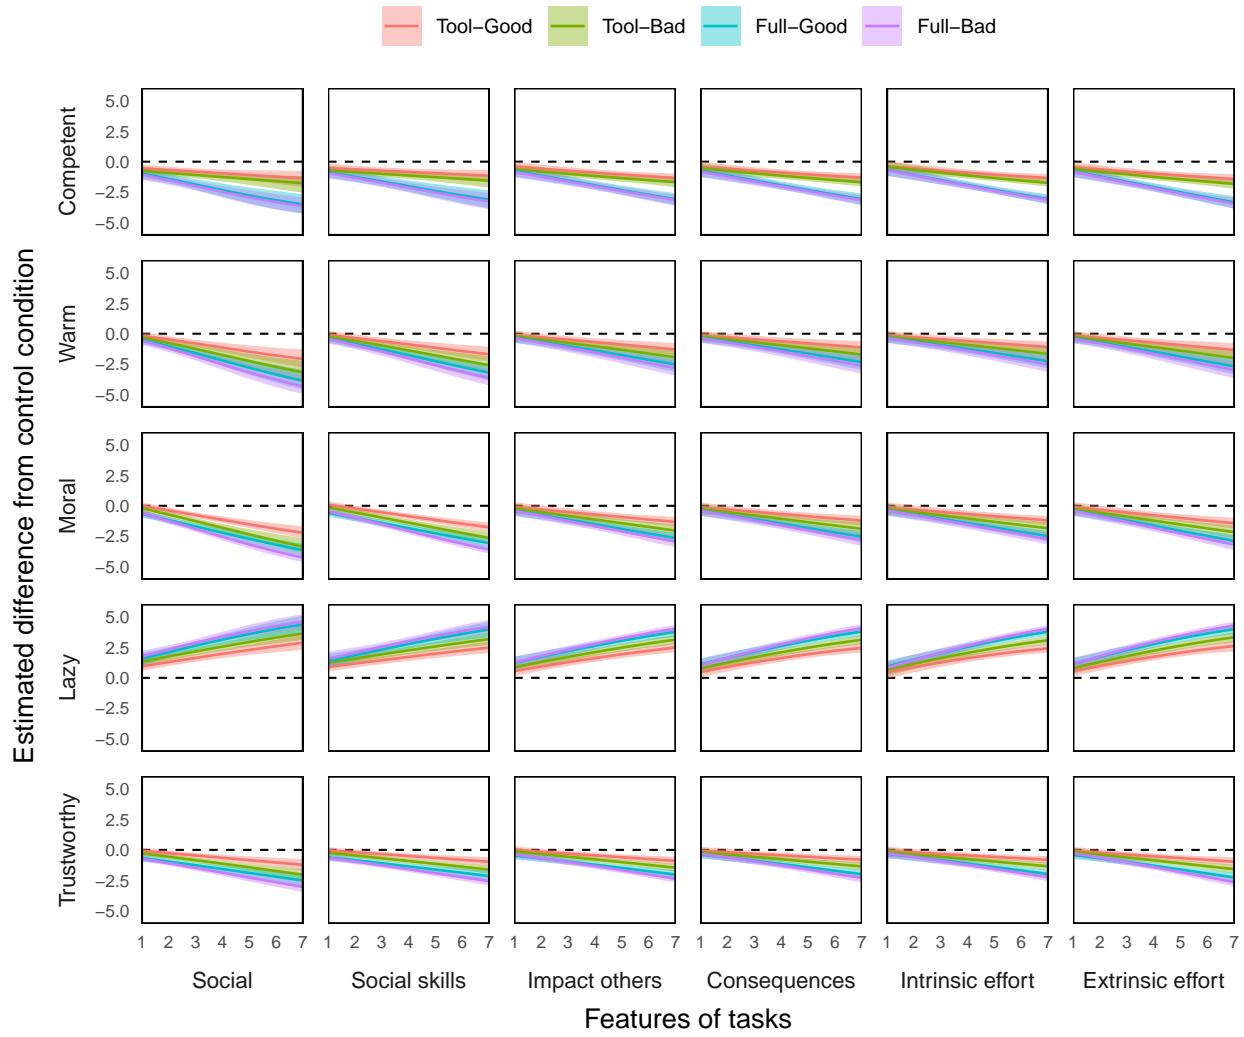
Supplementary Figure 3: Character evaluations in Study 2. Participants in the control condition, the tool outsourcing condition, and the full outsourcing condition evaluated the “other participant” on (a) competence, (b) warmth, (c) morality, (d) laziness, and (e) trustworthiness. Jittered points represent participant responses to the questions, split by whether the writing task was a non-social task (red) or a social task (blue). Point ranges are estimated marginal means from the fitted model, pooling over essay answers. Points and line ranges represent posterior medians and 95% credible intervals, respectively.



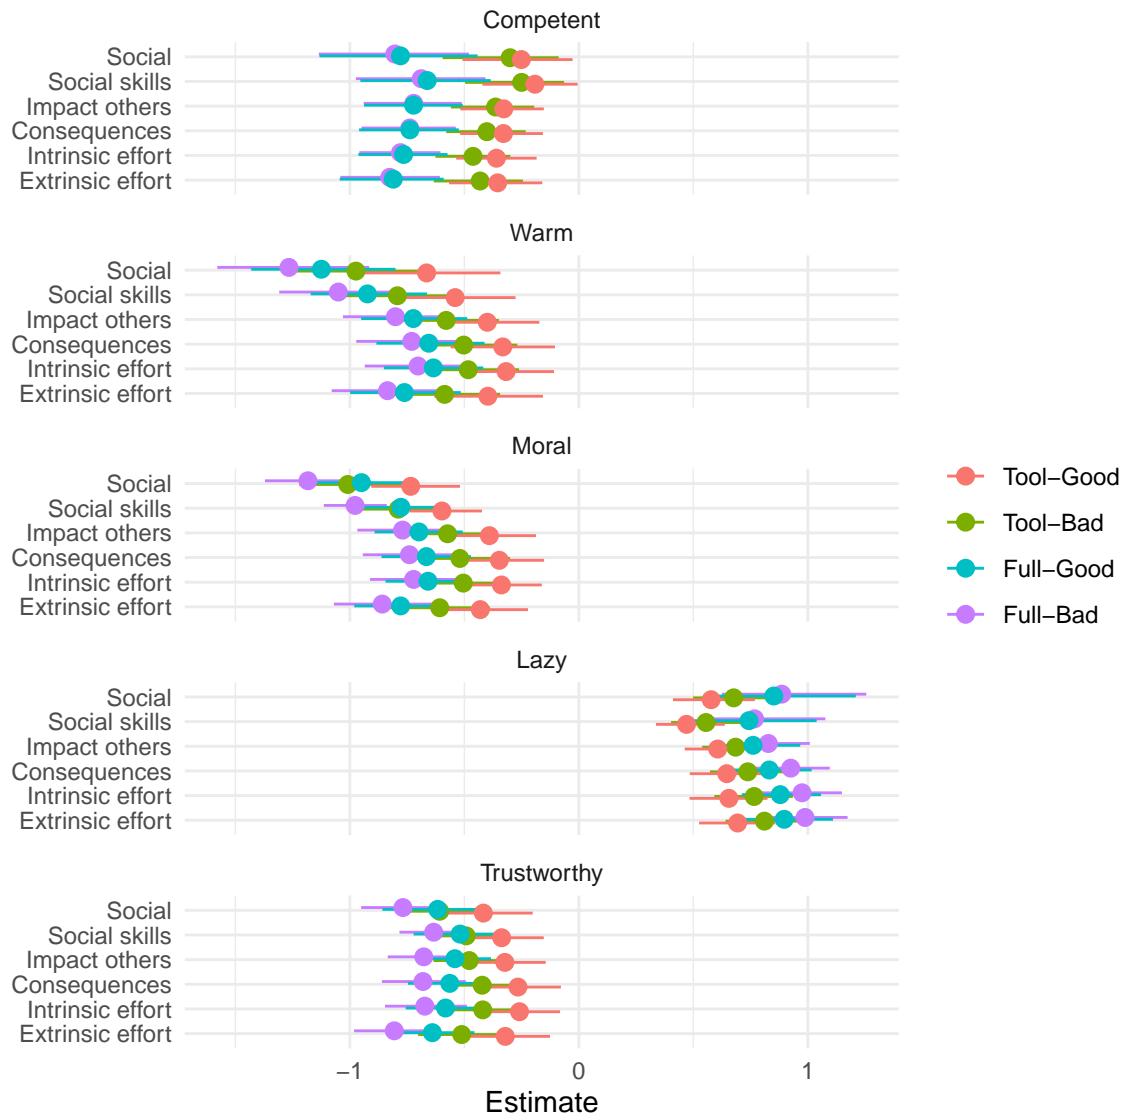
Supplementary Figure 4: The impact of task-specific features (e.g., being a social task) on the causal effects of outsourcing to AI compared to the control condition in Study 3. The y-axis reflects the estimated differences between the experimental conditions and the control condition (dashed line) on a 7-point Likert scale. Lines and shaded areas represent posterior medians and 95% credible intervals, respectively. The patterns indicate, for example, more negative effects of outsourcing on character evaluations for tasks that are rated as more social.



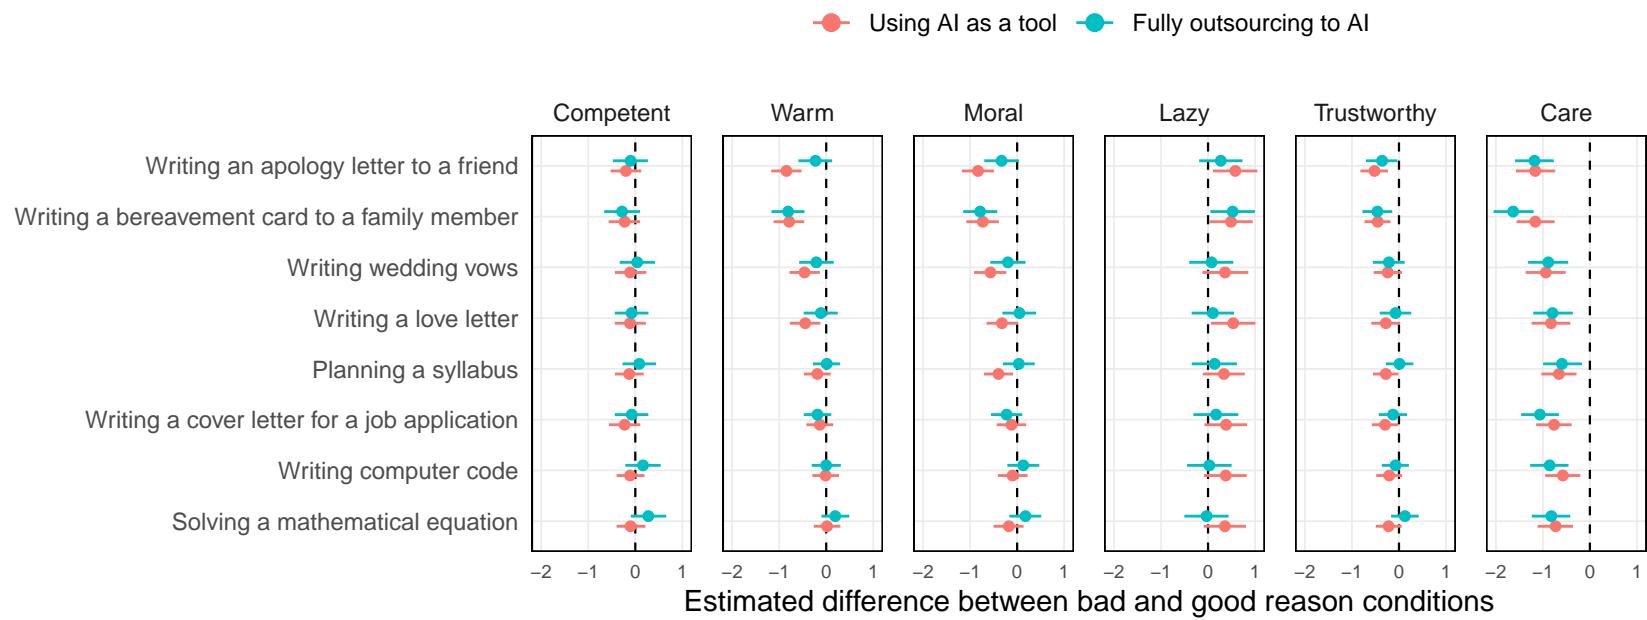
Supplementary Figure 5: Interaction parameters from models including task-specific features as moderators of the causal effects of AI outsourcing compared to the control condition in Study 3. Points and line ranges represent posterior medians and 95% credible intervals, respectively.



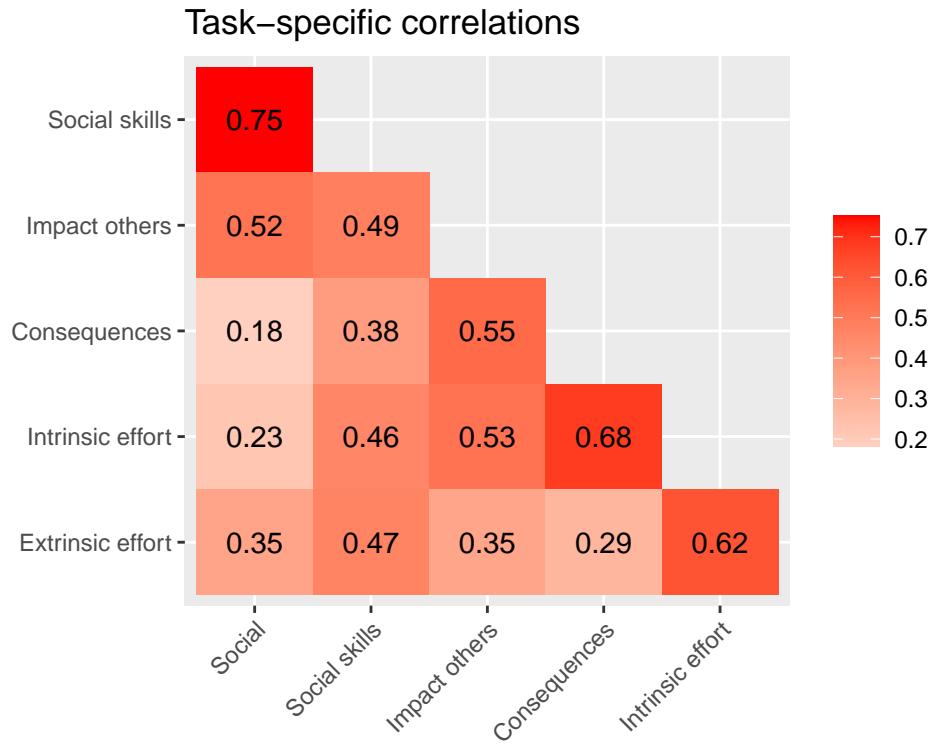
Supplementary Figure 6: The impact of task-specific features (e.g., being a social task) on the causal effects of outsourcing to AI compared to the control condition in Study 4. The y-axis reflects the estimated differences between the experimental conditions and the control condition (dashed line) on a 7-point Likert scale. Lines and shaded areas represent posterior medians and 95% credible intervals, respectively. The patterns indicate, for example, more negative effects of outsourcing on character evaluations for tasks that are rated as more social.




Supplementary Figure 7: Interaction parameters from models including task-specific features as moderators of the causal effects of AI outsourcing compared to the control condition in Study 4. Points and line ranges represent posterior medians and 95% credible intervals, respectively.




Supplementary Figure 8: Variation in the effect of reasons across tasks in Study 4. Tasks are ordered from most social (top) to least social (bottom) according to ratings from a pilot study. Point ranges are differences in marginal means on a 7-point Likert scale between the “bad reason” and “good reason” conditions, split by outsourcing type. Points and ranges represent posterior medians and 95% credible intervals, respectively.




Supplementary Figure 9: Model-estimated task-specific correlations between all six questions in the first pilot study. Values are posterior median correlations. A positive correlation indicates that tasks that are rated highly on one question tend to be rated highly on another question.




Supplementary Figure 10: Model-estimated means for the question “Is this a social task?” across all 20 tasks in the first pilot study. Grey points represent participant responses to the question, jittered for easier viewing. Black points are estimated means from the fitted model, pooling over participants. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.




Supplementary Figure 11: Model-estimated means for the question “Does this task require social skills?” across all 20 tasks in the first pilot study. Grey points represent participant responses to the question, jittered for easier viewing. Black points are estimated means from the fitted model, pooling over participants. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.




Supplementary Figure 12: Model-estimated means for the question “Does this task impact other people?” across all 20 tasks in the first pilot study. Grey points represent participant responses to the question, jittered for easier viewing. Black points are estimated means from the fitted model, pooling over participants. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.




Supplementary Figure 13: Model-estimated means for the question “How important are the consequences of this task?” across all 20 tasks in the first pilot study. Grey points represent participant responses to the question, jittered for easier viewing. Black points are estimated means from the fitted model, pooling over participants. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.



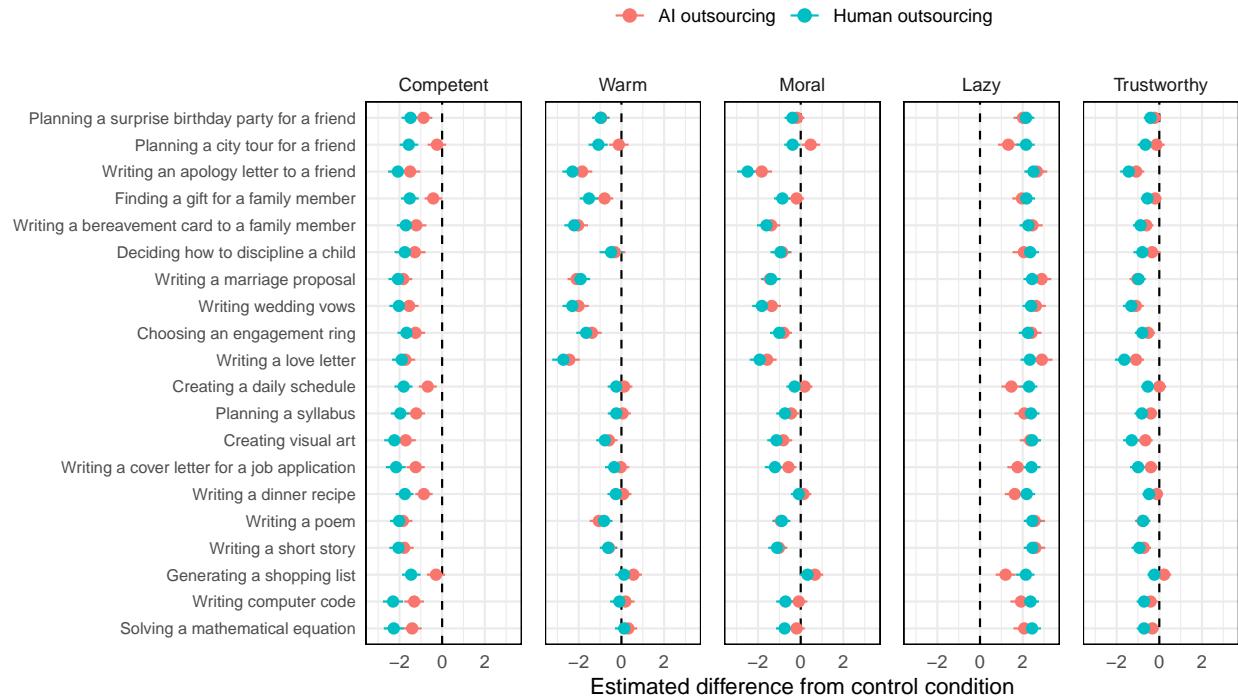
Supplementary Figure 14: Model-estimated means for the question “How important is it that effort goes into this task?” across all 20 tasks in the first pilot study. Grey points represent participant responses to the question, jittered for easier viewing. Black points are estimated means from the fitted model, pooling over participants. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.



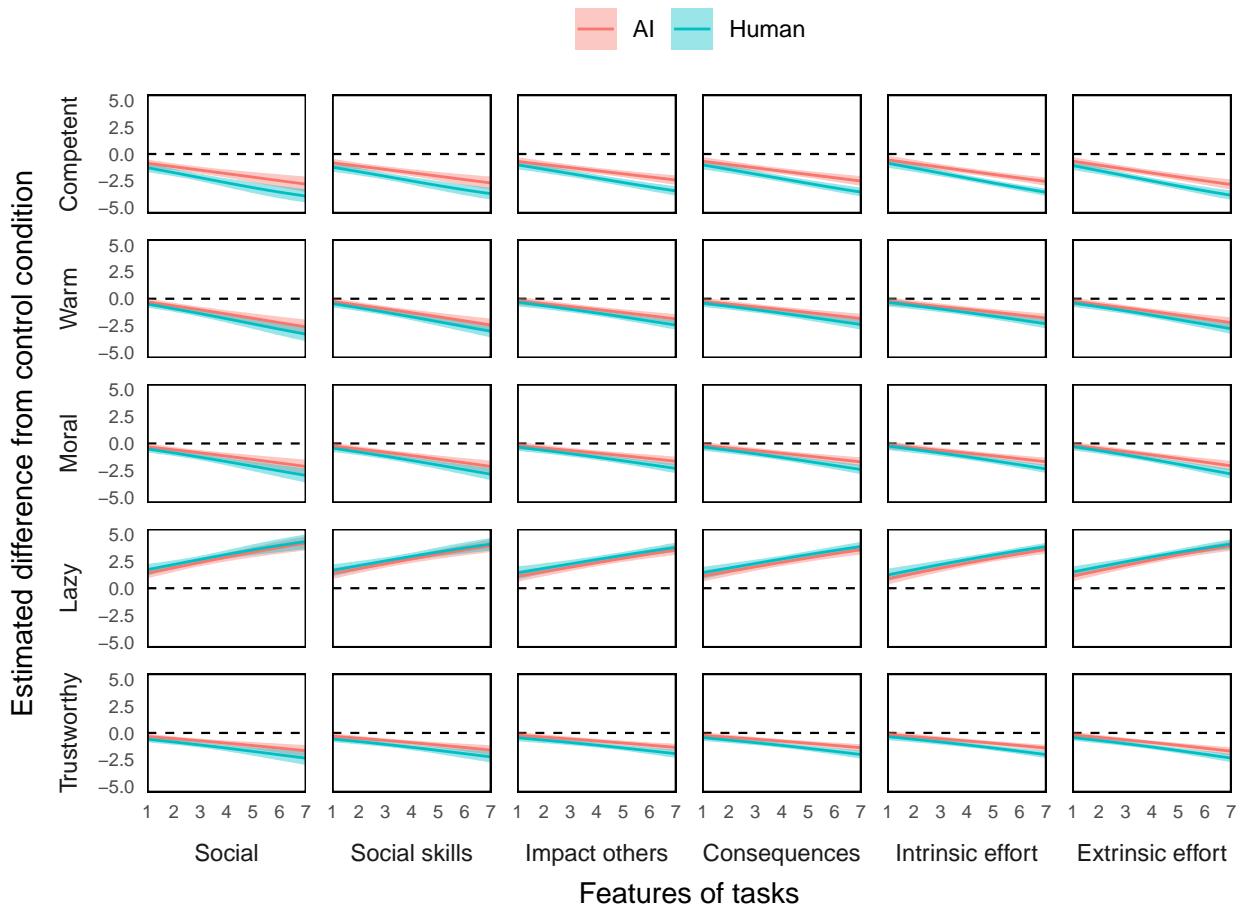
Supplementary Figure 15: Model-estimated means for the question “How important is it that others see the effort that goes into this task?” across all 20 tasks in the first pilot study. Grey points represent participant responses to the question, jittered for easier viewing. Black points are estimated means from the fitted model, pooling over participants. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.



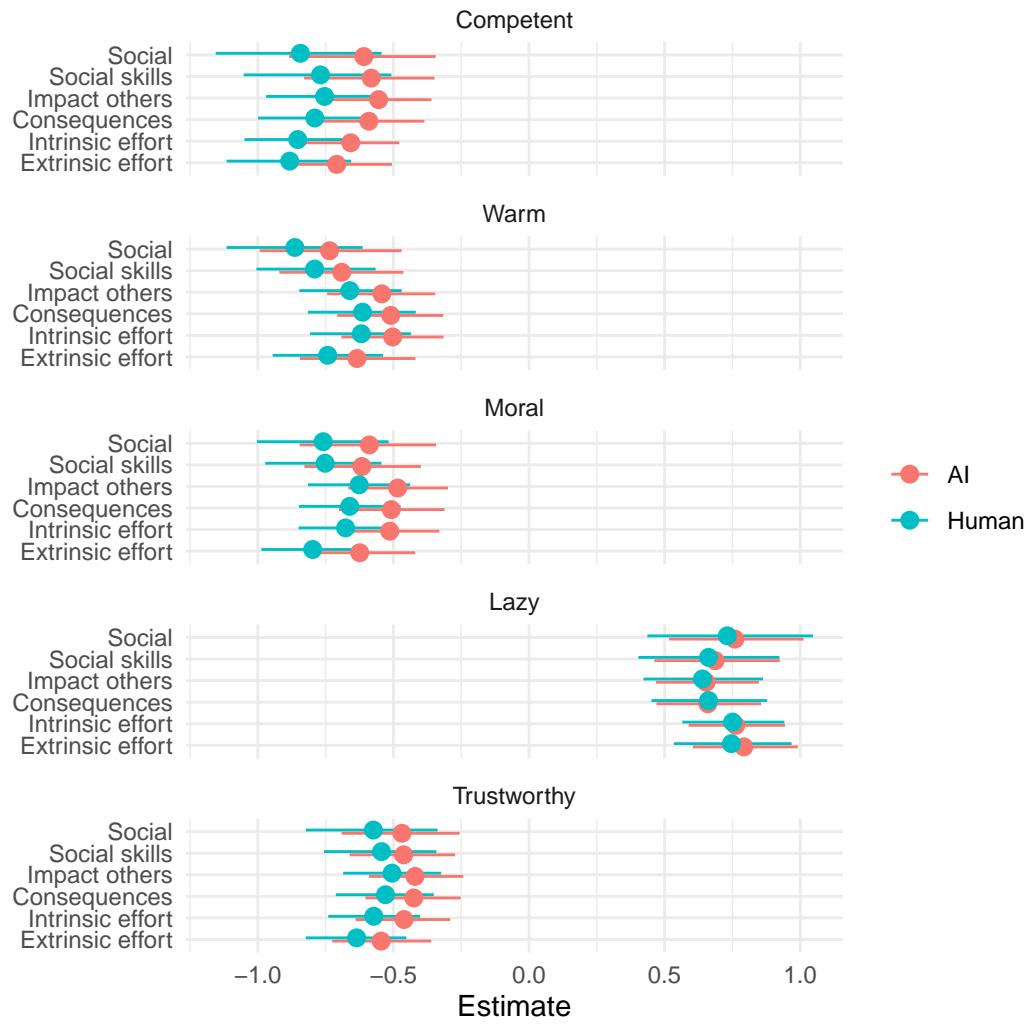
Supplementary Figure 16: Character evaluations in the second pilot study. Participants in the control condition, the AI outsourcing condition, and the human outsourcing condition evaluated people in the scenarios on (a) competence, (b) warmth, (c) morality, (d) laziness, and (e) trustworthiness. Coloured points represent participant responses to the questions, jittered for easier viewing. Black points are estimated marginal means from the fitted model, pooling over participants and tasks. Black points and line ranges represent posterior medians and 95% credible intervals, respectively.



Supplementary Figure 17: Variation in the effects of outsourcing across tasks in the second pilot study. Tasks are ordered from most social (top) to least social (bottom) according to ratings from the first pilot study. Point ranges are differences in marginal means on a 7-point Likert scale for the AI outsourcing condition (red) and the human outsourcing condition (blue) compared to the control condition. Points and ranges represent posterior medians and 95% credible intervals, respectively.



Supplementary Figure 18: The impact of task-specific features (e.g., being a social task) on the causal effects of outsourcing to AI (red) and humans (blue) compared to the control condition in the second pilot study. The y-axis reflects the estimated differences between the experimental conditions and the control condition (dashed line) on a 7-point Likert scale. Lines and shaded areas represent posterior medians and 95% credible intervals, respectively. The patterns indicate, for example, more negative effects of outsourcing on character evaluations for tasks that are rated as more social.



Supplementary Figure 19: Interaction parameters from models including task-specific features as moderators of the causal effects of AI outsourcing (red) and human outsourcing (blue) compared to the control condition in the second pilot study. Points and line ranges represent posterior medians and 95% credible intervals, respectively.

Supplementary Tables

Supplementary Table 1: Tasks included in the studies.

Task	Pilot Study 1	Pilot Study 2	Study 1	Study 2	Study 4
Writing wedding vows	✓	✓	✓	✓	✓
Writing a love letter	✓	✓	✓	✓	✓
Writing a marriage proposal	✓	✓	✓	✓	
Choosing an engagement ring	✓	✓	✓		
Finding a gift for a family member	✓	✓	✓		
Deciding how to discipline a child	✓	✓	✓		
Writing a bereavement card to a family member	✓	✓	✓	✓	✓
Writing an apology letter to a friend	✓	✓	✓	✓	✓
Planning a city tour for a friend	✓	✓	✓	✓	
Planning a surprise birthday party for a friend	✓	✓	✓	✓	
Writing a cover letter for a job application	✓	✓	✓	✓	✓
Writing computer code	✓	✓	✓	✓	✓
Solving a mathematical equation	✓	✓	✓	✓	✓
Planning a syllabus	✓	✓	✓	✓	✓
Writing a short story	✓	✓	✓	✓	
Writing a poem	✓	✓	✓	✓	
Creating visual art	✓	✓	✓		
Creating a daily schedule	✓	✓	✓	✓	
Generating a shopping list	✓	✓	✓	✓	
Writing a dinner recipe	✓	✓	✓	✓	

Supplementary Table 2: Pre-generated essay answers to the social prompt in Study 2.

Answer	Text
Father	<p>My dad is one of the most important people in my life. He's always been someone I look up to and rely on. Throughout my whole life, he's been there to guide me, teach me, and support me in everything I do. What makes my dad special is how much he cares about our family. He works hard every day to make sure we have what we need, but no matter how busy he is, he always makes time for us. My dad is emotionally strong. Even though he doesn't show his emotions a lot, I can tell how much he cares by how much he does for us. When things get hard, he stays calm and steady, and that helps me feel better. One of my favorite things about my dad is how much he loves to teach. He knows so much and is always happy to share what he knows. He explains things in a way that makes sense and is easy to understand. I also love my dad's sense of humor. He always knows how to make me laugh with a joke or a funny story. His laughter makes everything feel lighter and happier. My dad has taught me so much about working hard, being kind, and staying strong when life is tough. I'm so thankful for everything he's done for me, and I'm proud to have him as my dad!</p>
Sister	<p>My sister is one of the most important people in my life. She is special because she always supports me. She has a way of making me feel confident, even when I'm unsure of myself. Whenever I'm scared to try something new, she's the first to remind me of what I can do. Her belief in me helps me believe in myself. My sister also has a really kind heart. She always thinks about others and does her best to help. She's always putting others first, whether it's being there for a friend or helping out with family. Her kindness is something I look up to and try to follow. Another thing I love about my sister is how funny she is. She has a great sense of humor and always knows how to make people laugh, even in serious moments. If I'm ever feeling down, she can cheer me up with a joke or a funny story. Her laughter makes everything feel lighter and happier. What I admire most about my sister is how strong she is. She's faced tough times but never lets them hold her back. Her strength gives me courage to keep going when life gets hard. My sister is more than just a family member — she's my role model and my rock!</p>
Friend	<p>My best friend is one of the most amazing people I know. She's someone I can count on no matter what. What makes her so special is her kindness. She always makes people feel important and cared for. Whether it's helping someone she just met or being there for her friends, she's the first to offer support. She never hesitates to help me, whether I'm upset or just having a bad day. She also has a great sense of humor that can cheer anyone up. She finds ways to laugh about even the smallest things, and her laugh is so contagious! Her laughter makes everything feel lighter and happier. What I admire most about her is how strong she is. Life hasn't always been easy for her, but she never gives up. She stays calm and keeps going, no matter what happens. Watching her face challenges in adulthood has taught me to be brave and not let hard times hold me back. My best friend has shown me what it means to be loyal, caring, and strong. I feel so lucky to have her in my life. I try to be as good of a friend to her as she is to me. She inspires me to be a better person!</p>

Supplementary Table 3: Pre-generated essay answers to the non-social prompt in Study 2.

Answer	Text
The Hobbit	I will focus on describing the book “The Hobbit” by Tolkien. The Hobbit is a fantasy adventure story about Bilbo Baggins. Bilbo is a quiet hobbit who lives in the Shire. His life changes when Gandalf the wizard and a company of dwarves ask him to join their quest to take back treasure stolen by a dragon. At the beginning of the journey, Bilbo and the dwarves are nearly eaten by trolls, but Gandalf saves them. Then later, in the Misty Mountains, Bilbo meets a creature called Gollum and finds a magical ring that makes him invisible. This ring later becomes very important in “The Lord of the Rings”. As they travel, the group fights goblins, giant spiders in Mirkwood forest, and they get captured by Wood-elves. Bilbo shows his bravery by saving the group several times. Finally, they reach the Lonely Mountain where the dragon Smaug lives. Bilbo sneaks into the dragon’s lair and finds a weak spot in Smaug’s armor. The dragon gets angry and attacks the nearby town by a lake. Eventually, Smaug is killed. With the dragon dead, humans, elves, and dwarves all want the treasure. This leads to the “Battle of the Five Armies”. Tolkien doesn’t describe the battle in too much detail, but we later learn that the leader of the dwarves Thorin has fought bravely and died from his wounds. At the end of the story, Bilbo returns home to the Shire, richer and wiser from his adventure. He is happy to be back in his quiet life, and sets out to write a book of his adventures - which sets the stage for the sequel, The Lord of the Rings.
Buffy the Vampire Slayer	I will focus on describing the TV show “Buffy the Vampire Slayer”. Buffy is a TV show that completely flips the script on traditional high school dramas and supernatural horror. It’s about a teenager, Buffy Summers, who’s tasked with being the Slayer – basically a chosen one who hunts vampires and other demons. But what sets the show apart is how Buffy struggles to balance her responsibility with the regular teenage experience. She’s not just fighting creatures of the night, she’s also balancing school and friendships at the same time. One of the most striking things about Buffy is how layered the characters are. Buffy is tough and witty, but she’s also vulnerable. She’s faced with loss, guilt, and trying to make sense of her life outside of the supernatural chaos. And then there’s her team. Willow is the nerdy, sweet heart of the group, Xander is the funny loyal friend, and Giles (Buffy’s Watcher) is the stern mentor who’s also loving. Each character feels real, with their own flaws and growth arcs. The show has this incredible ability to mix humor, heart, and horror seamlessly. The dialogue is sharp and full of clever pop culture references. Yet, the writing isn’t afraid to get serious, exploring themes like trauma and growing up. The monsters Buffy faces often mirror real-life challenges, making the stakes feel personal. I love Buffy. It’s a show that’s smart and emotional, blending witty banter with moments of real depth. It’s got a cult following for a reason!
Titanic	I will focus on describing the film “Titanic”. The genre is a mix of romance, disaster, and historical tragedy. The film tells the love story of Jack and Rose, two passengers from different social classes aboard the passenger ship Titanic. Jack is a poor artist, but he manages to win a ticket to the ship’s maiden voyage. Rose is a young upper-class woman who is feeling trapped in her engagement to her fiance. Jack and Rose cross paths on the ship, and they fall in love. The film balances the spectacle of the ship’s design and atmosphere with the tension that gradually builds as the audience knows what fate awaits. The Titanic sails into the icy waters of the Atlantic and strikes an iceberg. Chaos immediately erupts. The film allows viewers to experience the terror, confusion, and heartbreak of the tragedy, showcasing both personal stories and the broader catastrophe. At its core, the film is a romance. But Titanic also touches on themes of class and fate. It highlights the disparity between the elite and the working-class passengers who are doomed to different fates. The film also explores the sense of inevitability that comes with knowing the ship’s doom. The most iconic scene from the film is arguably the scene where Jack and Rose stand together at the bow, arms outstretched. They seem free, but the scene also foreshadows the devastating crash to come. The film is truly heartbreakingly tragic!

Supplementary Table 4: Reading times and comprehension rates for the essay answers in Study 2. Expected reading times were calculated based on an estimated reading speed of 275 words per minute. Comprehension rates are the percentage of participants who answered the comprehension question correctly.

Prompt	Answer	Number of words	Expected reading time (secs)	Average reading time (secs)	Comprehension (%)
Social	Father	234	51.05	47.25	100.00
	Friend	211	46.04	49.91	98.50
	Sister	218	47.56	50.25	99.21
Non-social	Buffy	251	54.76	62.46	100.00
	Hobbit	278	60.65	65.97	99.28
	Titanic	239	52.15	63.90	100.00

Supplementary Table 5: Percentage of participants in Study 2 who passed the manipulation check and reported that they believed the manipulation, split by condition.

Condition	Pass manipulation check (%)	Believe manipulation (%)
Control	98.10	71.10
Tool outsourcing	96.11	77.99
Full outsourcing	100.00	86.07

Supplementary Table 6: Pairwise contrasts for character evaluations in Study 2. Numbers reflect differences in marginal means on a 7-point Likert scale. Estimates are pooled over essay answers. The bottom rows represent the interactions between outsourcing type and task type. Main numbers are posterior medians, numbers in the square brackets are 95% credible intervals.

	Response				
	Competent	Warm	Moral	Lazy	Trustworthy
Effect of outsourcing type					
Task type = Social					
Tool Social - Control Social	-0.97 [-1.33 -0.45]	-0.74 [-1.25 -0.12]	-0.81 [-1.16 -0.30]	1.31 [0.59 1.76]	-0.71 [-1.07 -0.21]
Full Social - Control Social	-0.99 [-1.96 -0.17]	-0.75 [-1.55 -0.03]	-0.78 [-1.63 -0.04]	0.87 [0.05 1.96]	-0.73 [-1.50 0.00]
Full Social - Tool Social	-0.03 [-1.13 0.92]	-0.03 [-1.06 0.92]	0.01 [-0.95 0.87]	-0.41 [-1.38 0.84]	-0.03 [-0.91 0.77]
Task type = Non-social					
Tool Non-social - Control Non-social	-1.21 [-1.55 -0.76]	-0.74 [-1.23 -0.10]	-0.91 [-1.28 -0.42]	1.82 [1.13 2.26]	-0.60 [-0.98 -0.17]
Full Non-social - Control Non-social	-1.48 [-2.41 -0.37]	-1.07 [-1.89 -0.11]	-1.15 [-2.11 -0.11]	1.48 [0.21 2.75]	-1.00 [-1.81 -0.04]
Full Non-social - Tool Non-social	-0.28 [-1.30 0.91]	-0.33 [-1.31 0.72]	-0.25 [-1.28 0.86]	-0.31 [-1.73 1.03]	-0.40 [-1.28 0.62]
Effect of task type					
Control Social - Control Non-social	-0.11 [-0.39 0.17]	0.57 [-0.16 0.99]	0.57 [0.15 0.88]	-0.24 [-0.55 0.11]	0.33 [-0.02 0.65]
Tool Social - Tool Non-social	0.13 [-0.30 0.61]	0.55 [-0.37 1.32]	0.66 [0.11 1.15]	-0.76 [-1.35 -0.15]	0.22 [-0.28 0.76]
Full Social - Full Non-social	0.34 [-0.37 1.07]	0.83 [-0.16 1.68]	0.92 [0.13 1.68]	-0.82 [-1.72 0.12]	0.56 [-0.14 1.27]
Interaction effect					
Interaction: Tool - Control	0.24 [-0.18 0.69]	-0.01 [-0.52 0.59]	0.10 [-0.32 0.57]	-0.52 [-1.06 0.00]	-0.11 [-0.53 0.38]
Interaction: Full - Control	0.46 [-0.21 1.13]	0.29 [-0.37 0.94]	0.35 [-0.33 1.05]	-0.57 [-1.41 0.28]	0.24 [-0.41 0.90]
Interaction: Full - Tool	0.21 [-0.60 1.02]	0.29 [-0.58 1.14]	0.25 [-0.53 1.09]	-0.06 [-1.03 0.99]	0.35 [-0.42 1.14]

Supplementary Table 7: Pairwise comparisons of word frequencies between conditions. LL = log likelihood.

Word	Control Freq.	Tool Freq.	Full Freq.	%DIFF Full vs Control	LL Full vs Control	%DIFF Tool vs Control	LL Tool vs Control	%DIFF Full vs Tool	LL Full vs Tool
Lazy	0	46	82	14138.18	97.16	6061.29	42.55	131.09	21.76
Genuine	36	11	12	-71.06	16.19	-79.54	25.91	41.42	0.69
Loves	9	9	0	-95.18	10.50	-33.03	0.72	-92.80	7.21
Romantic	0	7	0	-13.18	0.00	837.59	4.42	-90.74	5.16
Thoughtful	13	0	0	-96.66	16.27	-97.42	19.99	29.64	0.02
Caring	35	12	0	-98.76	49.01	-77.04	22.85	-94.60	10.36

Supplementary Table 8: Pairwise contrasts in the second pilot study. Numbers reflect differences in marginal means on a 7-point Likert scale, pooling over participants and tasks. Main numbers are posterior medians, numbers in the square brackets are 95% credible intervals.

	Response				
	Competent	Warm	Moral	Lazy	Trustworthy
AI - Control	-0.86 [-1.16 -0.55]	-0.30 [-0.57 -0.01]	-0.23 [-0.47 0.02]	1.80 [1.42 2.14]	-0.26 [-0.44 -0.05]
Human - Control	-1.68 [-1.98 -1.35]	-0.59 [-0.87 -0.28]	-0.54 [-0.80 -0.28]	2.26 [1.90 2.58]	-0.59 [-0.81 -0.39]
Human - AI	-0.81 [-1.15 -0.46]	-0.29 [-0.61 0.04]	-0.32 [-0.63 -0.02]	0.46 [0.04 0.90]	-0.33 [-0.58 -0.10]

Supplementary References

Boyd, R. L. (2019). BUTTER: Basic unit-transposable text experimentation resource. Available from <https://www.butter.tools/>

Gregson, R., Piazza, J., & Boyd, R. L. (2022). ‘Against the cult of veganism’: Unpacking the social psychology and ideology of anti-vegans. *Appetite*, 178, 106143.
doi:[10.1016/j.appet.2022.106143](https://doi.org/10.1016/j.appet.2022.106143)

Rayson, P., & Garside, R. (2000). Comparing corpora using frequency profiling. *Proceedings of the Workshop on Comparing Corpora - Volume 9*, 1–6. Presented in Hong Kong.
doi:[10.3115/1117729.1117730](https://doi.org/10.3115/1117729.1117730)